

Advanced Topics in Multicore Architecture & Software Systems

Cache Coherence

Administration

- HW#2 out tonight, due **December 11th**
 - 1 question on this lecture
 - 1 question on lecture #4
 - Max 3 pages
 - **Can be submitted in pairs**

Orientation

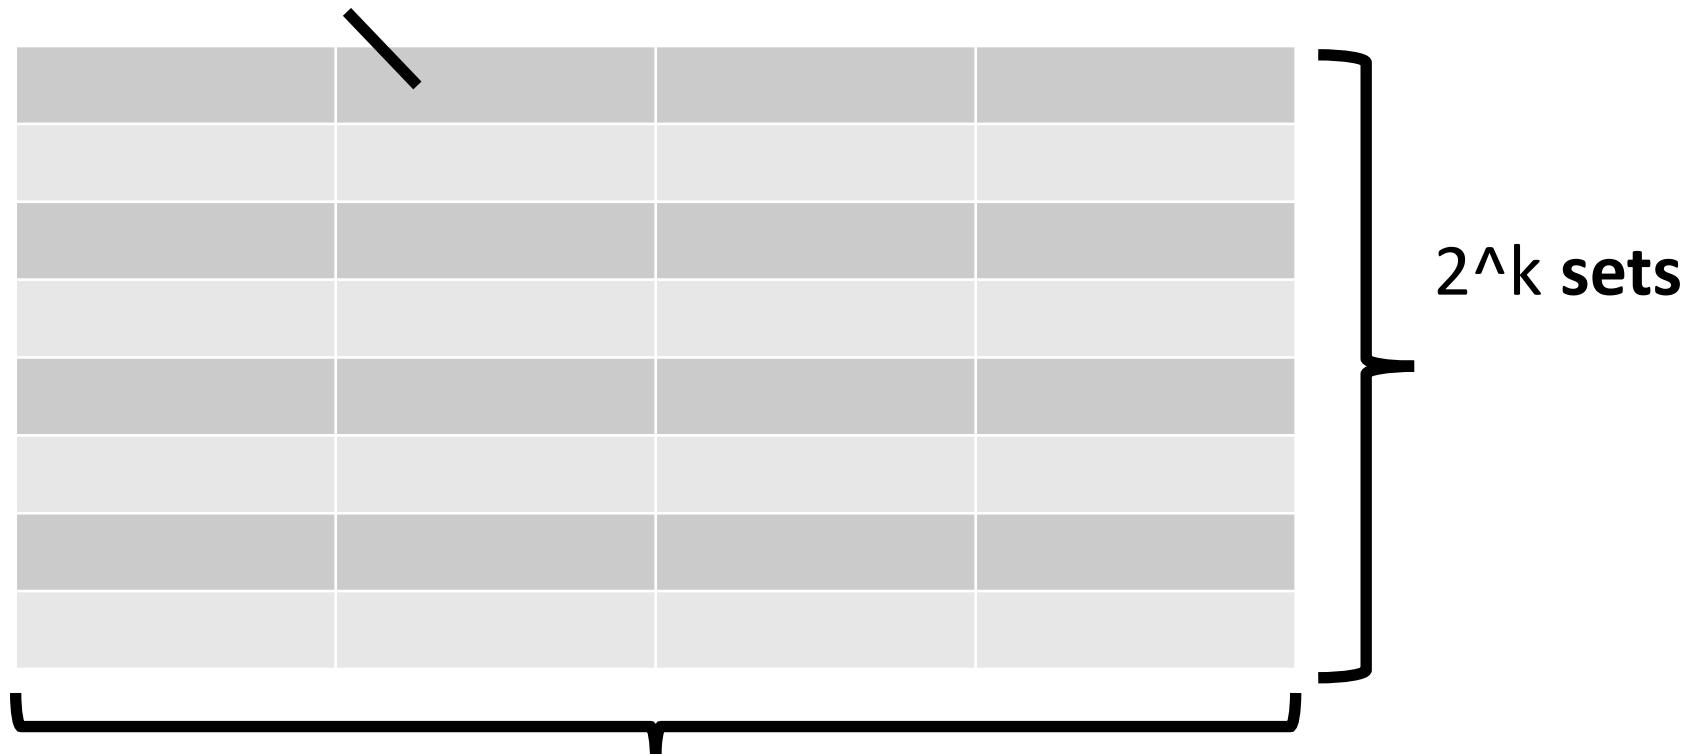
First part of the course: background

- Processors and OoO/speculative execution
- Linearizability
- **Cache coherence**
 - ⇒ Next week: Algorithmic interlude
- Memory consistency (PL & processors)

Background: CPU cache

Set-associative caches

line: unit of storage (usually 64 bytes)



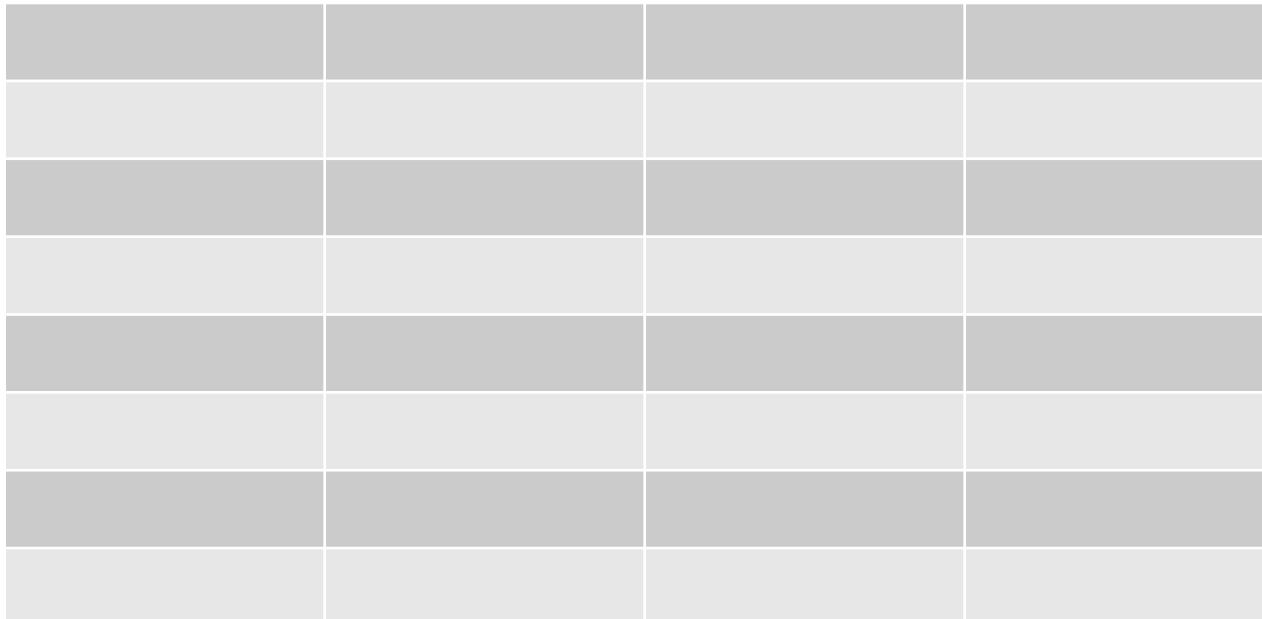
each set composed of 2^m ways

Background: CPU cache

physical
(physically
indexed cache)

Memory op: addr maps to set

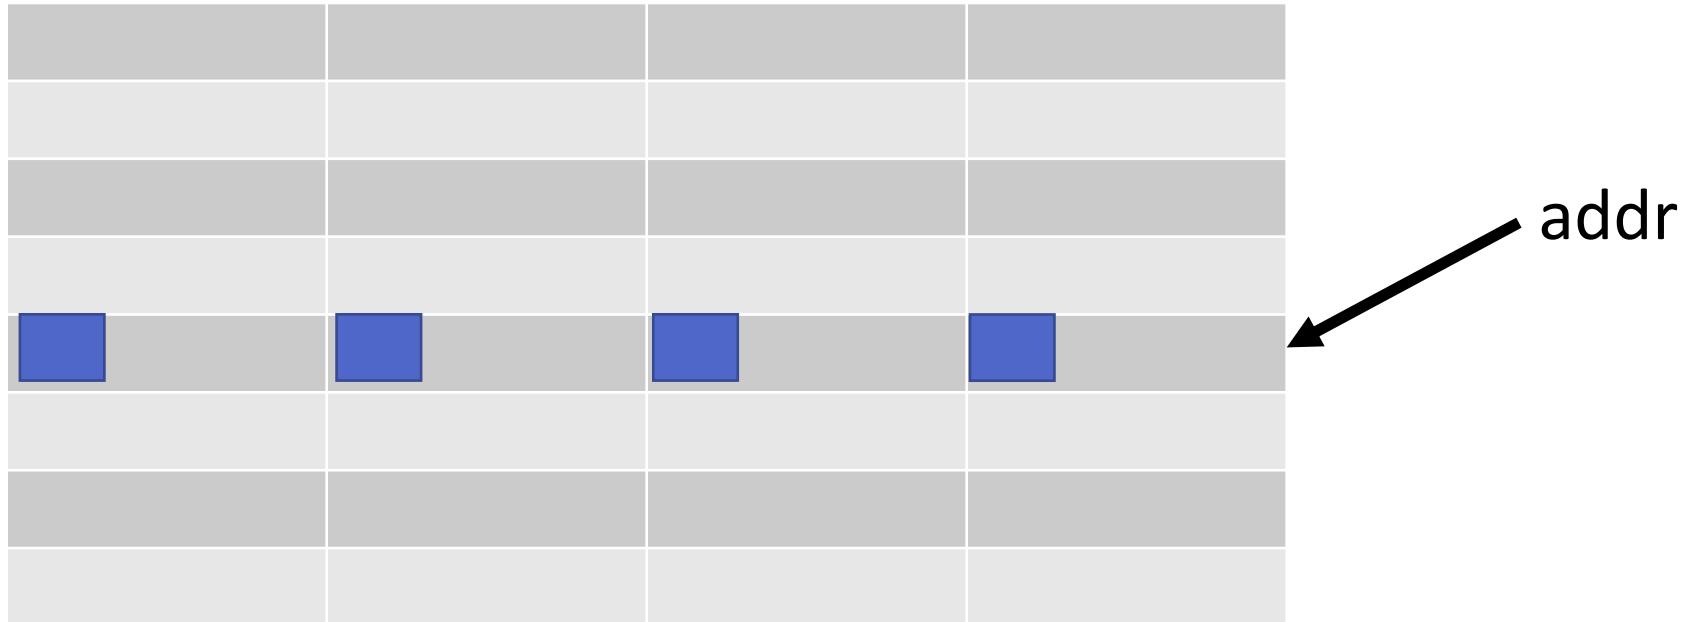
Typically: $(\text{addr} \gg \log(\text{line size})) \bmod \#\text{sets}$



addr

Background: CPU cache

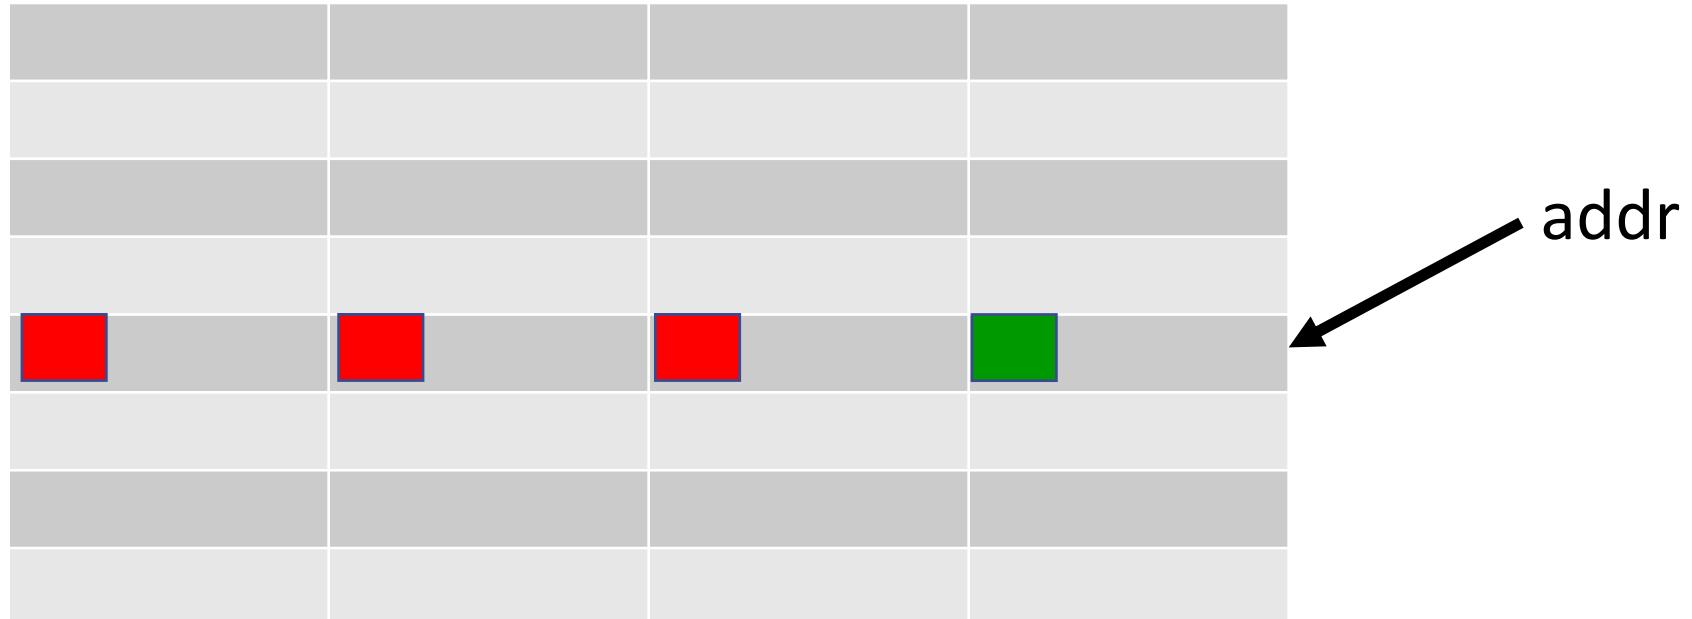
Tags in each way searched in parallel to figure out if addr is stored in cache or not



Tag contains high addr bits ($\text{addr} \gg \log(\text{line size})$)
Again, assume physical (physically tagged)

Background: CPU cache

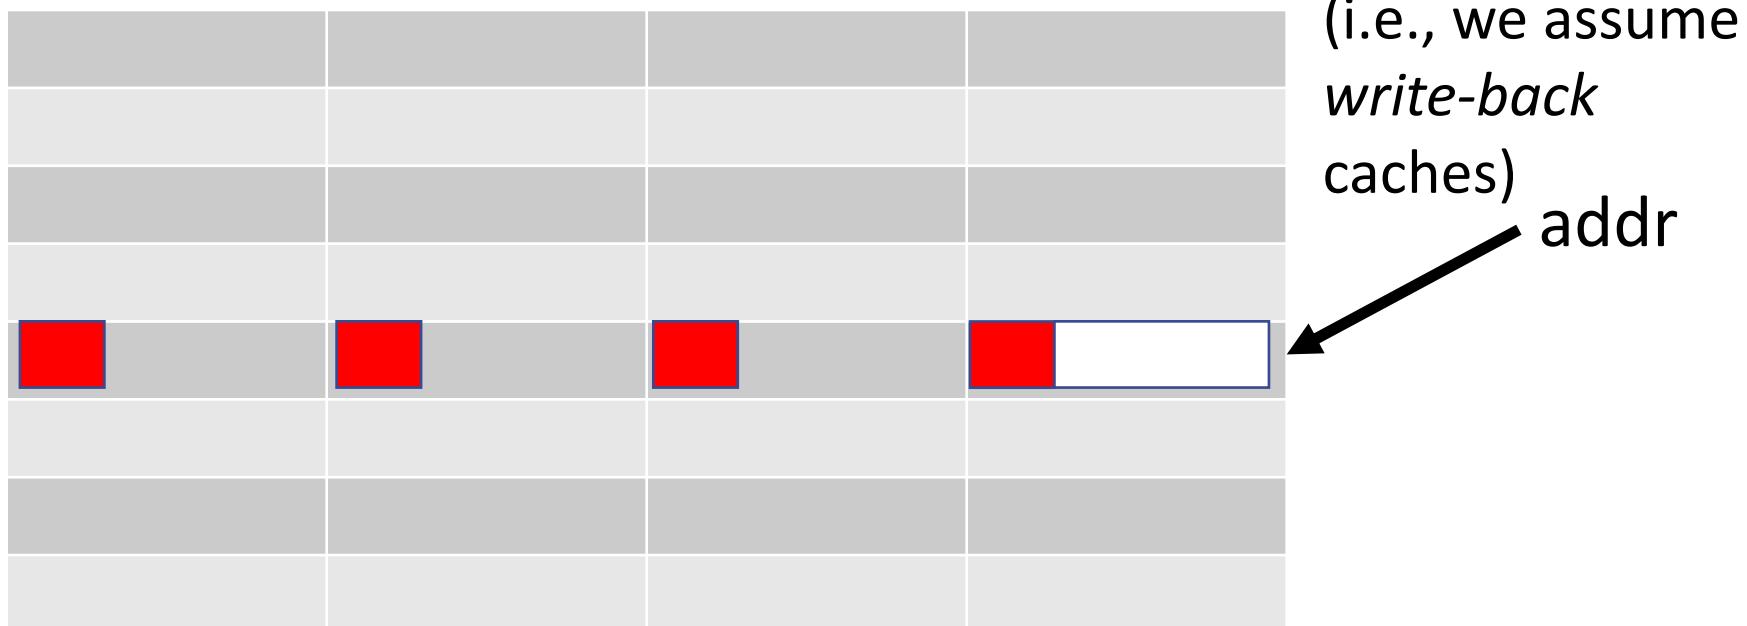
Hit: Line containing address in cache



Background: CPU cache

Miss: No line containing address in cache

Cache allocates way for data (even for writes)

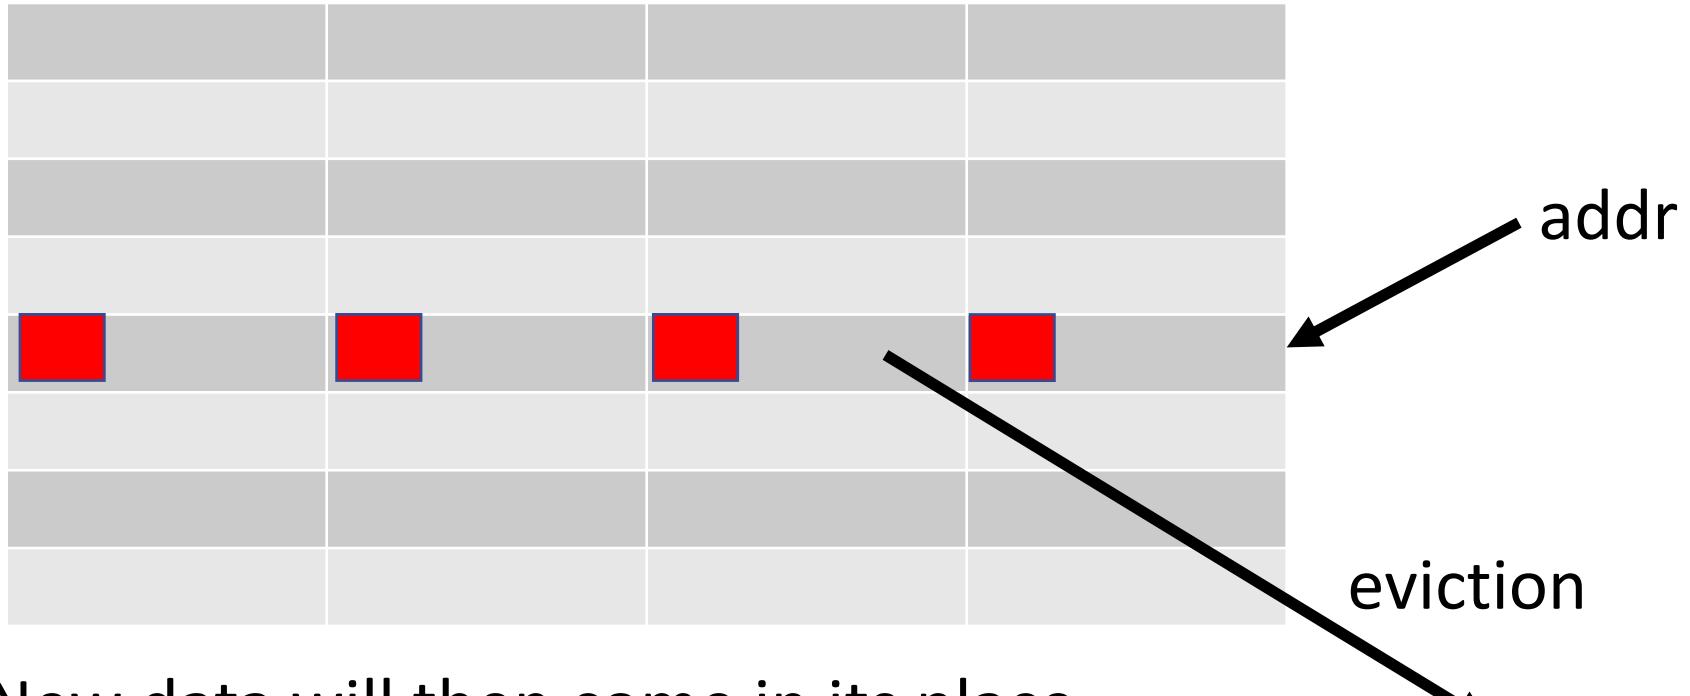


If some way in set is empty, data will go in there

Background: CPU cache

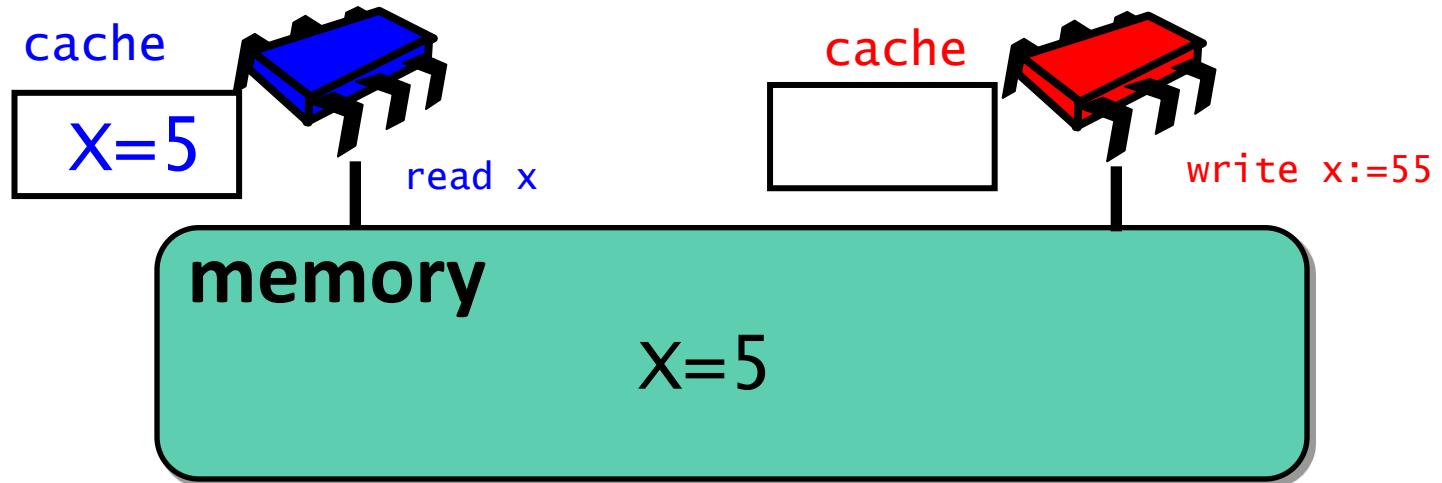
Typically, all ways have data

Replacement policy picks a line and *evicts* it

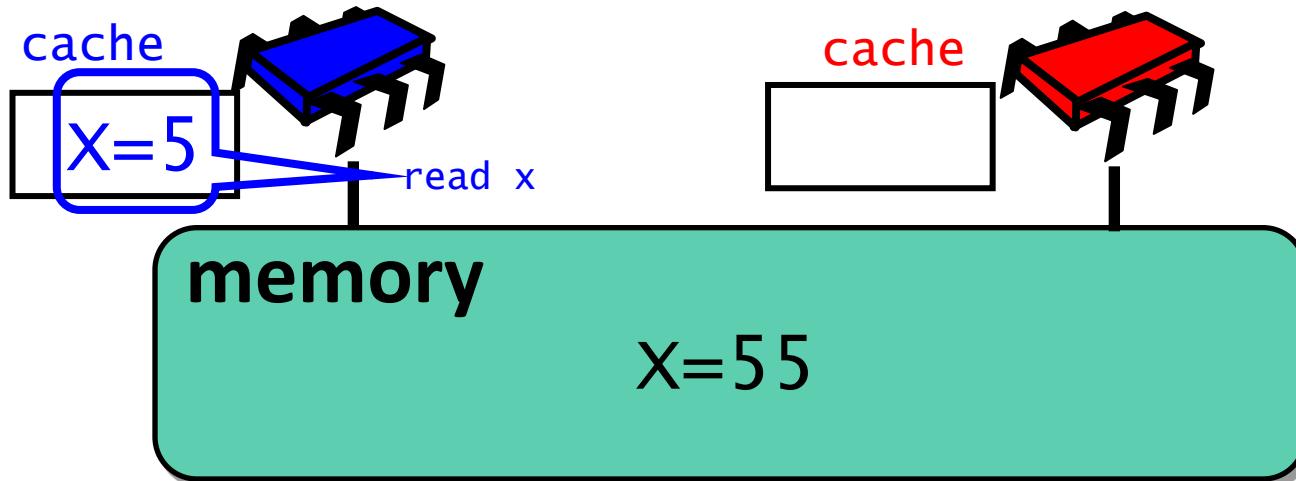


New data will then come in its place

Cache incoherence



Cache incoherence



Cache coherence

Cache coherence ("קשר הגיוני ורציף, "לכידות"):

- When multiple caches exist, want to prevent access to stale data

Applies to any type of caching system!

Cache (concept):

- Component storing data so that future requests for data can be served faster

Coherence abstraction

- In a cache coherent system, *caches are invisible*
- Any execution observed could have also been observed with just memory
- Example take-away: once a write leaves CPU (e.g., store buffer) it is *logically in memory*
 - Don't really care whether it's physically in DRAM or in some cache
 - Based on observing results of reads, no thread can distinguish the cases

Defining cache coherence

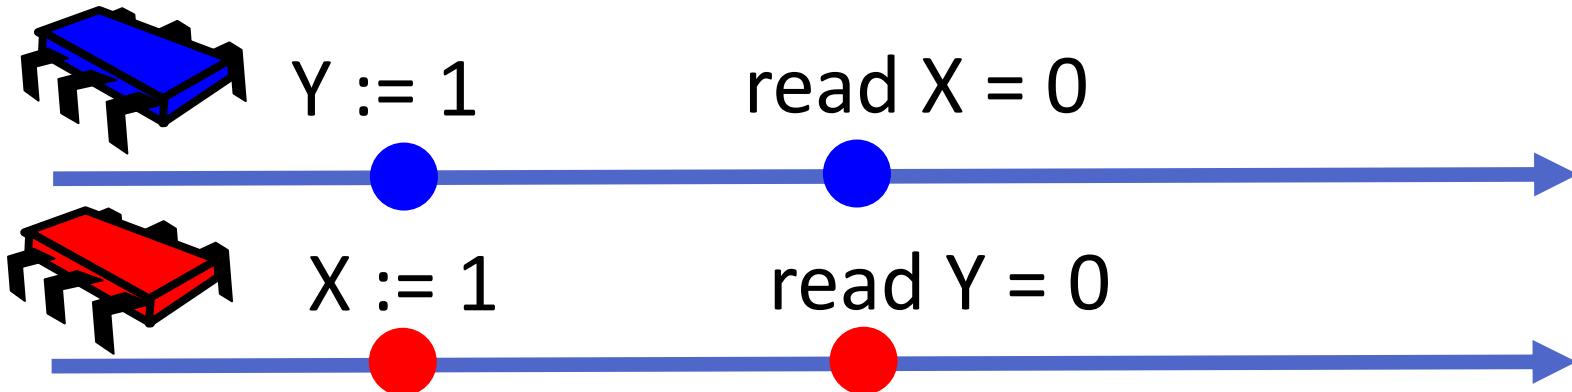
Important: cache coherence is defined per-location

A *consistency model* defines interactions between distinct locations

Defining cache coherence

One option:

- A coherent system must appear to execute all threads' loads and stores to a single memory location in a total order that respects the program order of each thread
- Like SC, but for a *single* location at a time
- Does *not* imply SC

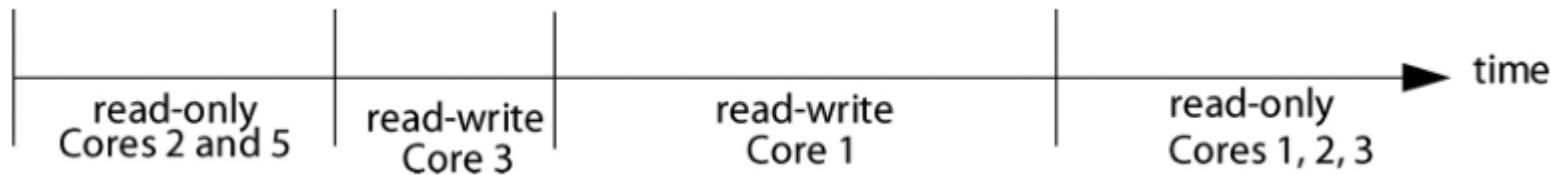


Defining cache coherence

Another option:

Single-writer-multi-reader. For any location, at any time, there is either a single core that may write it (and read it), or some number of cores that may read it

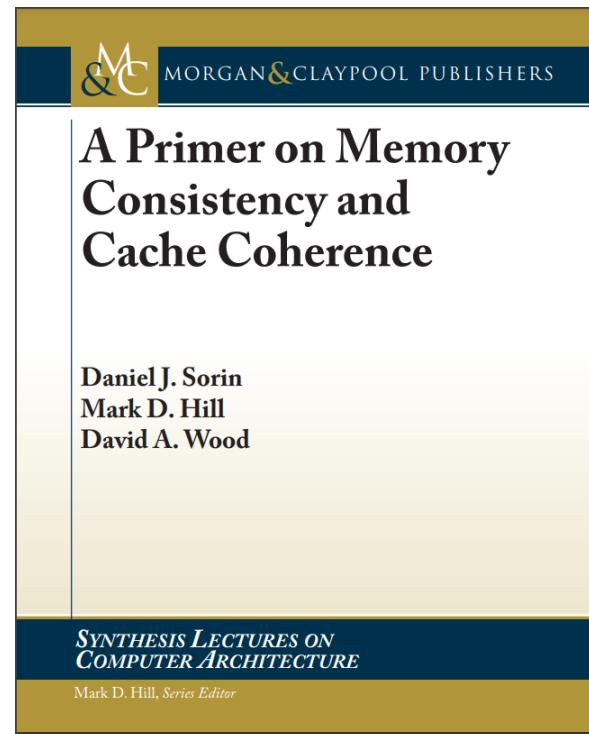
Equiv: memory location lifetime consists of *read/write epochs*



Value read at read epoch e is the value written in the prior write epoch

Lecture goals

- Understand multi-core cache coherence
- See details of a distributed algorithm
 - Multi-threaded algorithms have a similar taste
- Material based on book
(see link on moodle)

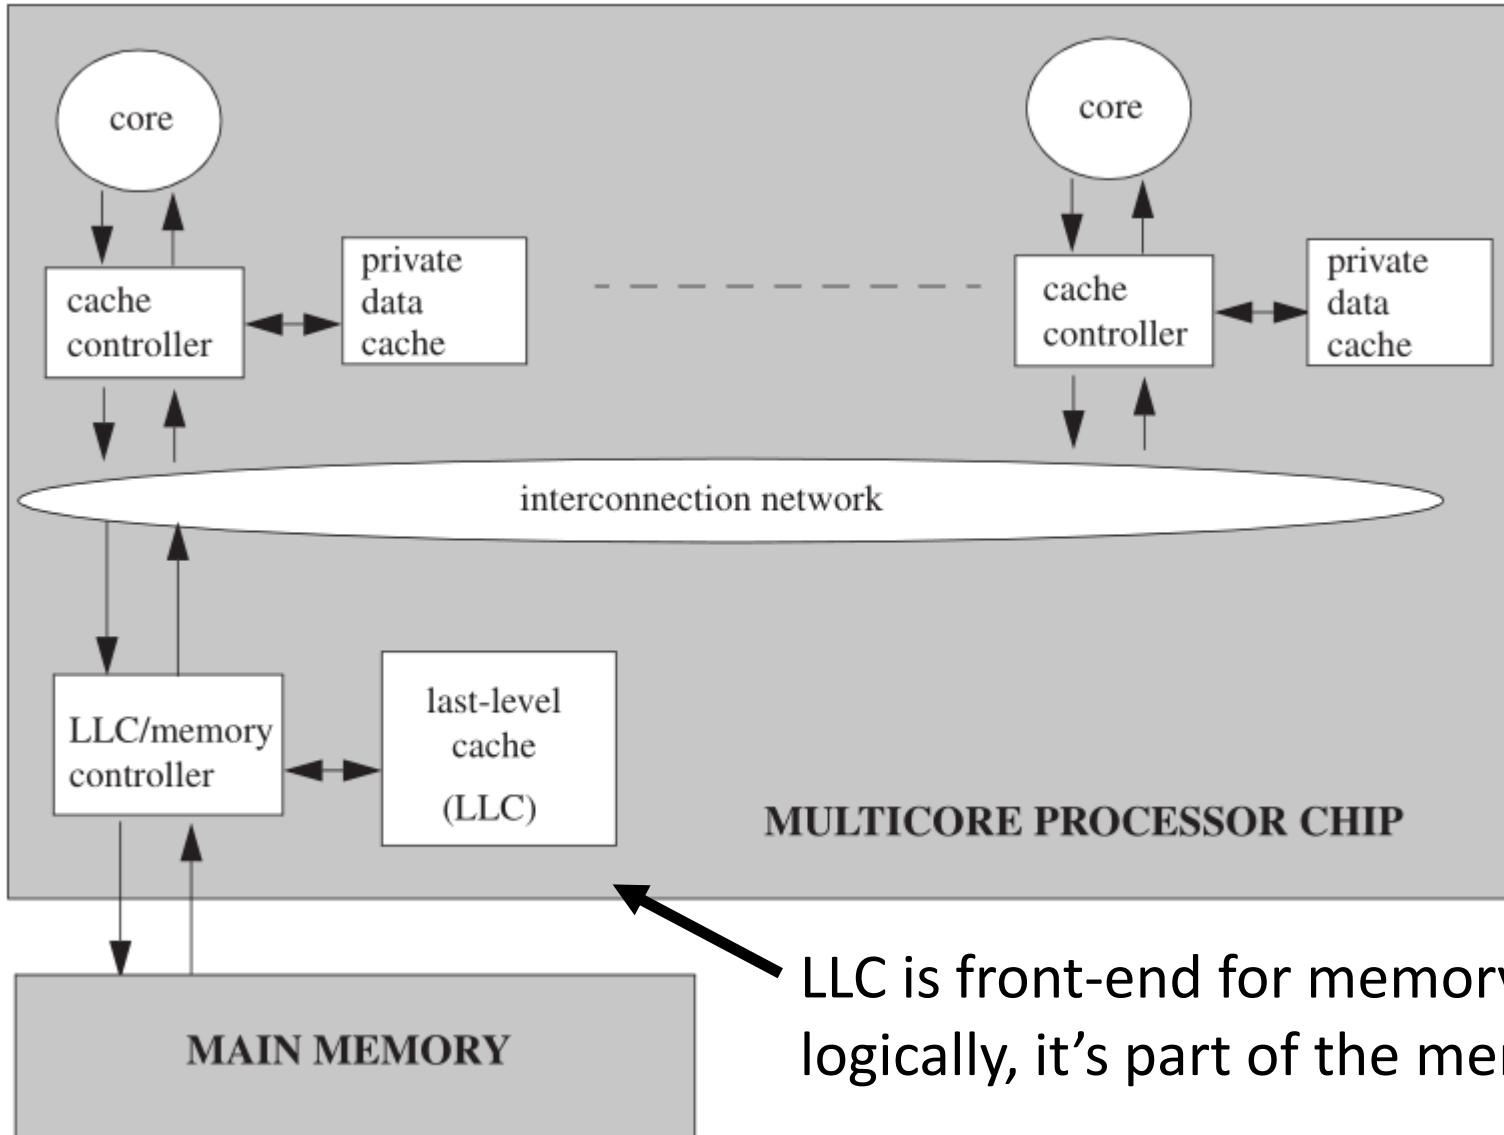


Coherence protocols

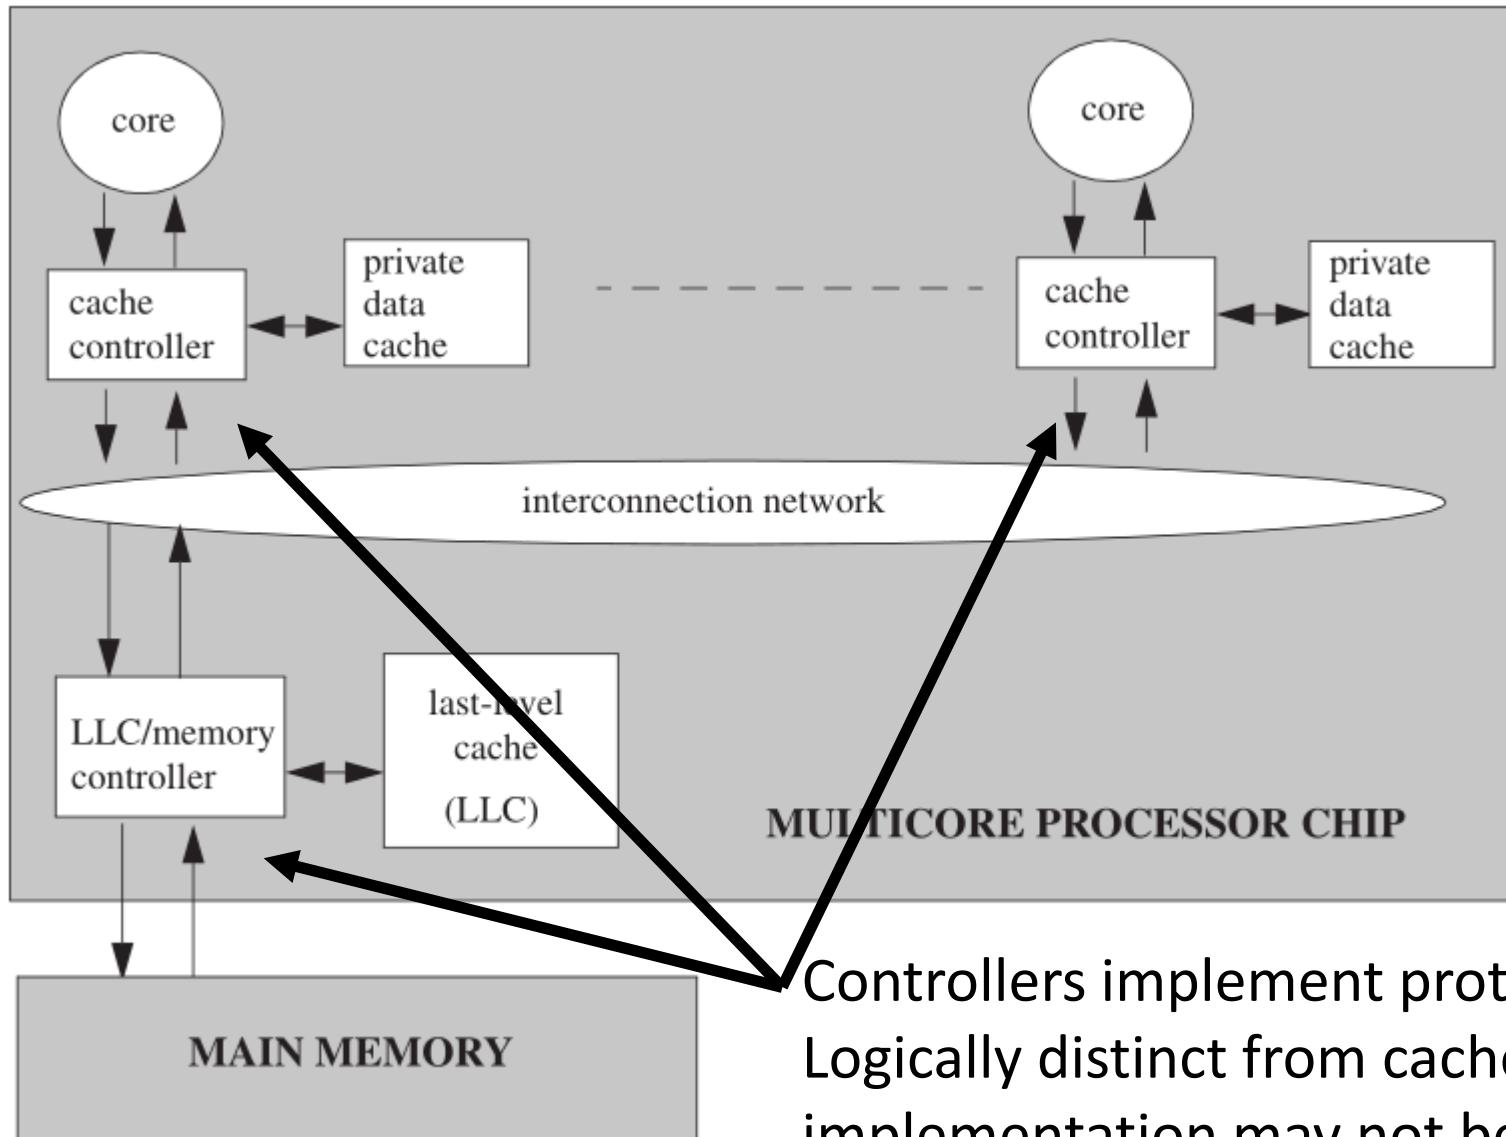
Coherence protocol

A set of rules implemented by the distributed set of caches in the system; if followed, provides cache coherence

System model

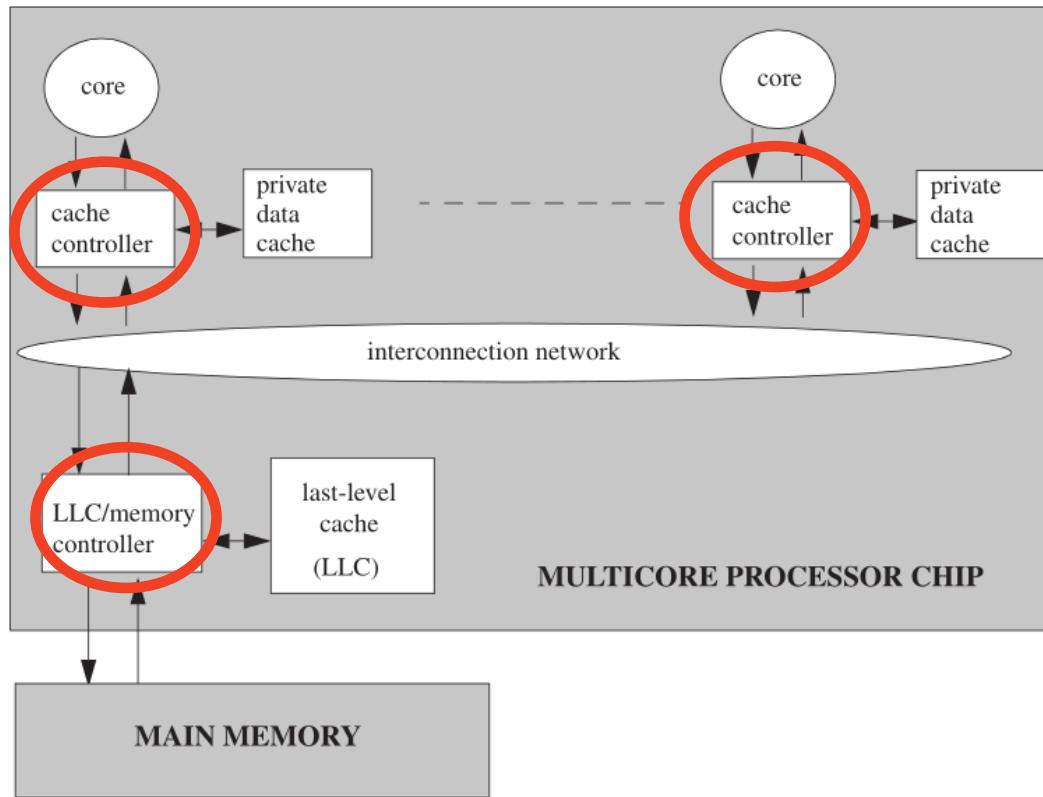


System model



Coherence protocol players

- Cache controllers
- Memory controllers



Cache controller

- Accepts loads/stores from core; returns loaded value to core
 - Cache miss issues a *coherence request* for the line
 - The request and all messages involved in satisfying it are called a *coherence transaction*
- Receives coherence requests and responses on the network

Memory controller

- Only has network connection
- Receives coherence requests on the network
 - Doesn't issue coherence requests
- We assume inclusive LLC
 - If line in private cache, LLC will also have state for it

Coherence protocol players

Cache controllers, memory controllers

- Implemented as state machine:
given state S and event E , take *actions* and
change state to S'
- Notice: state is *per-line*
(i.e., multiple state machines running concurrently)

Logical attributes of a line

Possible attributes of a line with respect to cache/LLC:

- *Valid* (cache has up to date value)
- *Exclusive* (cache has only privately cached line, but data may also be cached at the LLC)
- *Owned* (cache responsible for responding to coherence request; cannot be evicted without passing ownership)
- *Dirty* (up to date but different from value in memory/LLC; cache is responsible for updating them eventually)

Notice: These attributes are **not** mutually exclusive

MSI protocol

Possible cache controller states:

- M(odified)
- S(hared)
- I(nvalid)

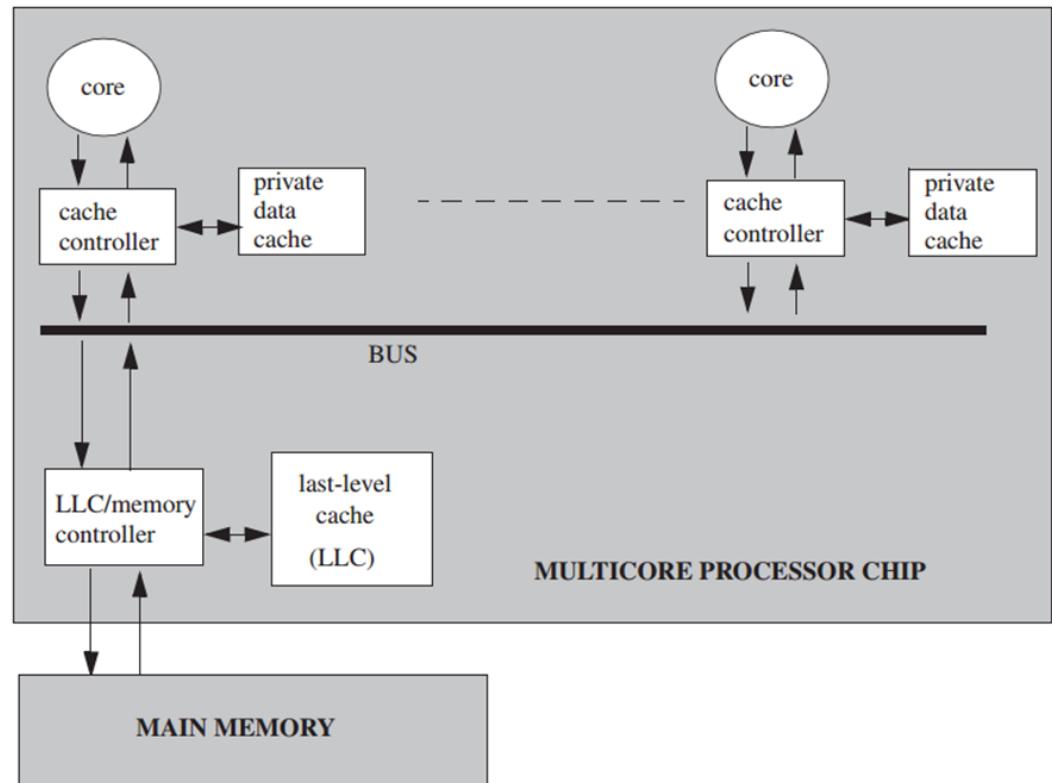
MSI protocol

Cache controller state:

- M(odified): valid, exclusive, owned, maybe dirty
- S(hared): valid, not exclusive, not dirty, not owned
 - Cache has read-only copy
 - LLC responsible for answering coherence requests
- I(nvalid): not valid
 - Cache doesn't have line or has a stale version, which it may not read or write

MSI snooping protocol

All controllers observe (*snoop*) coherence requests in the same order and can thus maintain coherence



MSI snooping protocol

- Coherence requests broadcast on *totally ordered* shared bus; observed in same order by all controllers (including sender)
- Controller sends and listens until seeing its request on the bus: may observe other requests in the meantime
- Bus acts as a serialization point
 - **Coherence transactions logically occurs when the request is serialized**
 - Responses don't need to go over shared bus

Shared bus analogy

Classroom with teacher

Raise hand: *Local event (send)*

Permission to speak: *Ordered on bus*

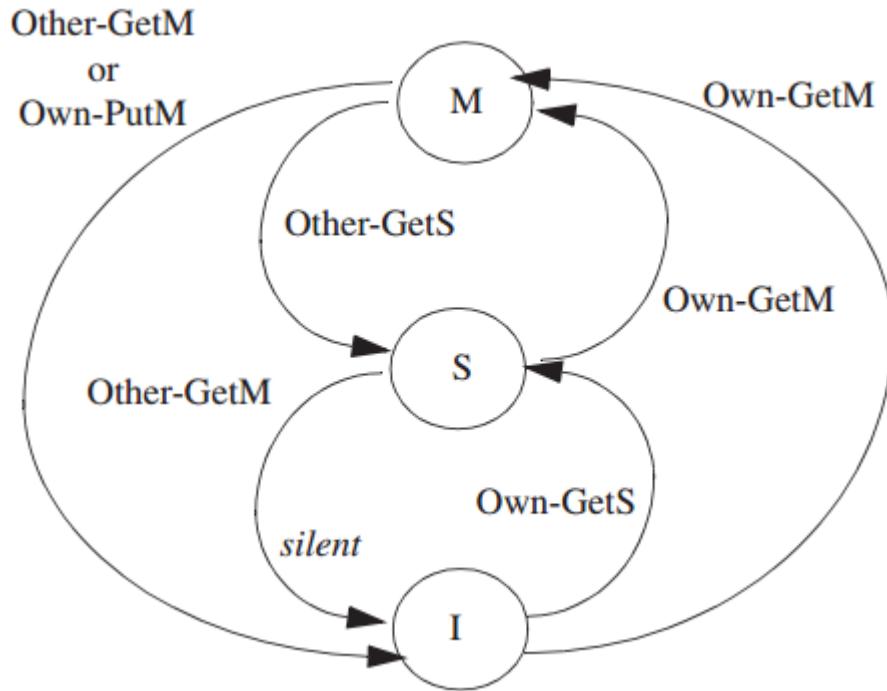
- Total order over all statements
- You hear other statements while waiting for your turn
- Everyone sees same events, so all agree about the state of the system and how it evolves

Simple protocol

Simplifying assumptions on system:

- *Atomic requests*: request ordered immediately
 - “Permission to speak immediately upon raising hand”
- *Atomic transactions*: a subsequent request for the same line cannot appear on the bus until the current transaction completes

MSI snooping protocol

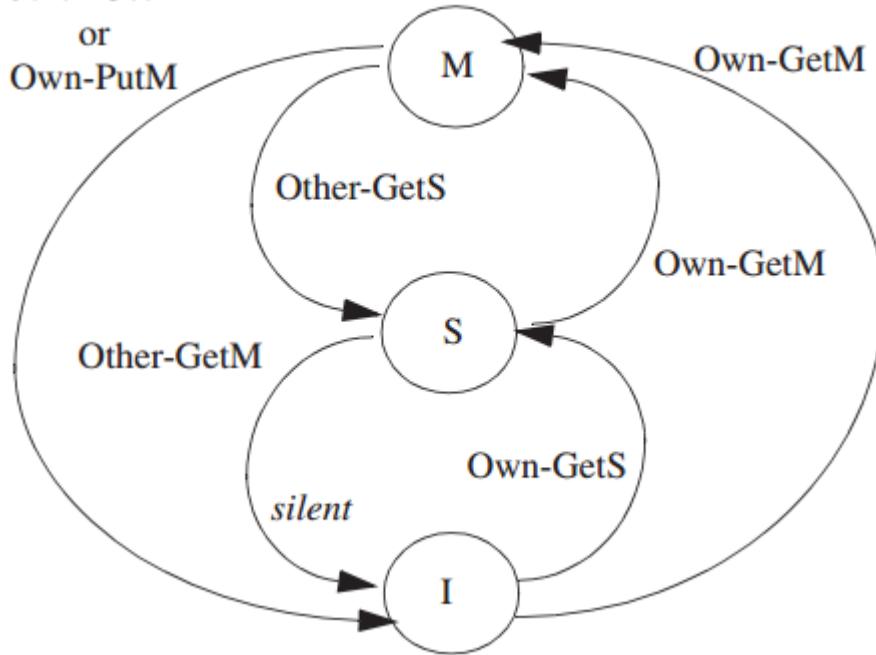


cache controller

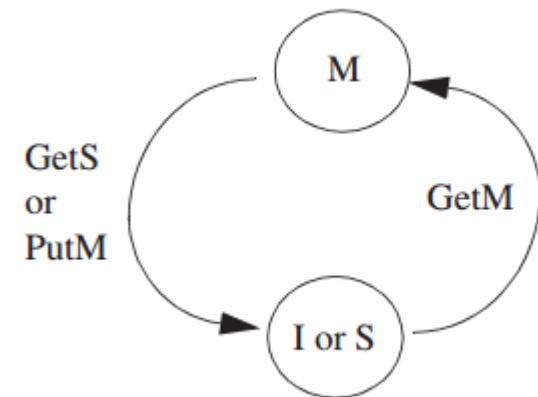
MSI snooping protocol

Other-GetM

or
Own-PutM



cache controller



memory controller

(state names reflect status of
line in caches)

Processor Core Events				Bus Events						
States	Own Transaction							Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	can happen any time									

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS ?									

Which state to go to after the GetS?

```

graph TD
    M((M)) -- "Other-GetM or Own-PutM" --> S((S))
    M -- "Own-GetM" --> I((I))
    S -- "Own-GetM" --> M
    S -- "Own-GetS" --> I
    S -- "silent" --> I
    I -- "Other-GetM" --> S
  
```

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS ?									

Which state to go to after the GetS?

Cannot go to S because don't have the data

States	Processor Core Events			Bus Events					
				Own Transaction				Transactions For Other Cores	
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM
I	issue GetS /IS ^D								
IS ^D									

Which state to go to after the GetS?

Cannot go to S because don't have the data

Transient state

States	Processor Core Events			Bus Events							
				Own Transaction				Transactions For Other Cores			
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D											
S											

Can't happen because of atomic transactions

Can't happen because of
atomic transactions

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M			
S										
M										

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)
S										
M										

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)
S	load hit									
M										

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)
S	load hit	issue GetM								
M										

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)
S	load hit	issue GetM /SM ^D								
SM ^D										
M										

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)
S	load hit	issue GetM /SM ^D	-/I							
SM ^D										
M										

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)
S	load hit	issue GetM /SM ^D	-/I							
SM ^D										
M										

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)
S	load hit	issue GetM /SM ^D	-/I						-/I	
SM ^D										
M										

States	Processor Core Events			Bus Events						
				Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM
I	issue GetS /IS ^D	issue GetM /IM ^D								
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)
S	load hit	issue GetM /SM ^D	-/I						-/I	
SM ^D	load hit									
M										

States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict								
M											

States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict				copy data into cache, store hit /M				
M											

States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
M											

States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
M	load hit										

States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
M	load hit	store hit									

States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
M	load hit	store hit	issue PutM, send Data to memory /I								

States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
M	load hit	store hit	issue PutM, send Data to memory /I								

States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
M	load hit	store hit	issue PutM, send Data to memory /I					send Data to req and memory /S	send Data to req /I		

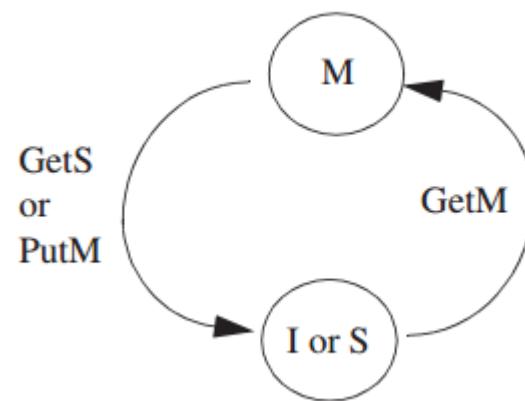
States	Processor Core Events				Bus Events						
					Own Transaction				Transactions For Other Cores		
	Load	Store	Replacement	Own-GetS	Own-GetM	Own-PutM	Data	Other-GetS	Other-GetM	Other-PutM	
I	issue GetS /IS ^D	issue GetM /IM ^D									
IS ^D	stall Load	stall Store	stall Evict				copy data into cache, load hit /S	(A)	(A)	(A)	
IM ^D	stall Load	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
S	load hit	issue GetM /SM ^D	-/I						-/I		
SM ^D	load hit	stall Store	stall Evict				copy data into cache, store hit /M	(A)	(A)	(A)	
M	load hit	store hit	issue PutM, send Data to memory /I					send Data to req and memory /S	send Data to req /I		

What's wrong with this table?

Memory controller

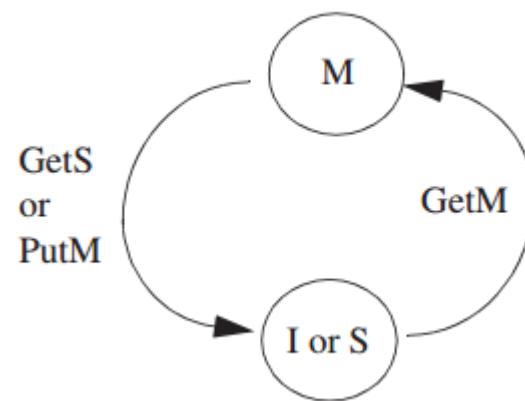
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
I or S				
M				



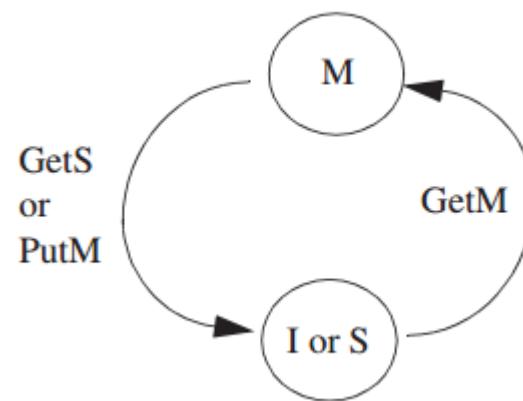
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS			
M				



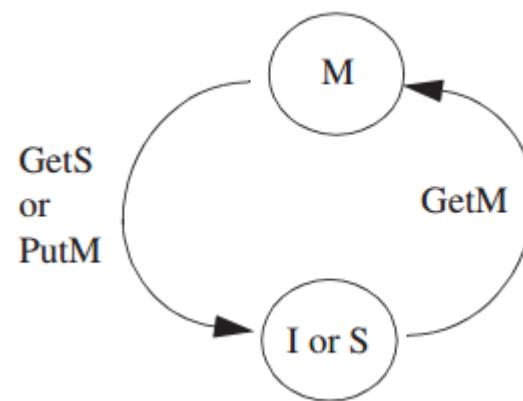
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
M				



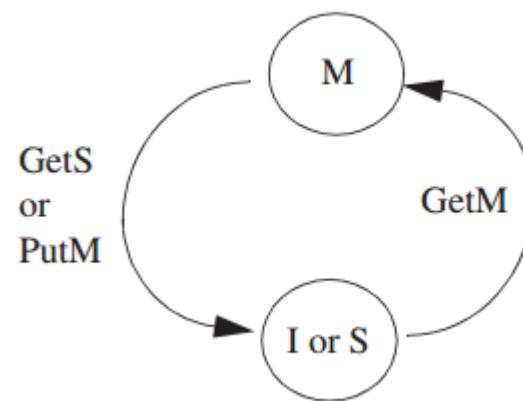
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
M				



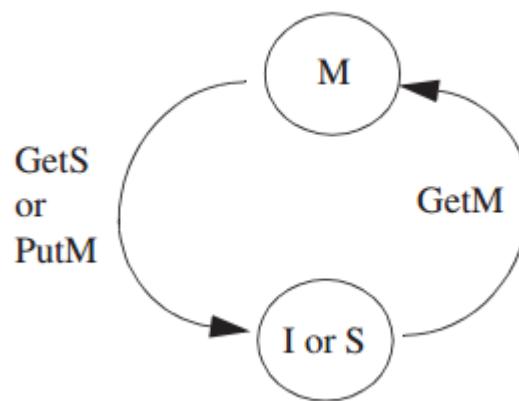
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
M	-/ ?			



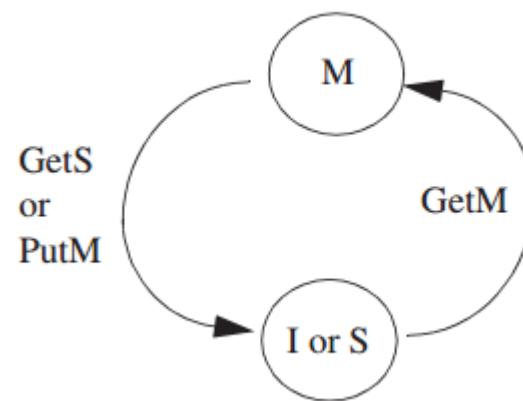
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
IorS ^D				
M	-/IorS ^D			



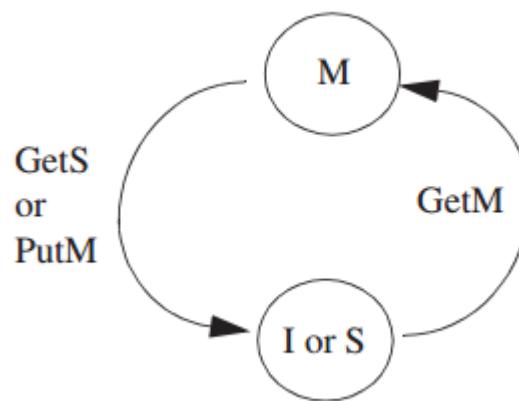
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
IorS ^D	(A)	(A)		
M	-/IorS ^D			



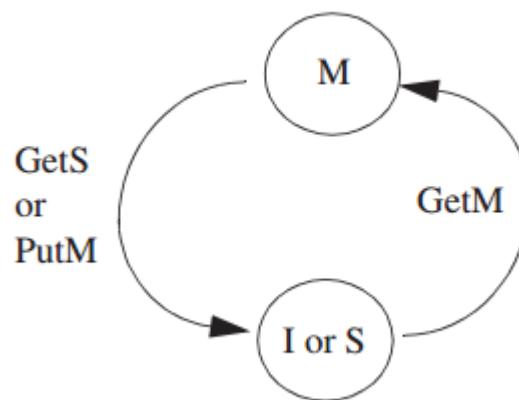
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
IorS ^D	(A)	(A)		update data block in memory/IorS
M	-/IorS ^D			



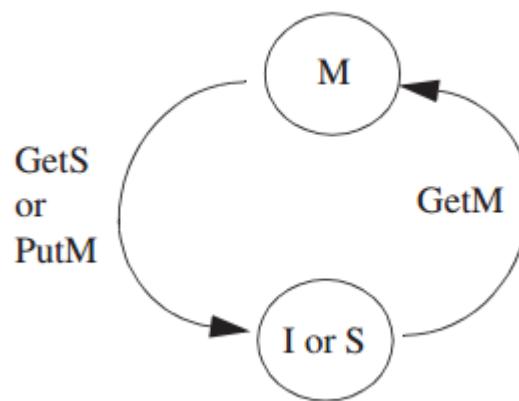
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
IorS ^D	(A)	(A)		update data block in memory/IorS
M	-/IorS ^D			



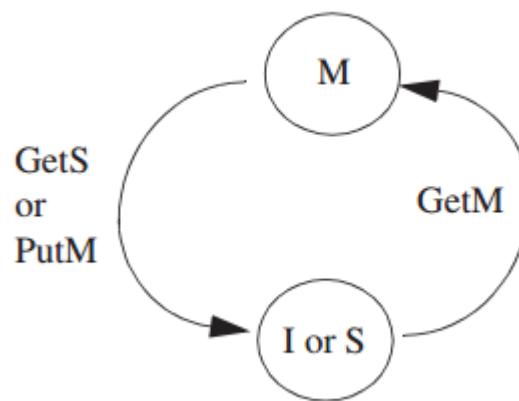
Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
IorS ^D	(A)	(A)		update data block in memory/IorS
M	-/IorS ^D		-/IorS ^D	



Memory controller

State	Bus Events			
	GetS	GetM	PutM	Data from Owner
IorS	send data block in Data message to requestor/IorS	send data block in Data message to requestor/M		
IorS ^D	(A)	(A)		update data block in memory/IorS
M	-/IorS ^D		-/IorS ^D	



Non-atomic requests

Time can pass (and transactions can get ordered) between making a request and seeing it on the bus

Non-atomic requests

Time can pass (and transactions can get ordered) between making a request and seeing it on the bus

I->S and I->M need another transient state:

- IS^{AD}: Logically still in I mode, waiting to be ordered
- IS^D: Logically in S mode, just don't have the data

	load	store	replacement	OwnGetS	OwnGetM	OwnPutM	OtherGets	OtherGetM	OtherPutM	Own Data response
I	issue GetS/IS ^{AD}	issue GetM/IM ^{AD}				-	-	-	-	
IS ^{AD}	stall	stall	stall	-/IS ^D		-	-	-	-	
IS ^D	stall	stall	stall			(A)	(A)			-/S
IM ^{AD}	stall	stall	stall		-/IM ^D	-	-	-	-	
IM ^D	stall	stall	stall			(A)	(A)			-/M
S										
M										

Non-atomic requests

Time can pass (and transactions can get ordered) between making a request and seeing it on the bus

I->S and I->M need another transient state:

- IS^{AD} : Logically still in I mode, waiting to be ordered
- IS^D : Logically in S mode, just don't have the data

S->M more complex:

- $S \xrightarrow{\text{store}} SM^{AD} \xrightarrow{\text{Can satisfy loads}} IM^{AD} \xrightarrow{\text{Other-GetM}} IM^{AD} \xrightarrow{\text{Can't satisfy loads}}$

	load	store	replacement	OwnGetS	OwnGetM	OwnPutM	OtherGets	OtherGetM	OtherPutM	Own Data response
I	issue GetS/IS ^{AD}	issue GetM/IM ^{AD}				-	-	-	-	
IS ^{AD}	stall	stall	stall	-/IS ^D		-	-	-	-	
IS ^D	stall	stall	stall			(A)	(A)			-/S
IM ^{AD}	stall	stall	stall		-/IM ^D	-	-	-	-	
IM ^D	stall	stall	stall			(A)	(A)			-/M
S	hit	issue GetM/SM ^{AD}	-/I			-	-/I	-		
SM ^{AD}	hit	stall	stall		-/SM ^D	-	-/IM ^{AD}	-		
SM ^D	hit	stall	stall			(A)	(A)			-/M
M										

Non-atomic requests

Time can pass (and transactions can get ordered) between making a request and seeing it on the bus

I->S and I->M need another transient state:

S->M more complex:

- $S \xrightarrow{\text{store}} SM^{AD} \xrightarrow{\text{loads OK}} IM^{AD} \xrightarrow{\text{Other-GetM}} \text{loads not OK}$

M->I even more complex:

- $M \xrightarrow{\text{data still owned}} MI^A \xrightarrow{\text{Other-GetM/ Send data to requestor}} II^A \xrightarrow{\text{Own-PutM/ Send NoData msg to memory}} I$

	load	store	replacement	OwnGetS	OwnGetM	OwnPutM	OtherGets	OtherGetM	OtherPutM	Own Data response
I	issue GetS/IS ^{AD}	issue GetM/IM ^{AD}				-	-	-	-	
IS ^{AD}	stall	stall	stall	-/IS ^D		-	-	-	-	
IS ^D	stall	stall	stall			(A)	(A)			-/S
IM ^{AD}	stall	stall	stall		-/IM ^D	-	-	-	-	
IM ^D	stall	stall	stall			(A)	(A)			-/M
S	hit	issue GetM/SM ^{AD}	-/I			-	-/I	-		
SM ^{AD}	hit	stall	stall		-/SM ^D	-	-/IM ^{AD}	-		
SM ^D	hit	stall	stall			(A)	(A)			-/M
M	hit	hit	issue PutM/MI ^A			send data to requestor and to memory/S	send data to requestor/I	-		
MI ^A										
II ^A										

	load	store	replacement	OwnGetS	OwnGetM	OwnPutM	OtherGets	OtherGetM	OtherPutM	Own Data response
I	issue GetS/IS ^{AD}	issue GetM/IM ^{AD}				-	-	-	-	
IS ^{AD}	stall	stall	stall	-/IS ^D		-	-	-	-	
IS ^D	stall	stall	stall			(A)	(A)			-/S
IM ^{AD}	stall	stall	stall		-/IM ^D	-	-	-	-	
IM ^D	stall	stall	stall			(A)	(A)			-/M
S	hit	issue GetM/SM ^{AD}	-/I			-	-/I	-		
SM ^{AD}	hit	stall	stall		-/SM ^D	-	-/IM ^{AD}	-		
SM ^D	hit	stall	stall			(A)	(A)			-/M
M	hit	hit	issue PutM/MI ^A			send data to requestor and to memory/S	send data to requestor/I	-		
MI ^A	hit	hit	stall			send data to memory/I	send data to requestor and to memory/II ^A	send data to requestor/II ^A		
II ^A										

	load	store	replacement	OwnGetS	OwnGetM	OwnPutM	OtherGets	OtherGetM	OtherPutM	Own Data response
I	issue GetS/IS ^{AD}	issue GetM/IM ^{AD}				-	-	-	-	
IS ^{AD}	stall	stall	stall	-/IS ^D		-	-	-	-	
IS ^D	stall	stall	stall			(A)	(A)			-/S
IM ^{AD}	stall	stall	stall		-/IM ^D	-	-	-	-	
IM ^D	stall	stall	stall			(A)	(A)			-/M
S	hit	issue GetM/SM ^{AD}	-/I			-	-/I	-		
SM ^{AD}	hit	stall	stall		-/SM ^D	-	-/IM ^{AD}	-		
SM ^D	hit	stall	stall			(A)	(A)			-/M
M	hit	hit	issue PutM/MI ^A			send data to requestor and to memory/S	send data to requestor/I	-		
MI ^A	hit	hit	stall			send data to memory/I	send data to requestor and to memory/II ^A	send data to requestor/II ^A		
II ^A	stall	stall	stall			send NoData to memory/I	-	-		

load	store	replacement	OwnGetS	OwnGetM	OwnPutM	OtherGets	OtherGetM	OtherPutM	Own Data response
------	-------	-------------	---------	---------	---------	-----------	-----------	-----------	-------------------

Why send to memory controller?

- It observes the PutM, so it waits for data. (It has no way of knowing when core issued PutM; it only sees it on the bus.)
- On the other hand, can't send real data, because ours might be stale by now!

II ^A	stall	stall	stall			send NoData to memory/I	-	-	
-----------------	-------	-------	-------	--	--	-------------------------------	---	---	--

Memory controller

	GetS	GetM	PutM	Data From Owner	NoData
IorS	send data to requestor	send data to requestor/M	-/IorS ^D		
IorS ^D	(A)	(A)		write data to LLC/memory /IorS	-/IorS
M	-/IorS ^D		-/M ^D		
M ^D					

Memory controller

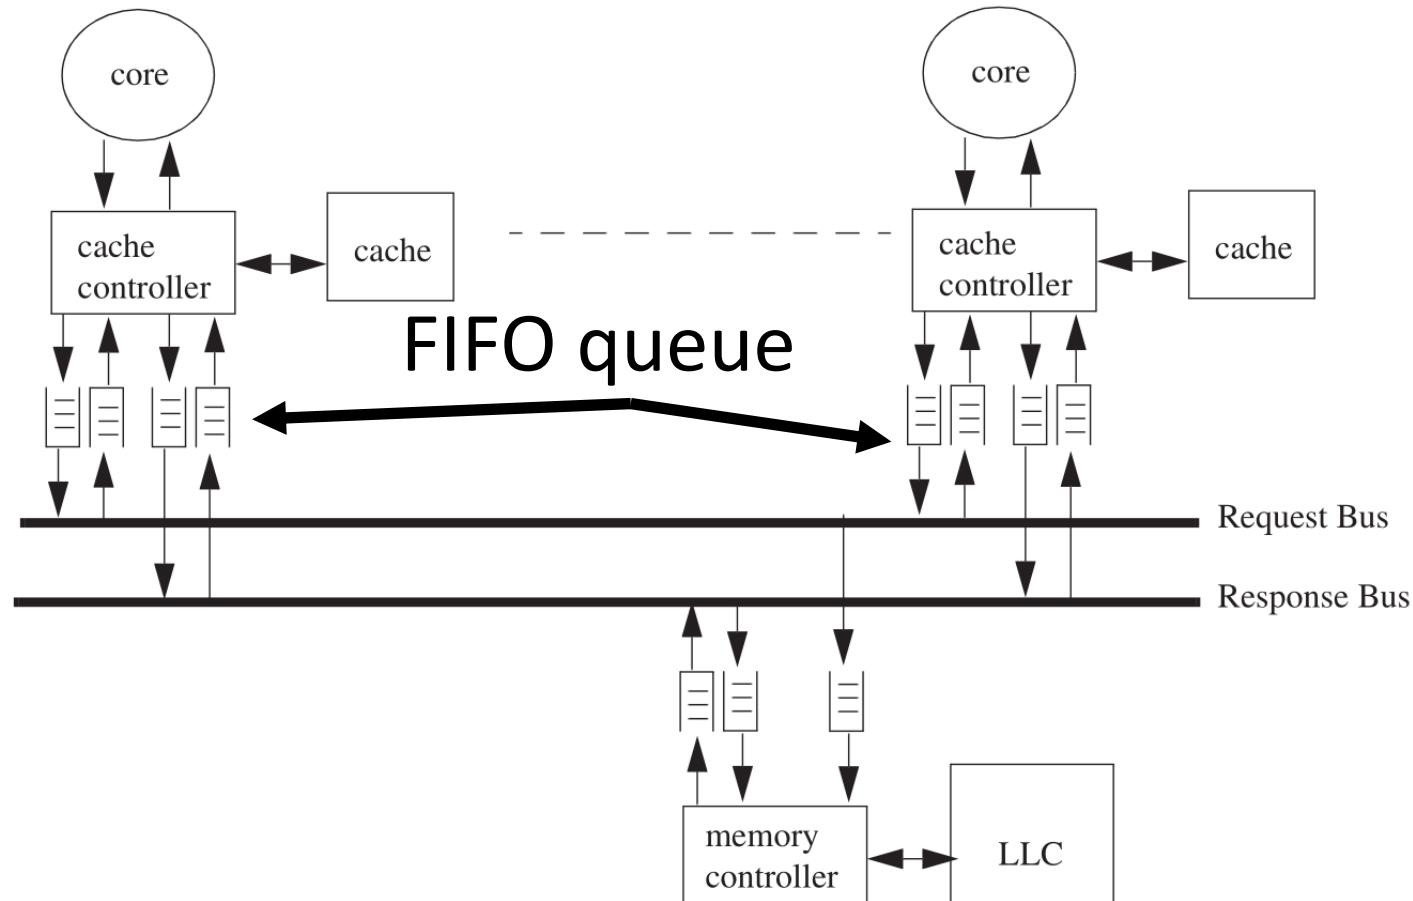
	GetS	GetM	PutM	Data From Owner	NoData
IorS	send data to requestor	send data to requestor/M	-/IorS ^D		
IorS ^D	(A)	(A)		write data to LLC/memory /IorS	-/IorS
M	-/IorS ^D		-/M ^D		
M ^D	(A)	(A)		write data to LLC/IorS	-/M

Upgrade transaction

- Current S->M issues a GetM, which waits for the data from the LLC
- But we have the data: an *upgrade* transaction would invalidate other sharers and then go to M state.
- Complicated by possibility of state changing before transaction gets ordered; see book

Non-atomic transactions

Can observe another transaction before getting response to mine



Non-atomic transactions

Example problem:

- Issued (& observed) GetM: line is logically mine
- Observe a GetS (but I don't have the data)

Non-atomic transactions

Example problem:

- Issued (& observed) GetM: line is logically mine
- Observe a GetS (but I don't have the data)

Simple solution:

- Stall response (buffer request and respond only when data arrives)

Non-atomic transactions

Example problem:

- Issued (& observed) GetM: line is logically mine
- Observe a GetS (but I don't have the data)

Simple solution:

- Stall response (buffer request and respond only when data arrives)
- Deadlock? No, because my request was already ordered, now just waiting for data

Non-stalling non-atomic txns

Instead of stalling requests, can add new transient states:

$IM^D S$: in I, going to M, waiting for data, and when
data arrives will go to S

$IS^D I$, $IM^D I$, $IM^D SI$

To avoid livelock, must perform the instruction that triggered the original request (e.g., store for $IM^D S$) when data arrives, before switching to the final state

- (Otherwise, in the $IM^D S$ case, we'd end up in S state with a store pending, which would trigger another GetM request... which can go on forever)

MESI protocol

- Optimizes for a common case of write-after-read; used in commercial processors
- MSI: GetS followed by GetM
- MESI: GetS can get line in E(xclusive) state, if no other cache has line in S state. Then, later store can silently upgrade from E to M
 - Requires: knowing that no other cache has line in S state
 - LLC can track this; hard to do precisely due to silent evictions, but “LLC state == I” is good enough

Atomic RMW instructions

No need to “lock the bus”

- Get line in M state
- Stall incoming coherence requests
- Complete instruction
- Resume processing incoming coherence transactions

Directories

Scalability problems with snooping:

- Bandwidth
- Having to process unrelated coherence requests
- Enforcing global total order

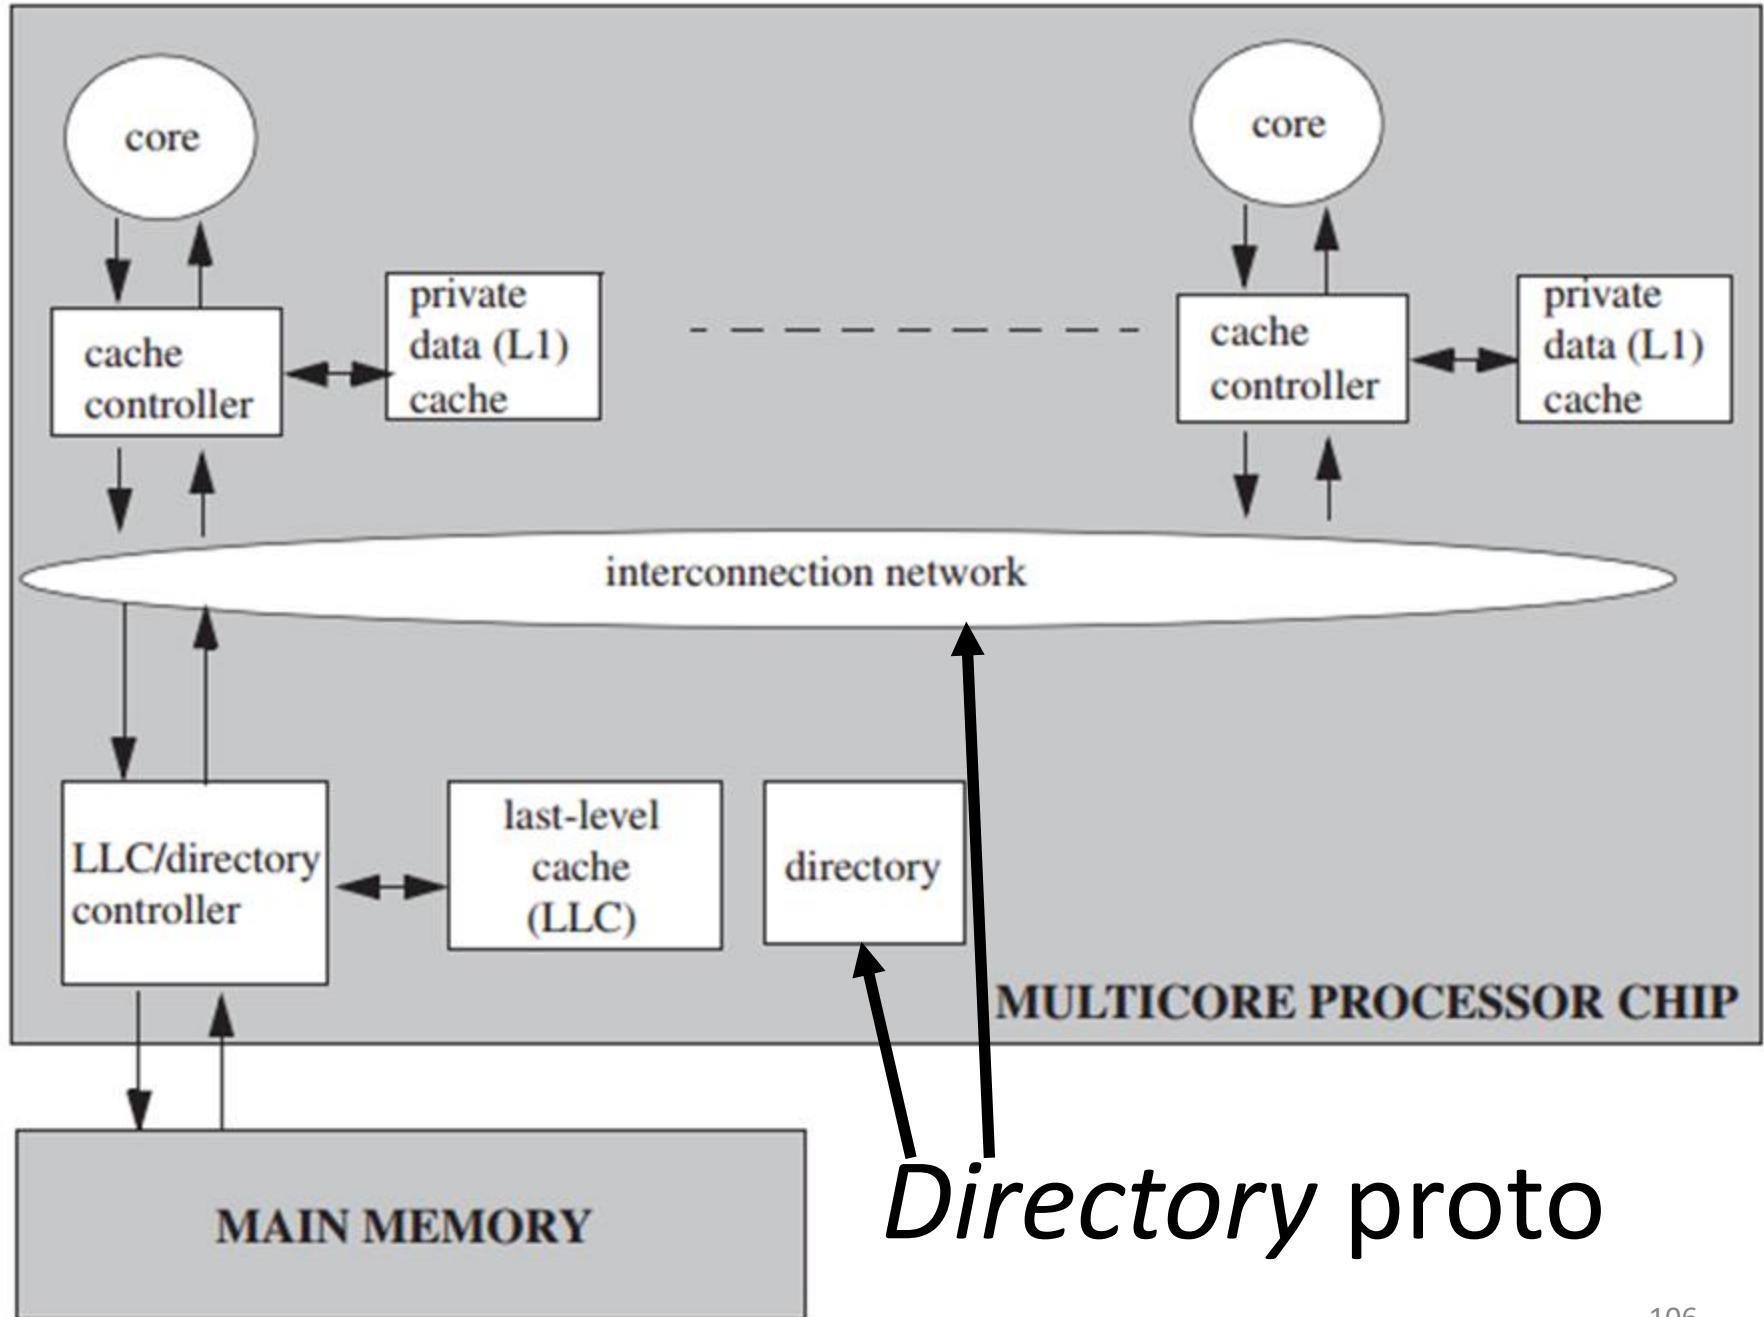
Directories

Scalability problems with snooping:

- Bandwidth
- Having to process unrelated coherence requests
- Enforcing global total order

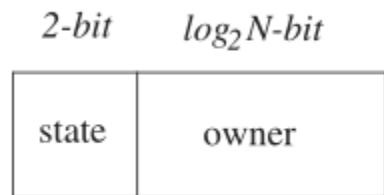
Directory:

- Maintains global view of coherence state instead of a distributed view
- Requests go to it rather than being broadcast
- Trade-off: latency, need for acks

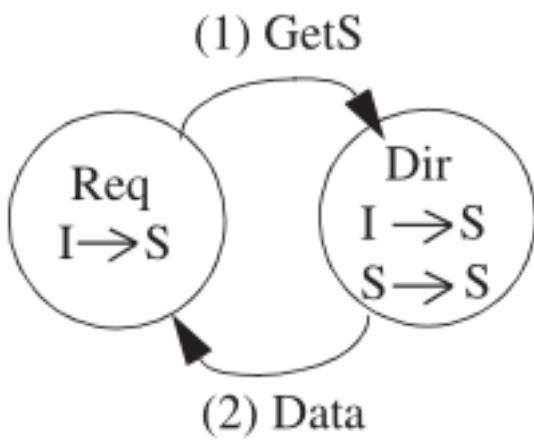


Directory state

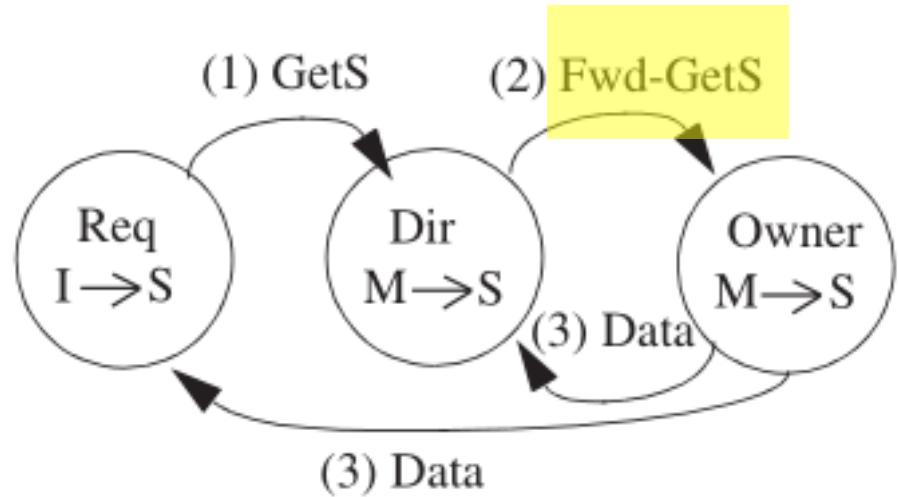
- Line is owned by directory (i.e., will respond with data) unless it's in M state at some core



I → S flow



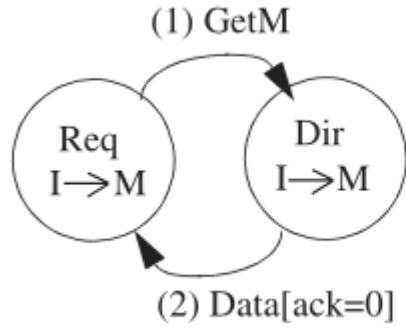
line owned by dir



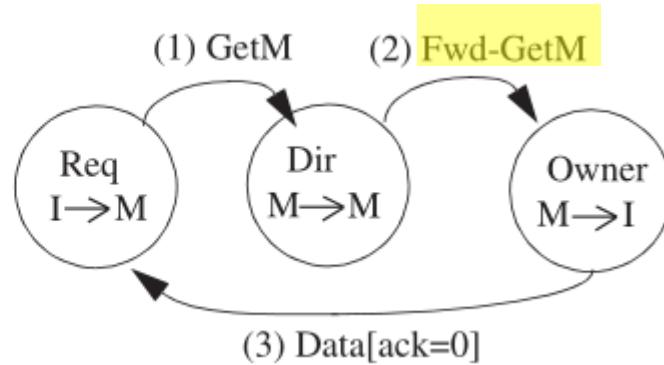
line owned by core *Owner*

new message types

I → M flow



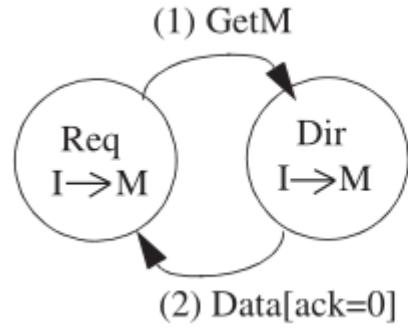
line only in mem



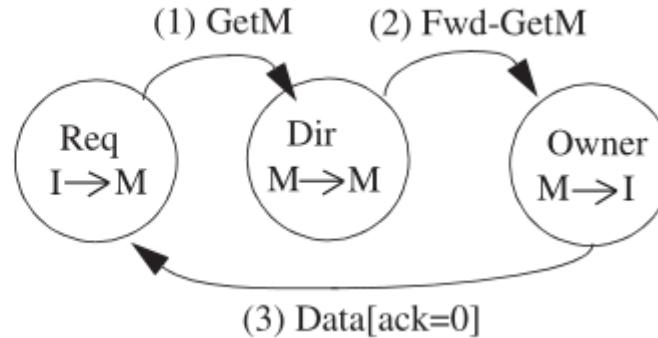
line M at *Owner*

new message types

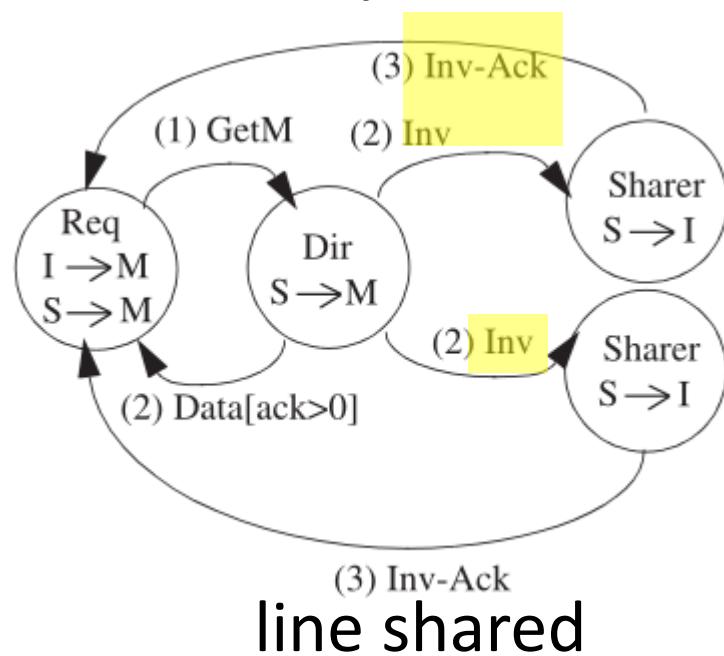
I → M flow



line only in mem



line M at *Owner*



if *Req* is only sharer,
Data[ack=0]

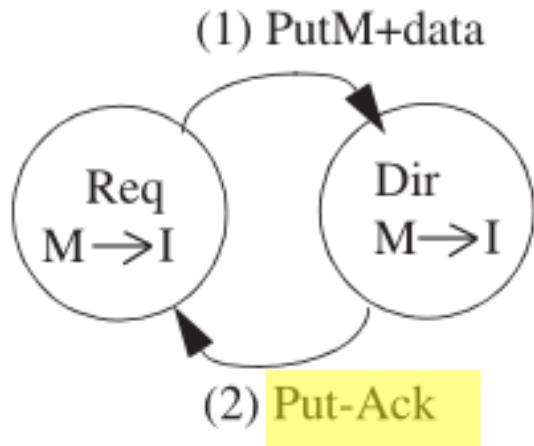
new message types

Directory state

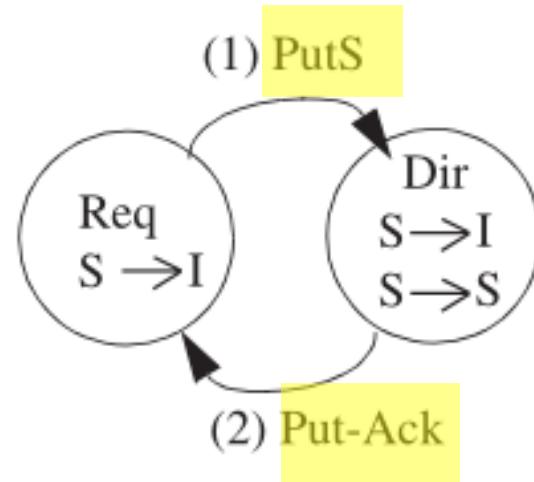
- Line is owned by directory (i.e., will respond with data) unless it's in M state at some core
- **Directory needs to know who has line in S state**
 - In snooping, sharers observed all requests; now they don't, and directory needs to forward requests to them

2-bit	$\log_2 N\text{-bit}$	$N\text{-bit}$
state	owner	sharer list (one-hot bit vector)

Evictions



line M at *Req*



line shared

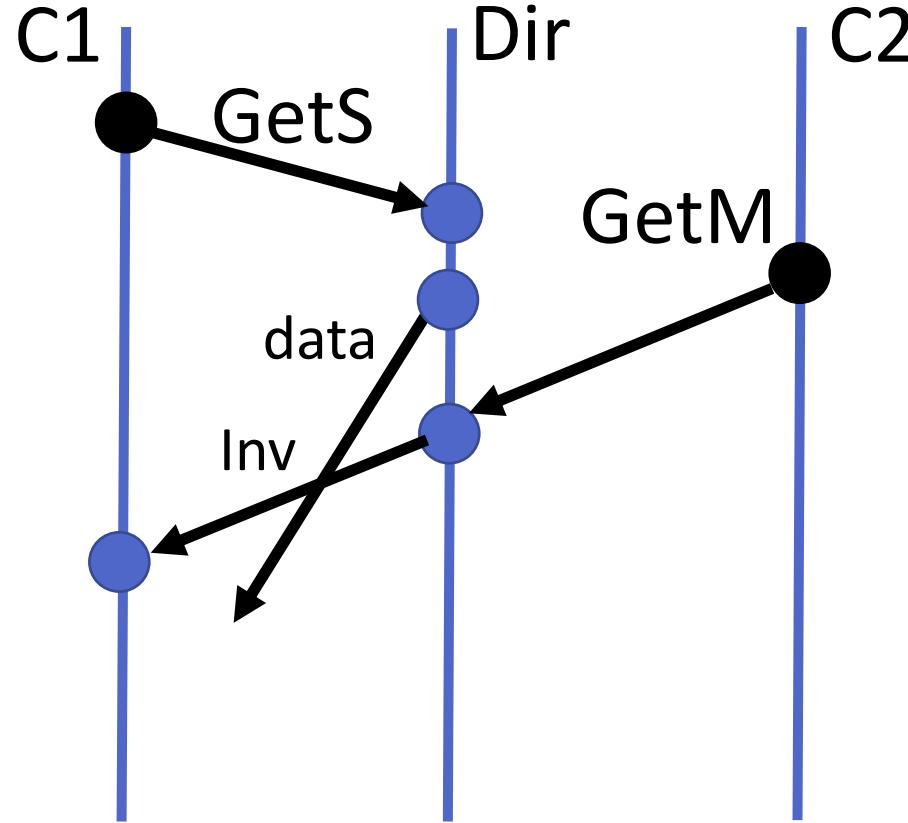
new message types

Evictions

- PutM now includes data
 - In snooping, sent two messages in parallel
- Evictions from S no longer silent
 - Directory needs to keep track of sharers
 - Avoids a race:
 - Cache silently evicts, then does a GetS
 - Now gets an Inv; was this Inv sent for the prior S state or the next one?

Races

- Multiple transactions in progress for a line
- Example:

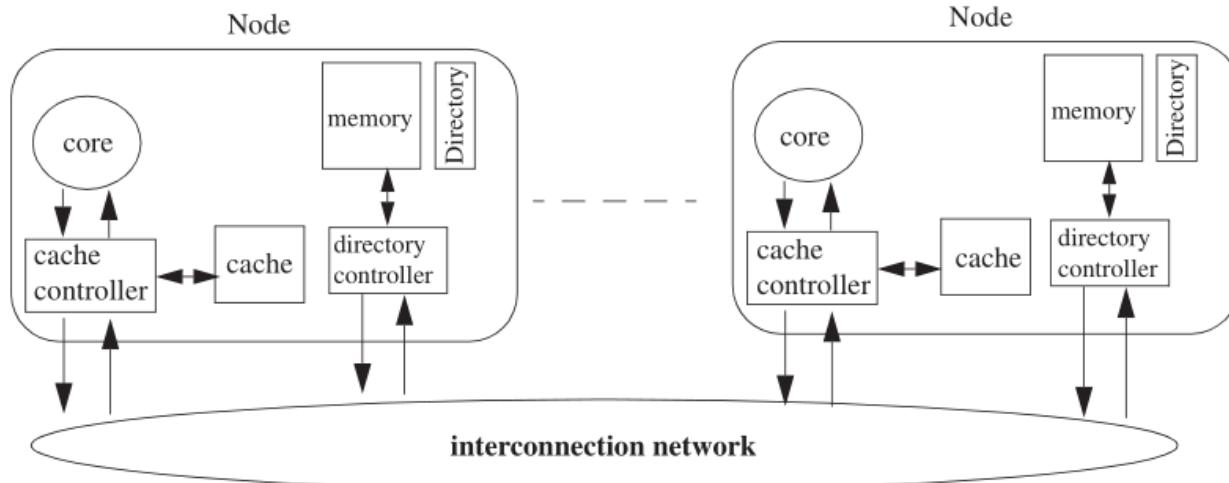


Races

- Multiple transactions in progress for a line
- Solutions:
 - Stall racy requests (costs performance)
 - More transient states

Distributed directories

- Everything works (and scales better) if we have several directories, and hash the line to obtain its directory
- Example: NUMA system



- Same idea applies within chip (banked LLC and directory)

Summary

- Cache coherence
- MSI & MESI
 - Commercial protocols are based on these
- Snooping and directory-based protocols
- **Important:** state machine way of looking at an algorithm