

# Advanced Topics in Multicore Architecture & Software Systems

Serializing efficiently

# Serializing efficiently

A part of the program is *serialized* if it cannot be run by multiple threads concurrently

- For example, a critical section

*Serialization* is the mechanism used to guarantee this property

- For example, a lock

# Motivation

A serialized code section doesn't exploit parallelism, and so is “slow”

Why care about making it efficient?

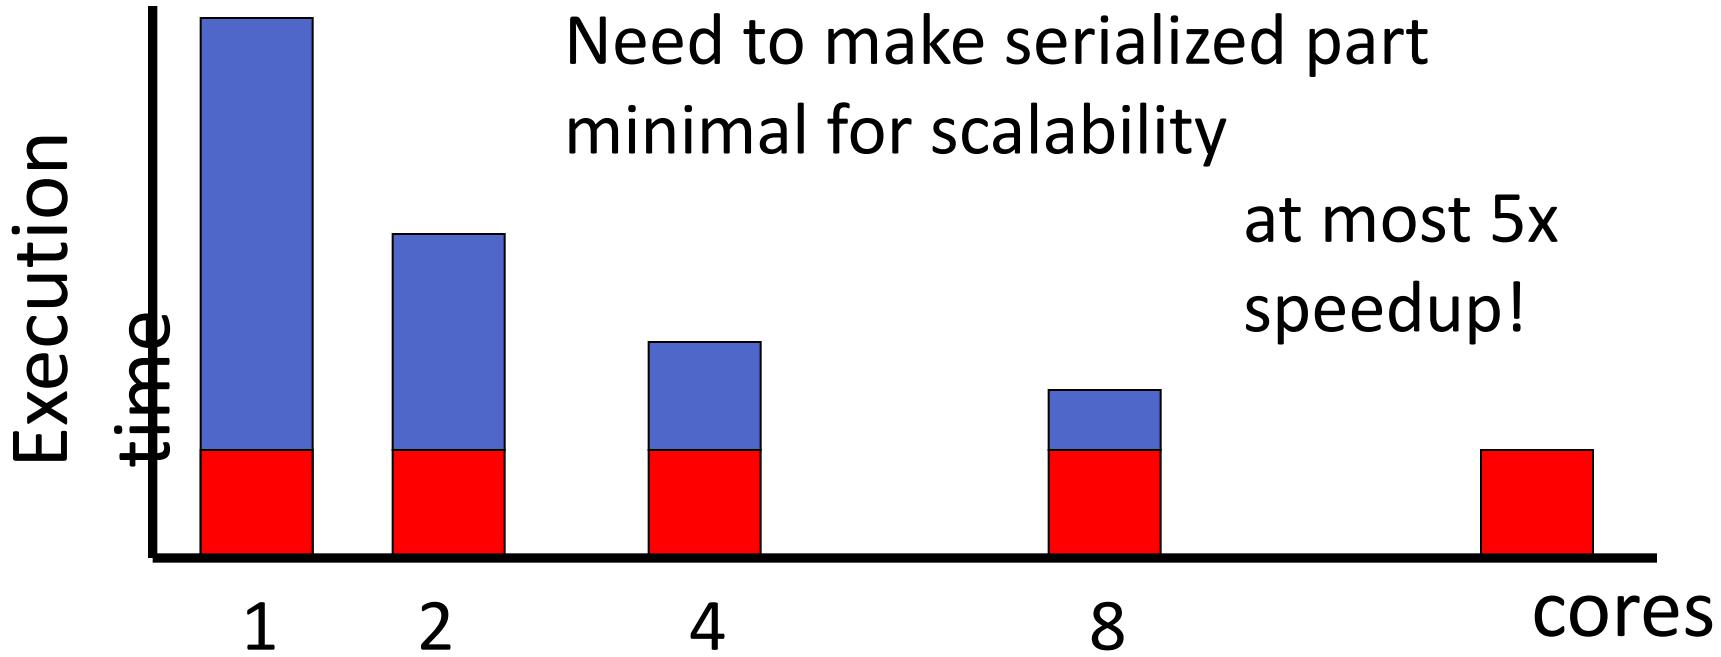
# Motivation: Amdahl's law

What is the effect of speeding up  $p$  of the execution time by  $s$ ?

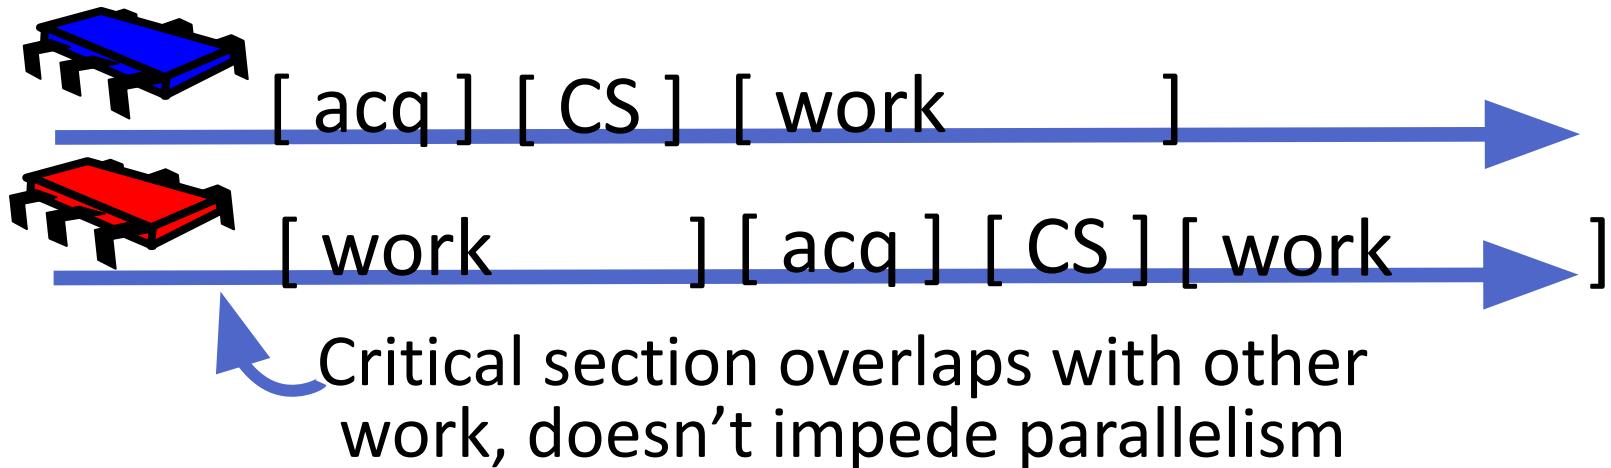
$$\text{speedup} = \frac{1}{(1-p) + p/s}$$

# Amdahl's law

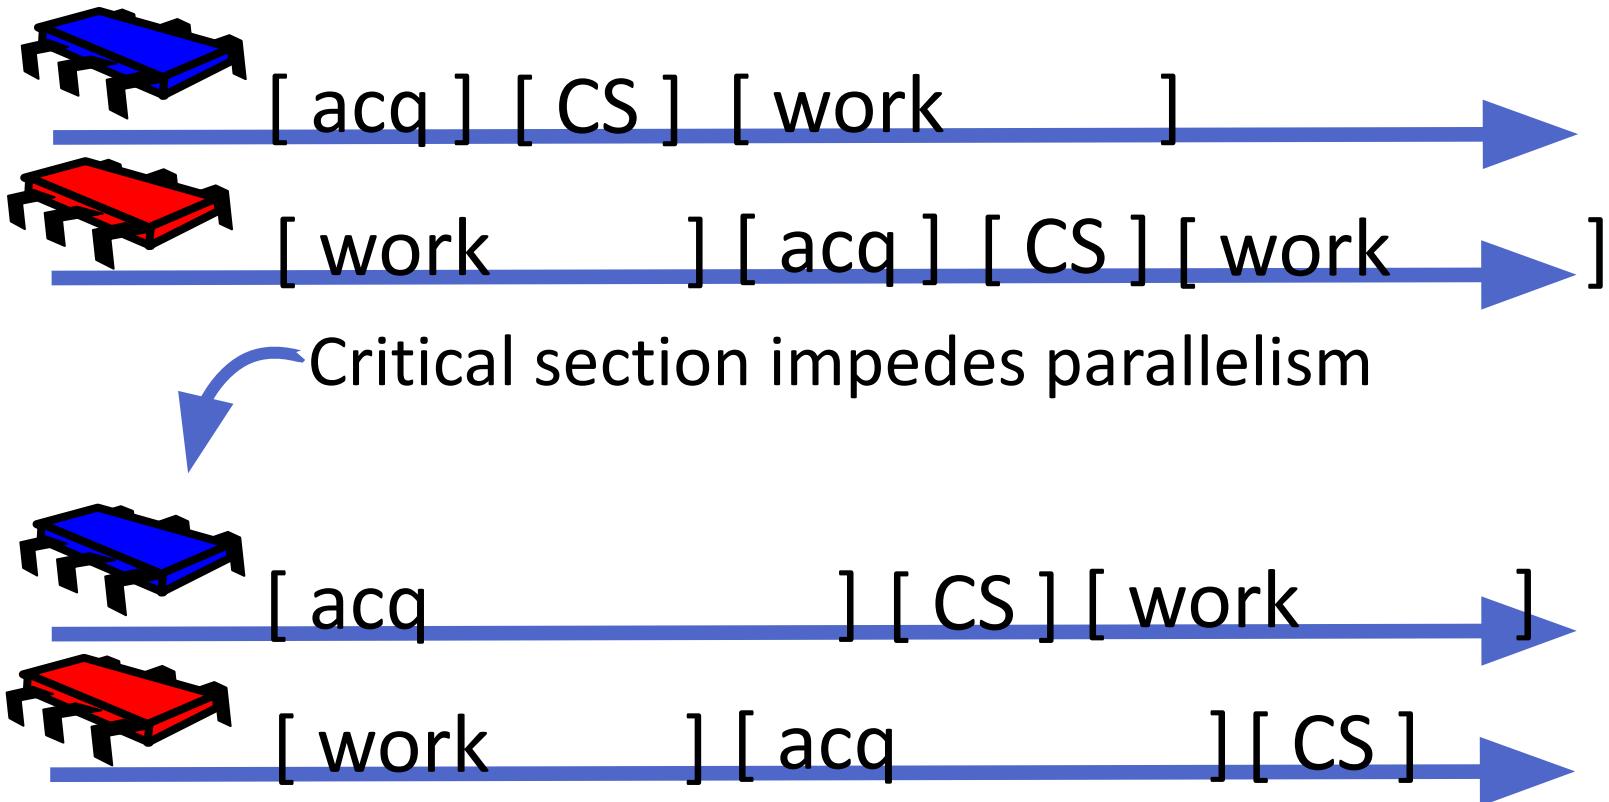
80% of execution parallelizable:



# Another way of looking at it



# Another way of looking at it

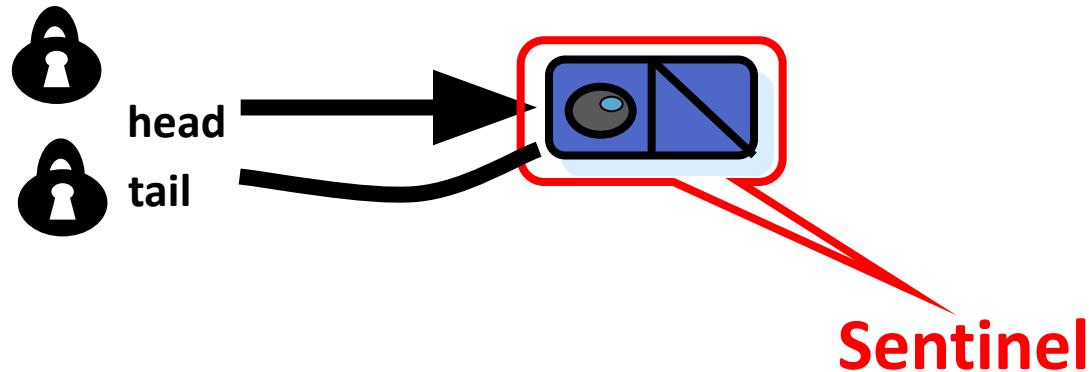


# Today

- Lock implementation issues
- Delegation locking
- Lock-free synchronization and CAS failures
- Running example: a FIFO queue (unbounded)

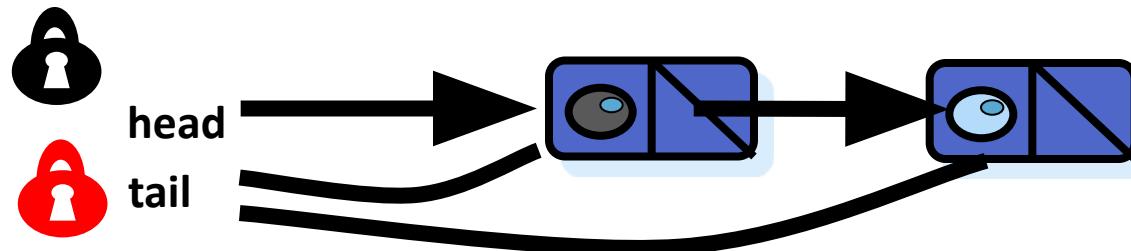
# Two-lock queue

[Michael & Scott '96]

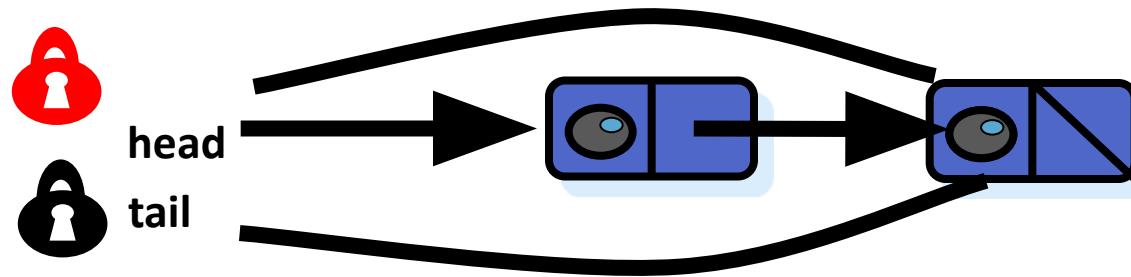


- Allows concurrent enqueue & dequeue operations in most situations

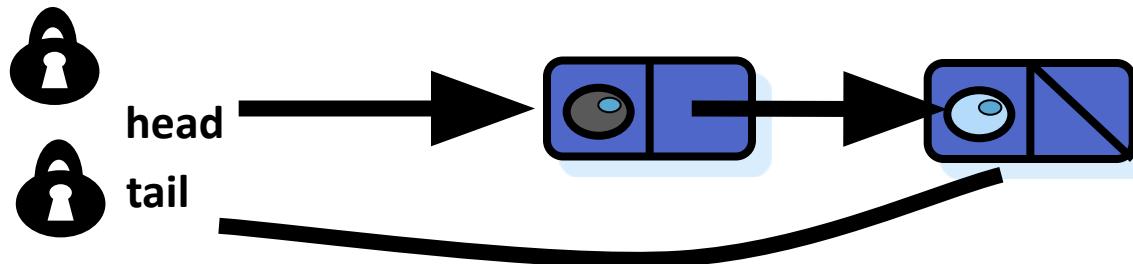
# Two-lock queue: enqueue



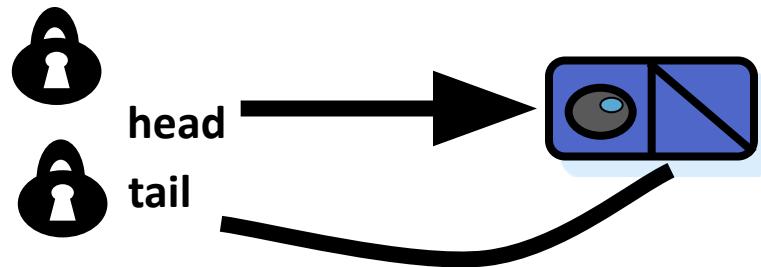
# Two-lock queue: dequeue



# Concurrent enqueue & dequeue



# Concurrent enqueue & dequeue



# Two-lock queue

```
Node { void* val; Node* next; };
```

Initially:

```
head = tail = new Node(NULL);  
head->next = NULL;
```

# Two-lock queue

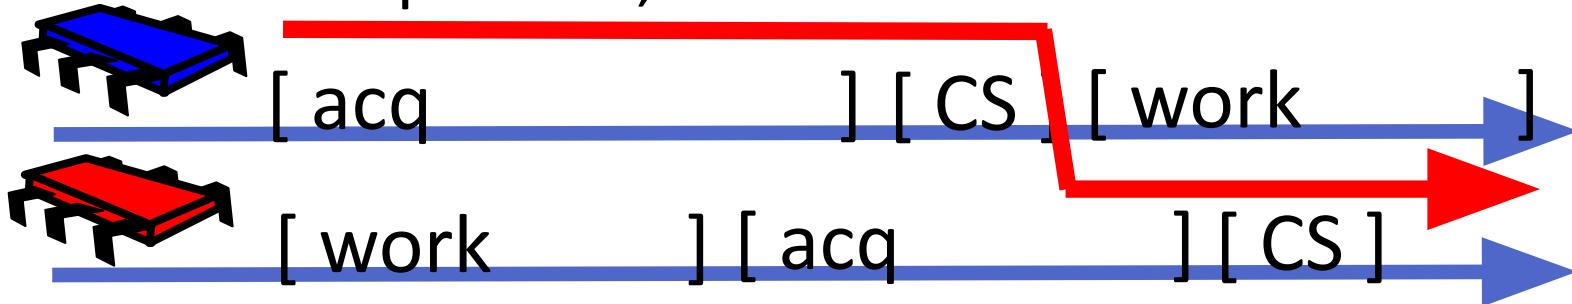
```
enqueue(void* v) {  
    Node* node = new Node(v);  
    node->next = NULL;  
    lock(&tail_lock);  
    tail->next = node;  
    tail = node;  
    unlock(&tail_lock);  
}
```

# Two-lock queue

```
void* dequeue() {  
    lock(&head_lock);  
    Node* node = head;  
    Node* nh = head->next;  
    if (!nh) { unlock(..); return 0; }  
    head = nh;  
    void* rv = nh->val;  
    unlock(&head_lock);  
    free(node);
```

# Locks and critical sections

- Assumption: critical section is short
  - If not, the CS will be the bottleneck
- Goal: Minimize length of the *critical path*
  - Lock acquisition, CS & handoff



# TAS lock

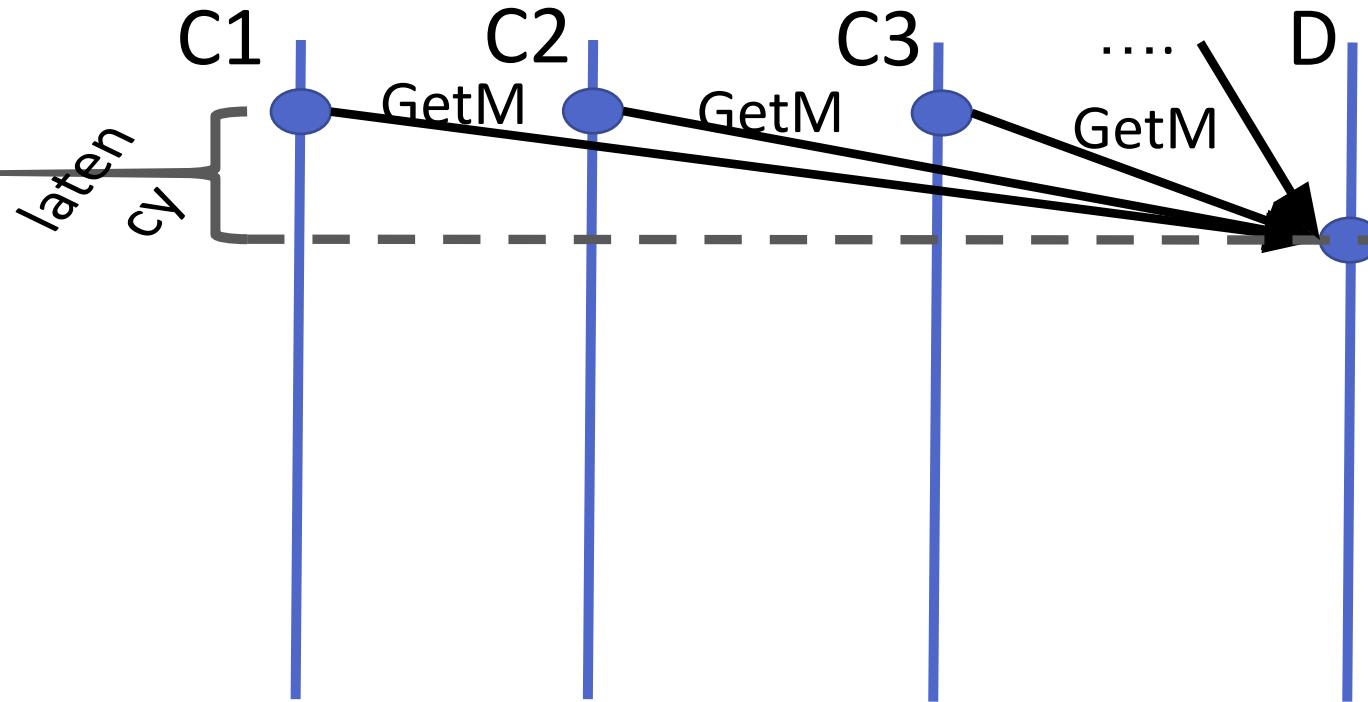
```
bool L;
```

```
lock() {
    while (!CAS(&L, 0, 1)) { }
}
```

```
unlock() {
    L = 0
}
```

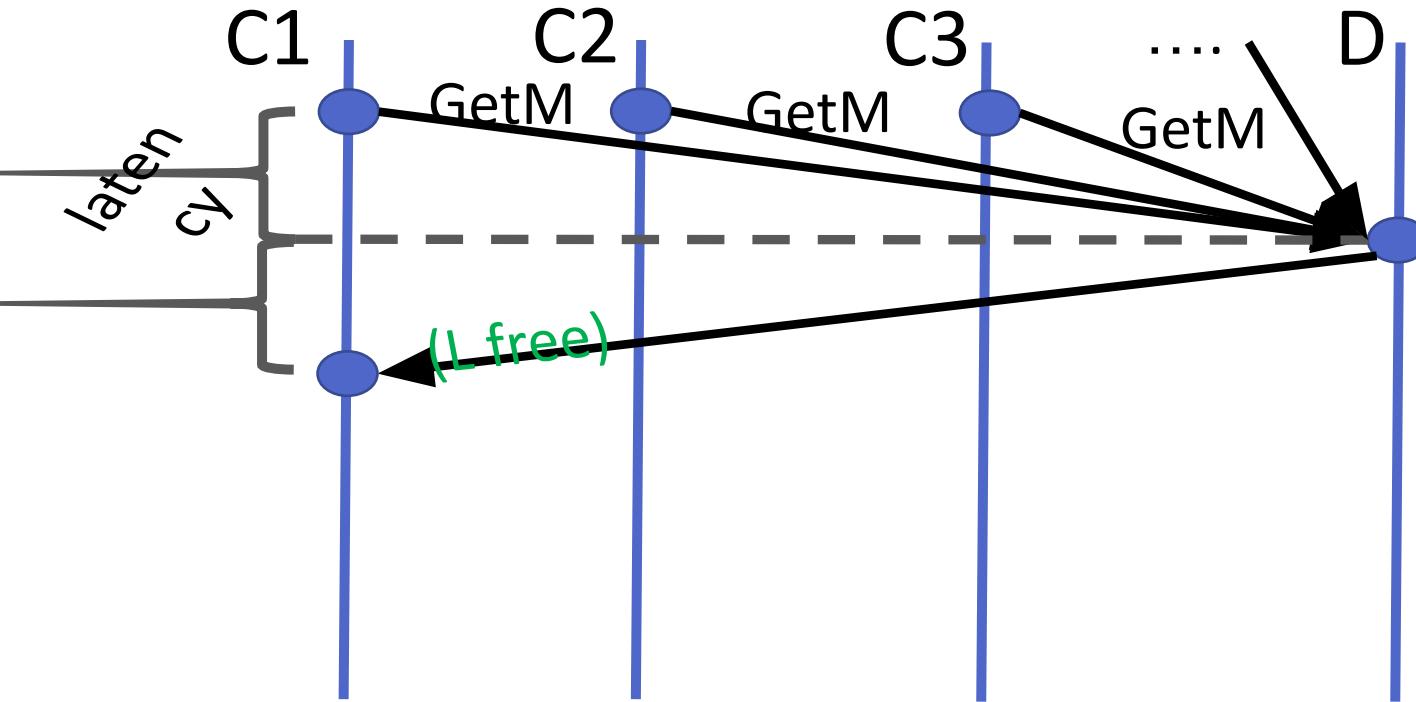
# TAS analysis

Cores try to acquire concurrently



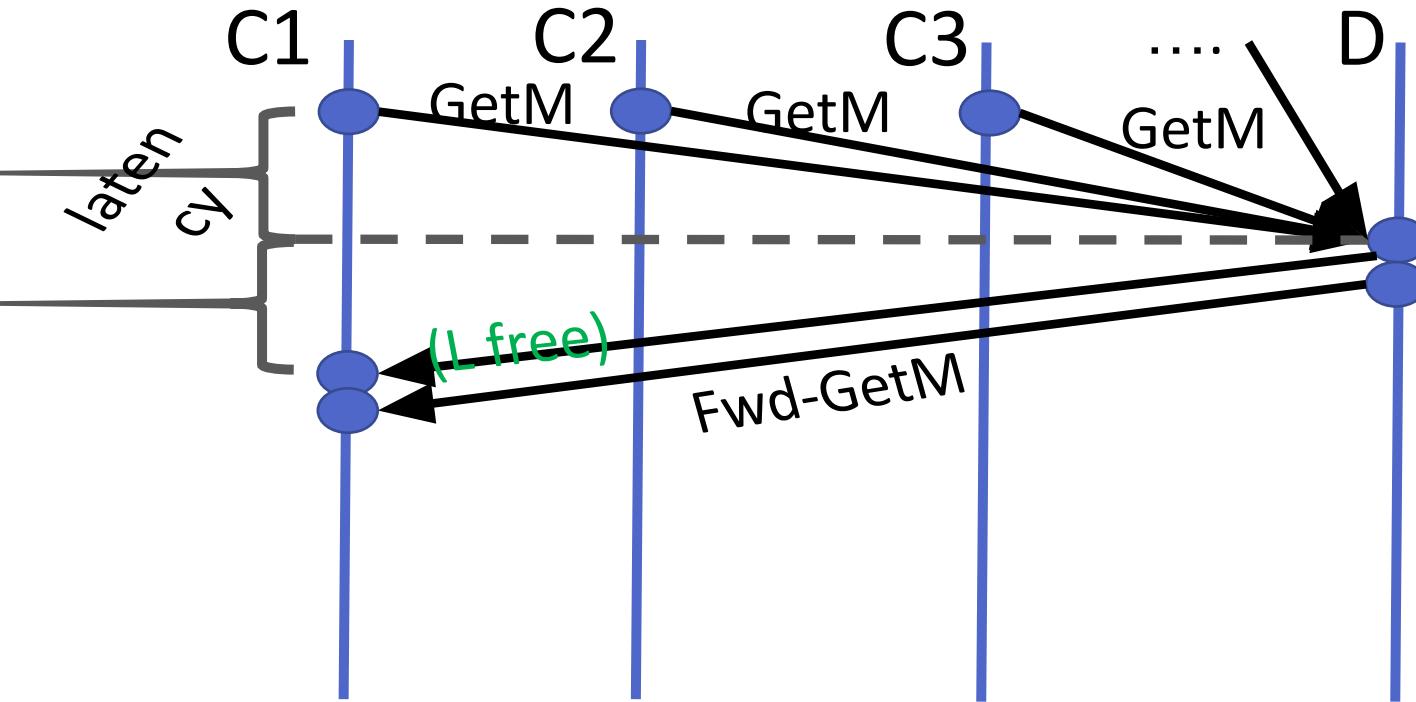
# TAS analysis

Directory serializes requests



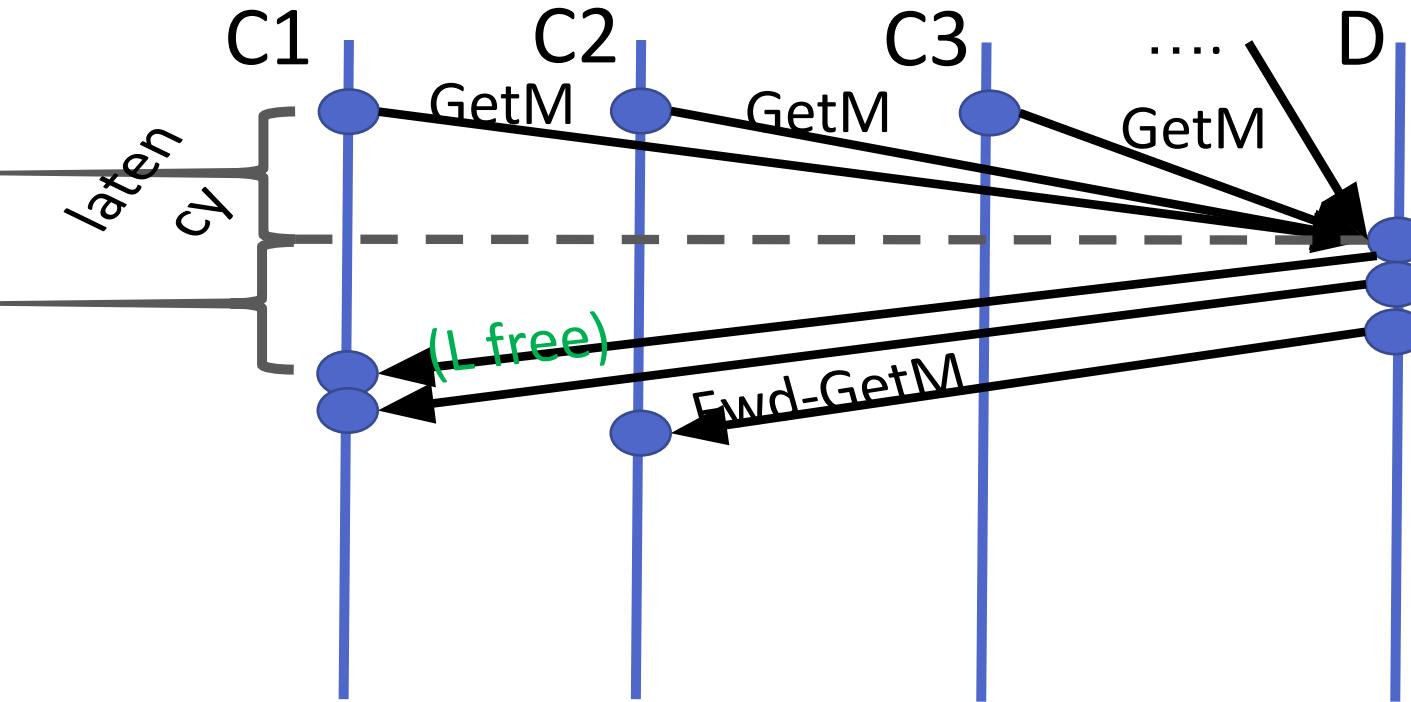
# TAS analysis

Directory serializes requests



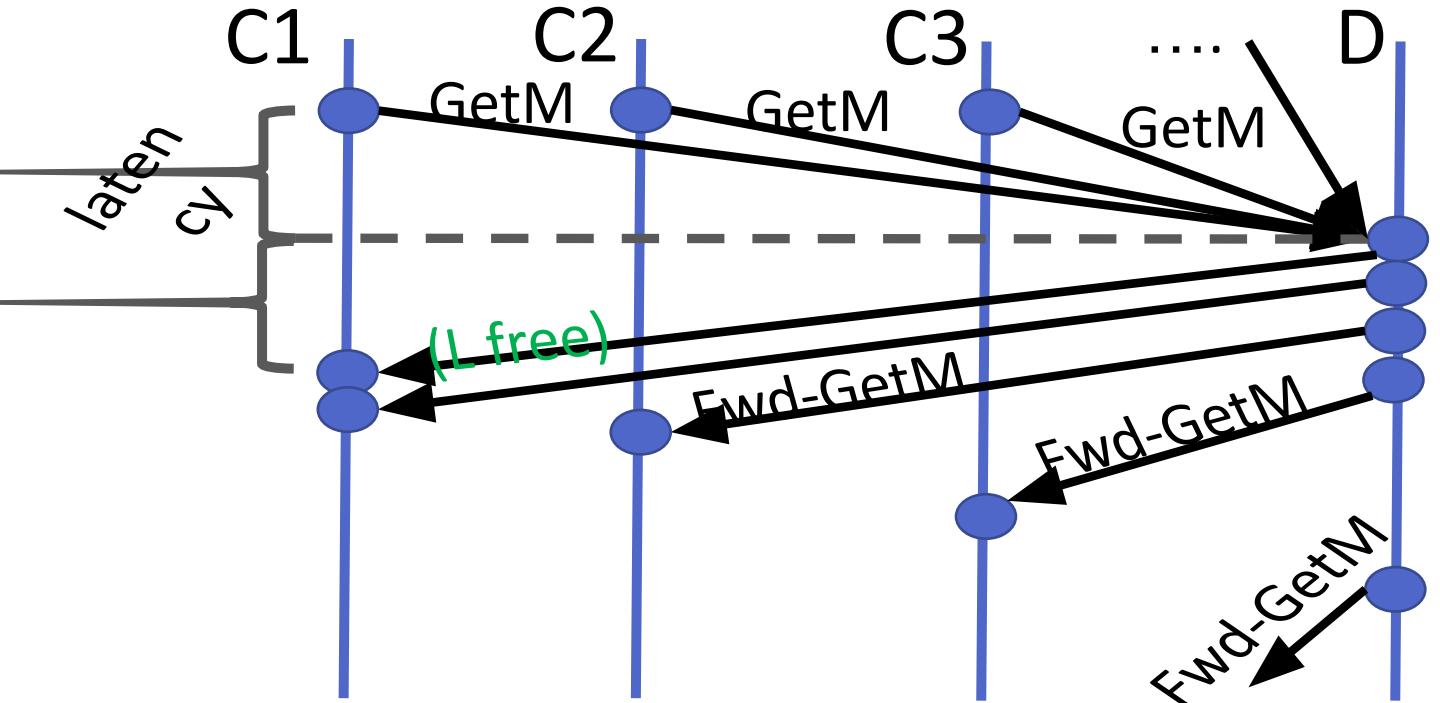
# TAS analysis

Directory serializes requests



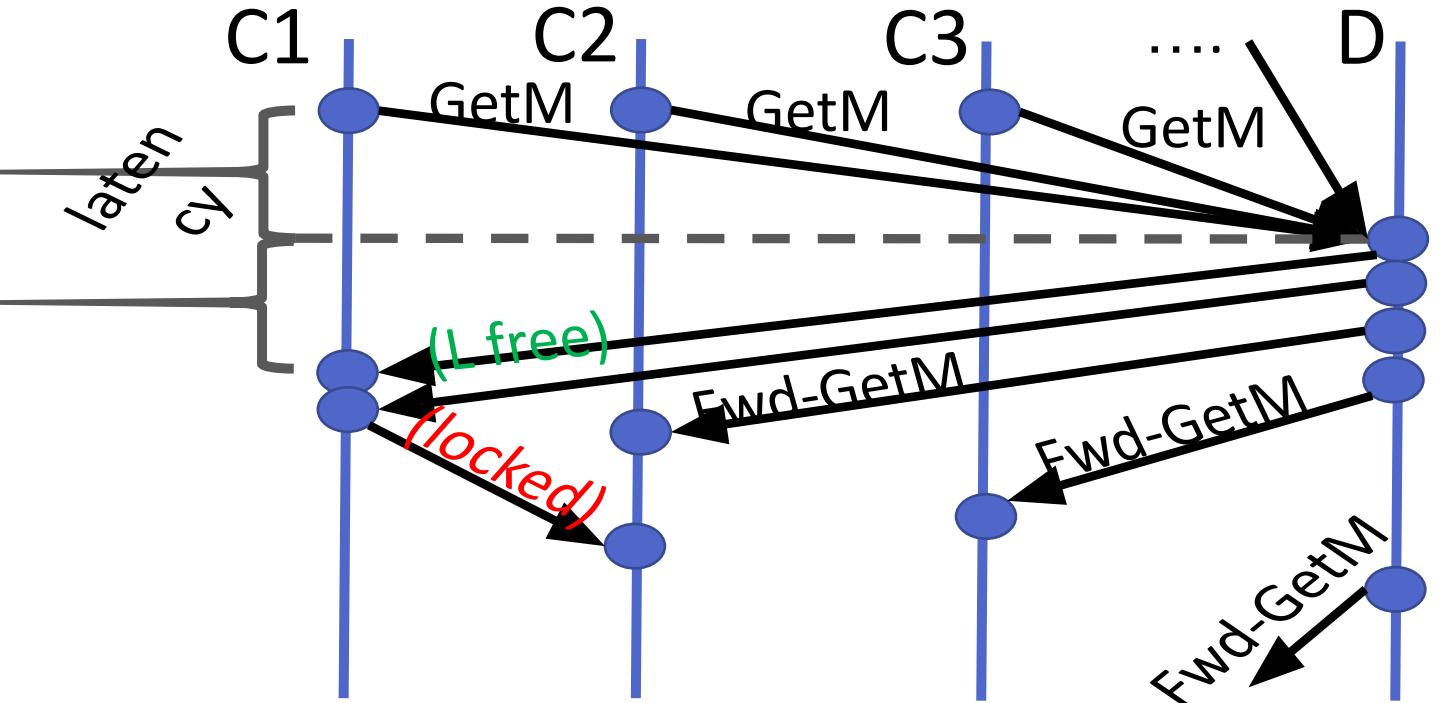
# TAS analysis

Directory serializes requests



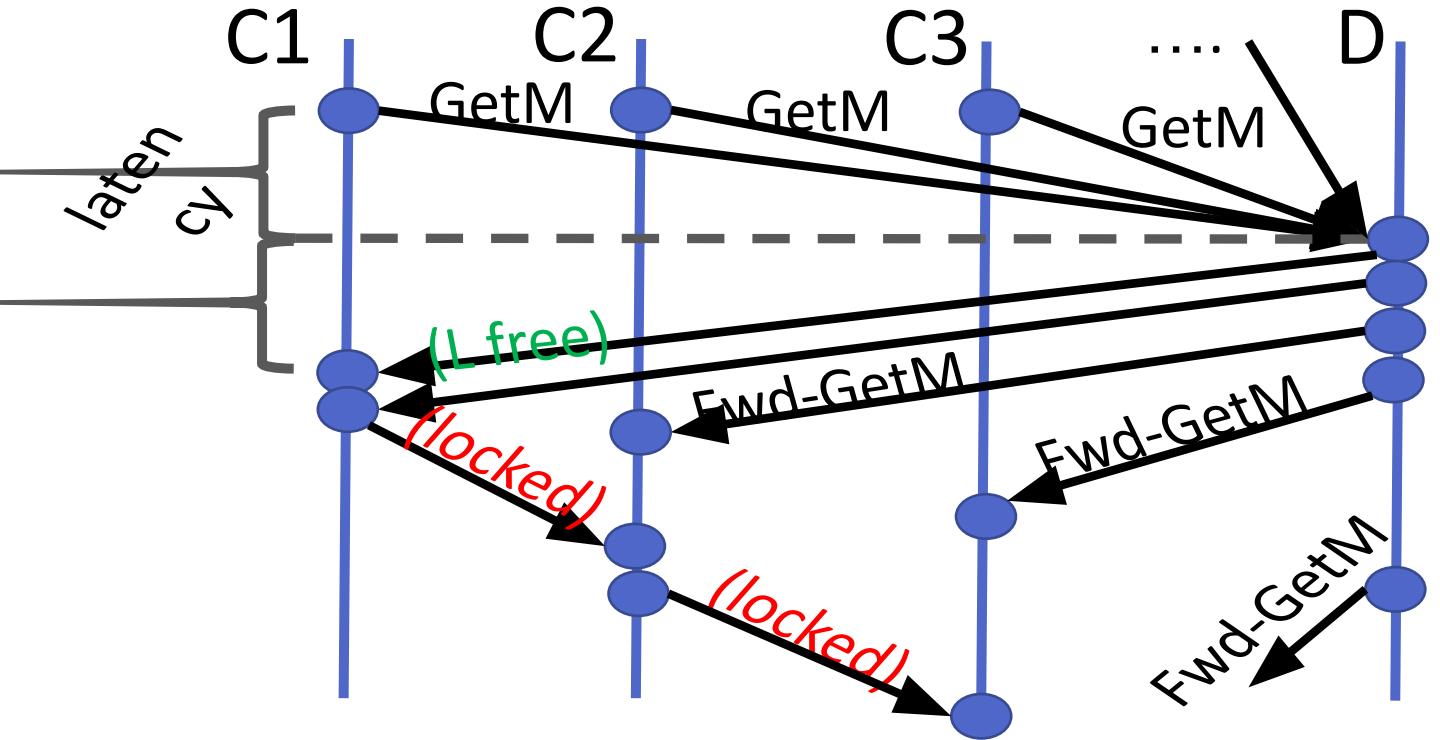
# TAS analysis

Owner gets invalidated



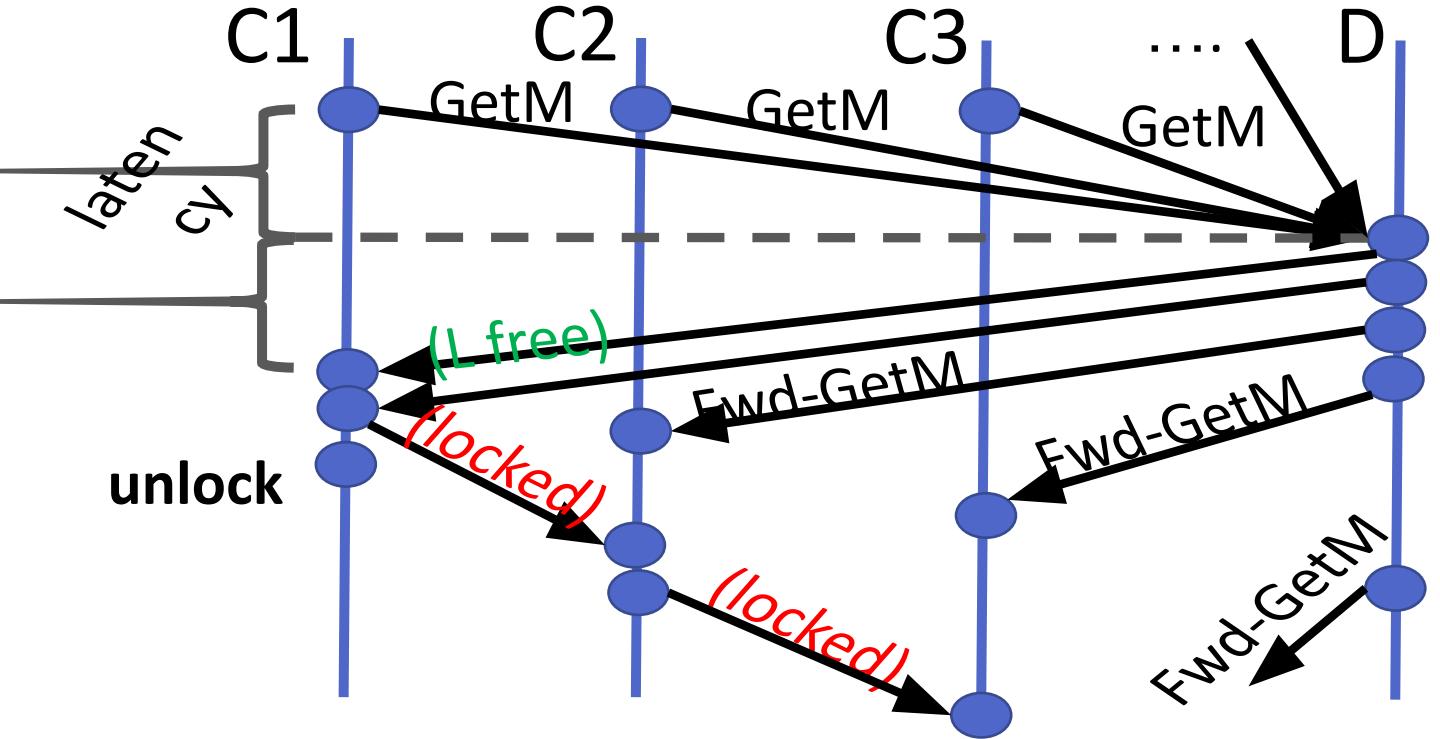
# TAS analysis

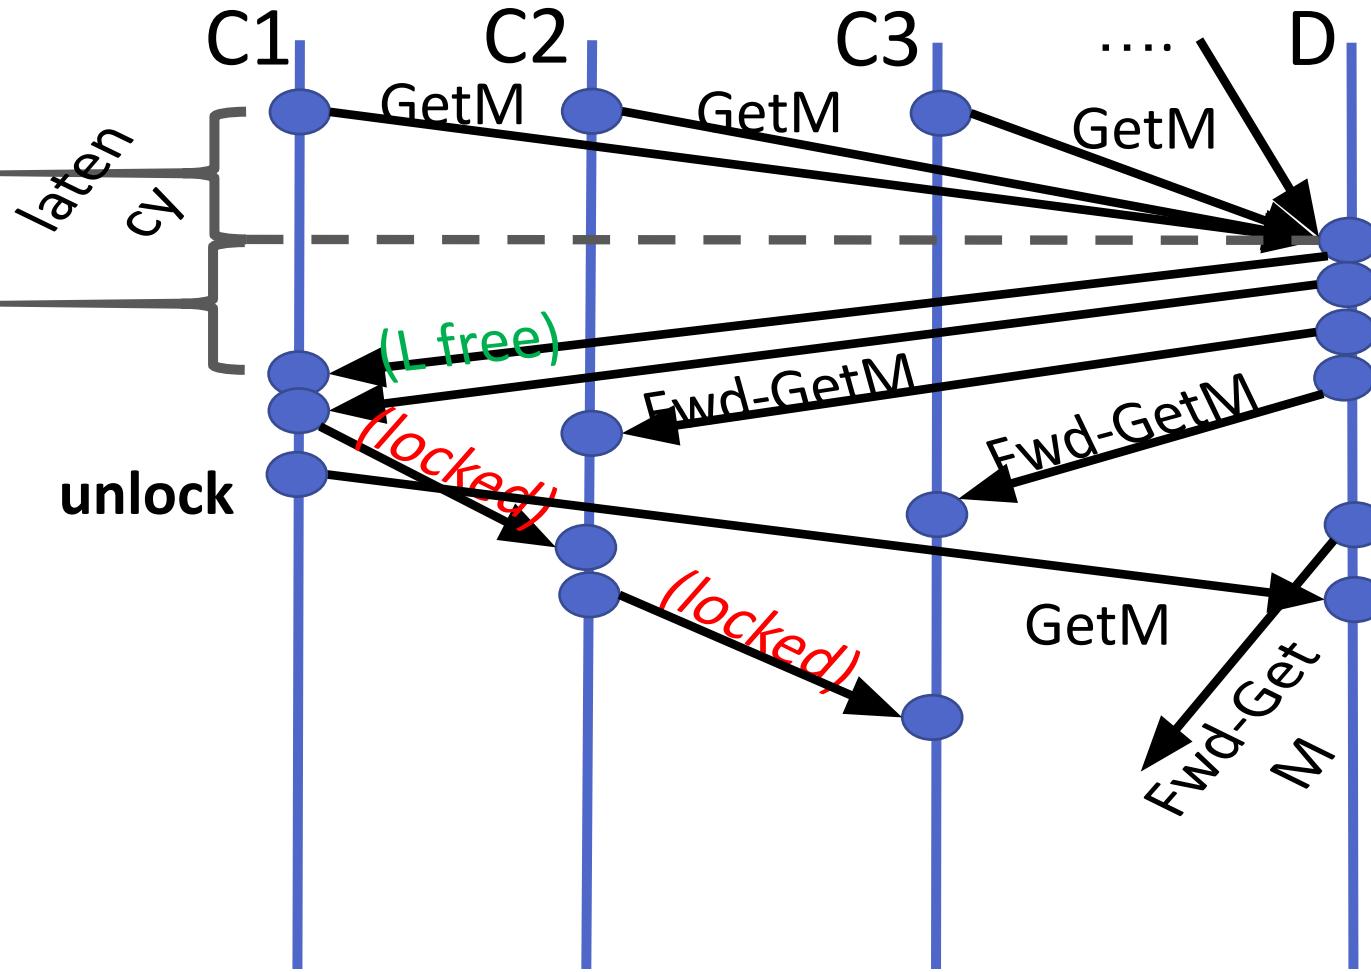
2<sup>nd</sup> in line gets invalidated



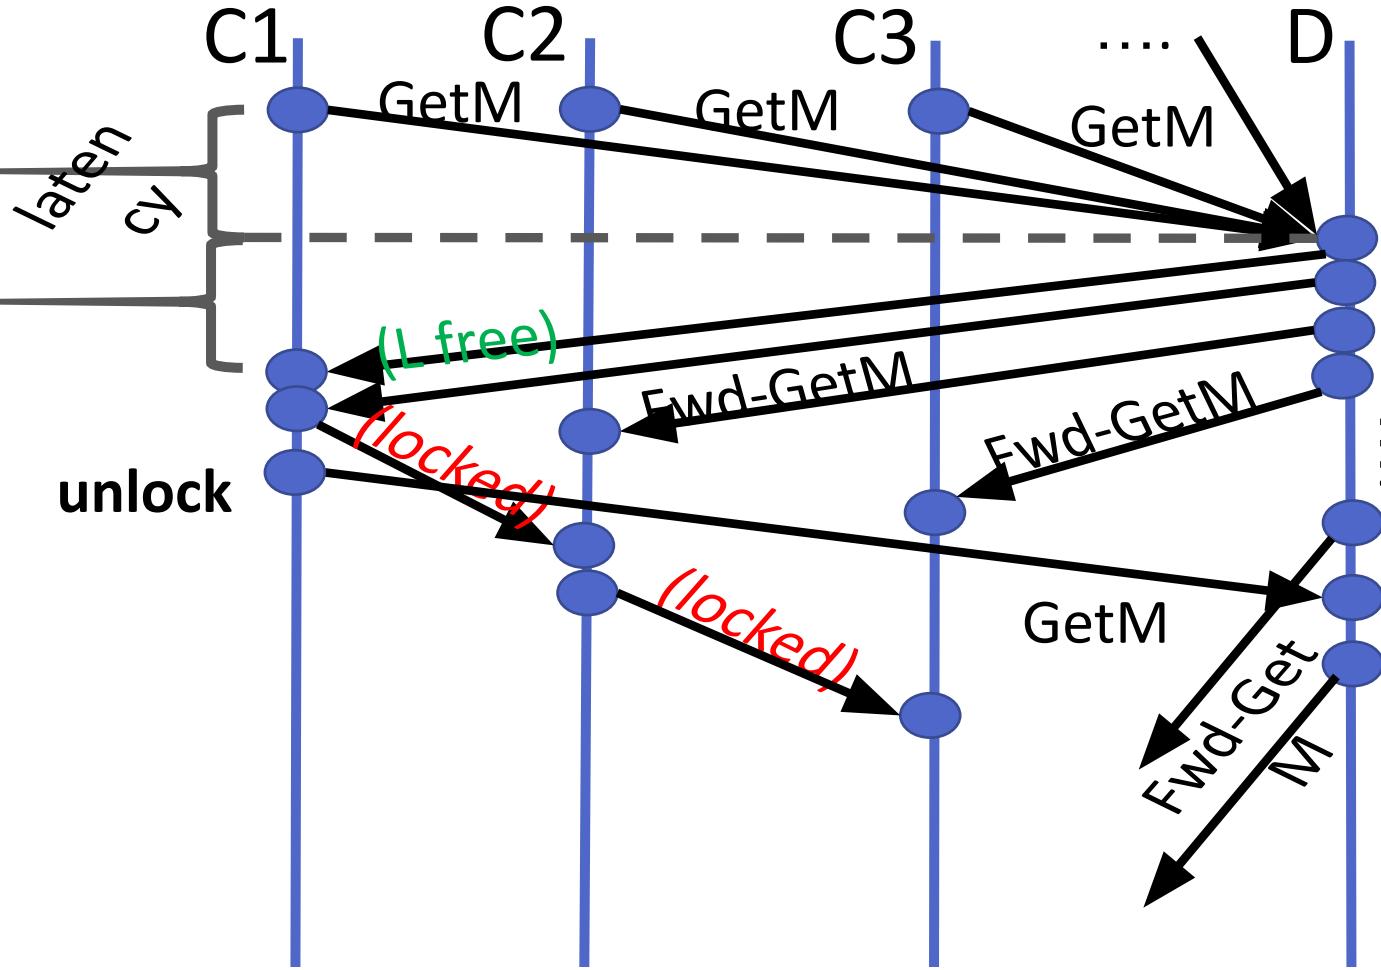
# TAS analysis

Owner finishes CS, tries to unlock





Unlock GetM is serialized after all acquire attempts



Unlock  
GetM is  
serialized  
after all  
acquire  
attempts

# TAS analysis

- Coherence transactions by (failing) lock acquisitions create an *invalidation storm*, which forms an implicit queue of cores waiting for ownership of the line
- An unlock gets put at the end of the queue, and thus takes linear ( $N*L$ ) time
  - $N = \#$  of contending cores
  - $L =$  invalidation latency
  - Example multi-core case:  $N=20, L=20$

# TTAS lock

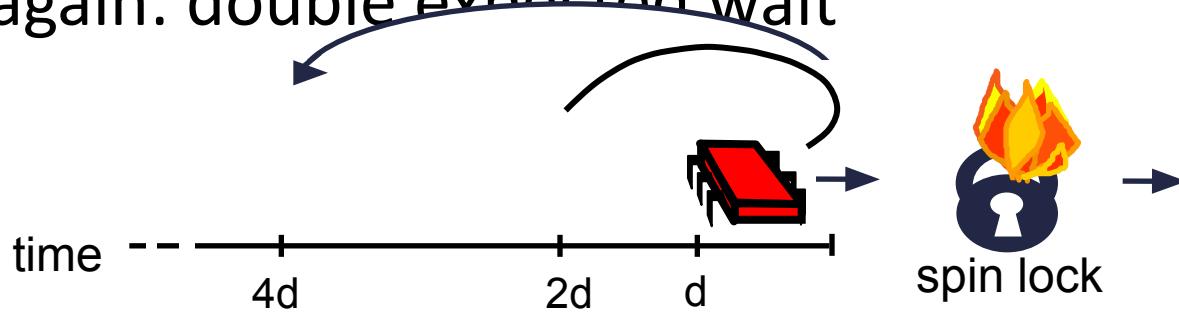
```
lock() {  
    while (true) {  
        while (L) { }  
        if (CAS(&L, 0, 1)) return;  
    }  
}  
  
unlock() {
```

# TTAS analysis

- If all cores see lock free, all try to acquire it
  - ⇒ An invalidation storm as in TAS lock
- If CS is short (< length of N invalidations), the invalidation storm delays the lock release
- And thus the next acquisition
- And therefore increases the critical path length

# Backoff

- Possible TTAS solution: exponential backoff
- If fail to lock, wait random duration before retry
- Fail again: double ~~expected~~ wait



- Very hard to get right in practice, requires tuning and

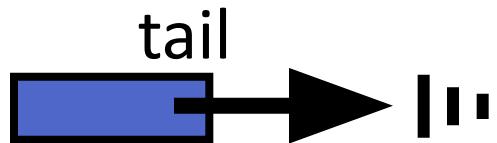
# MCS lock [Mellor-Crummey & Scott '91]

- Goal: lock handoff in  $O(1)$  invalidations
- Put threads in explicit queue, so that unlock can transfer ownership to thread waiting at queue head
  - owne
  - waiting
  - waiting



- Waiters do *local spinning* & don't get invalidated until reaching queue head

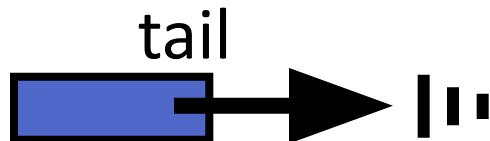
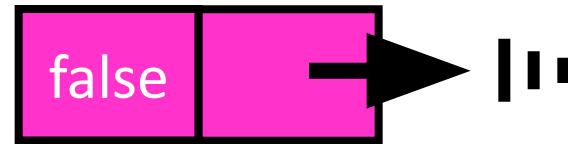
# MCS lock [Mellor-Crummey & Scott '91]



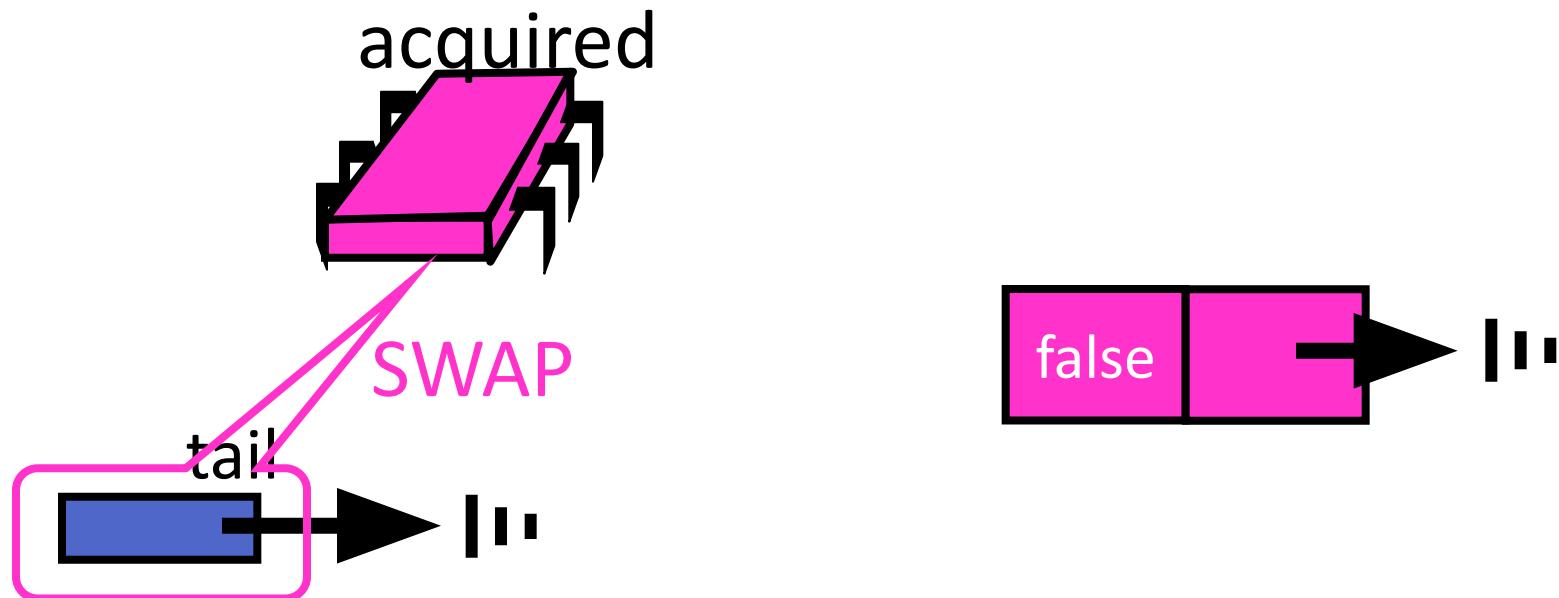
# MCS lock [Mellor-Crummey & Scott '91]



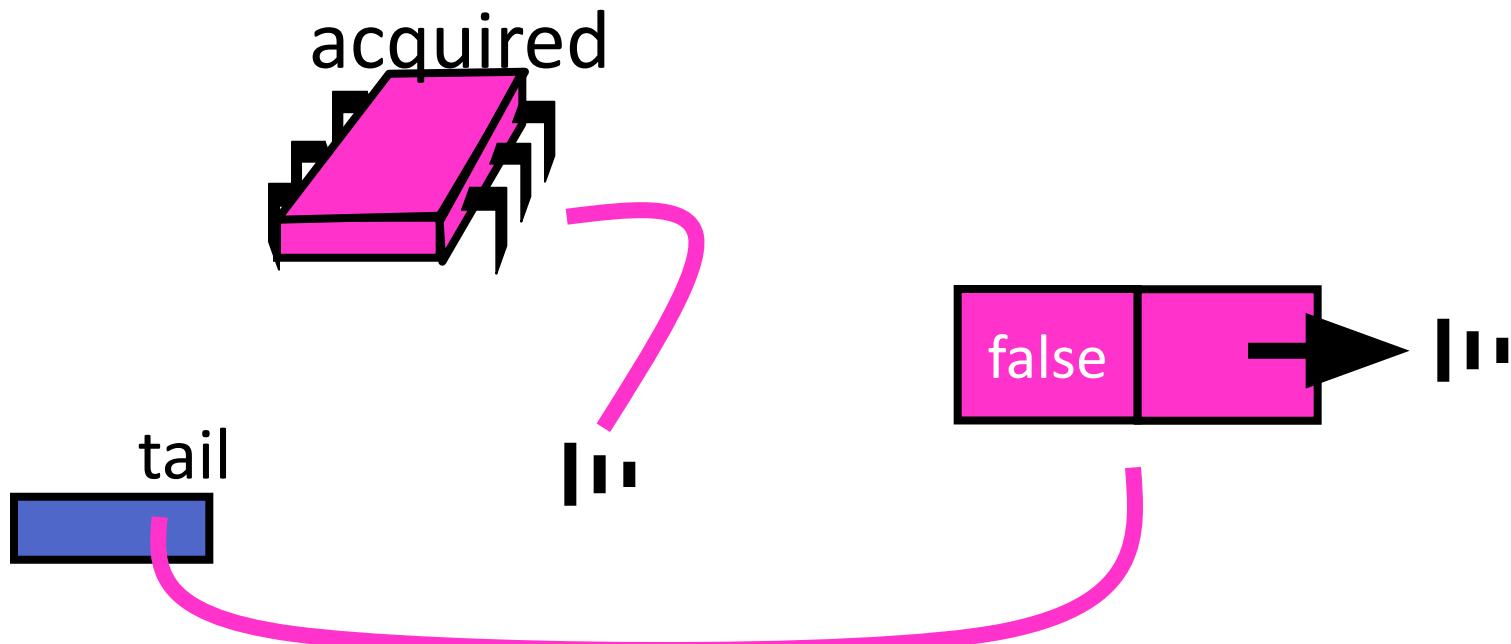
(allocate Qnode)



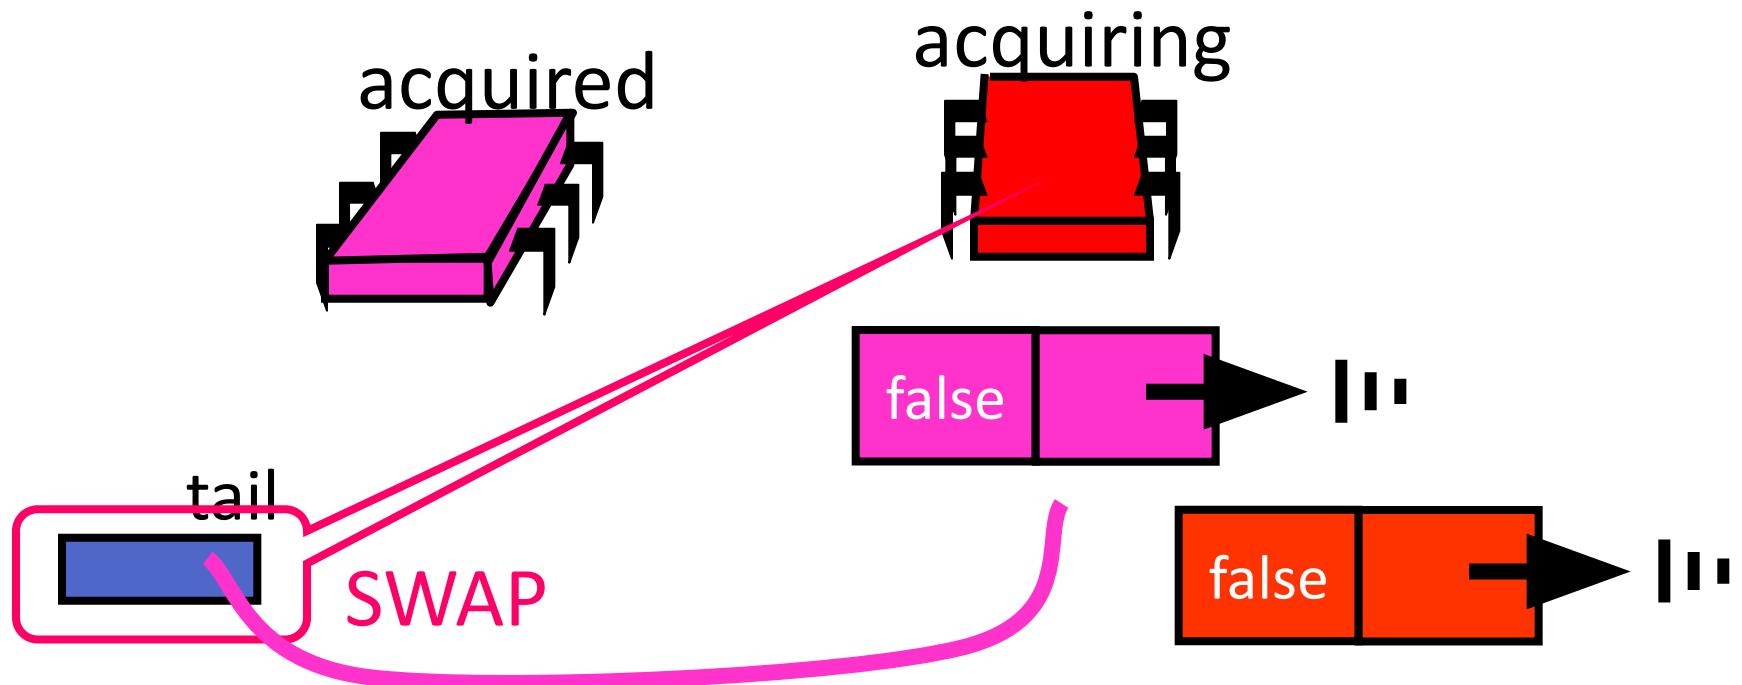
# MCS lock [Mellor-Crummey & Scott '91]



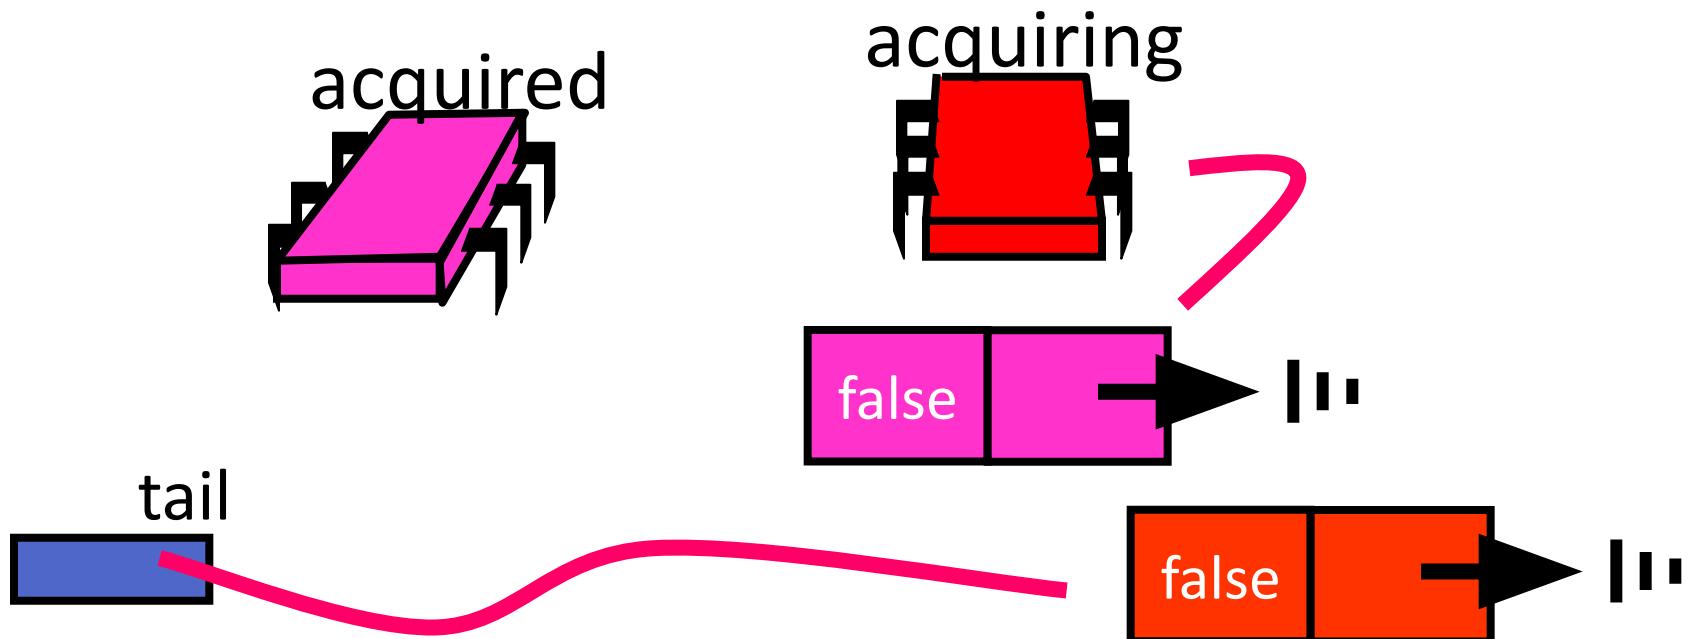
# MCS lock [Mellor-Crummey & Scott '91]



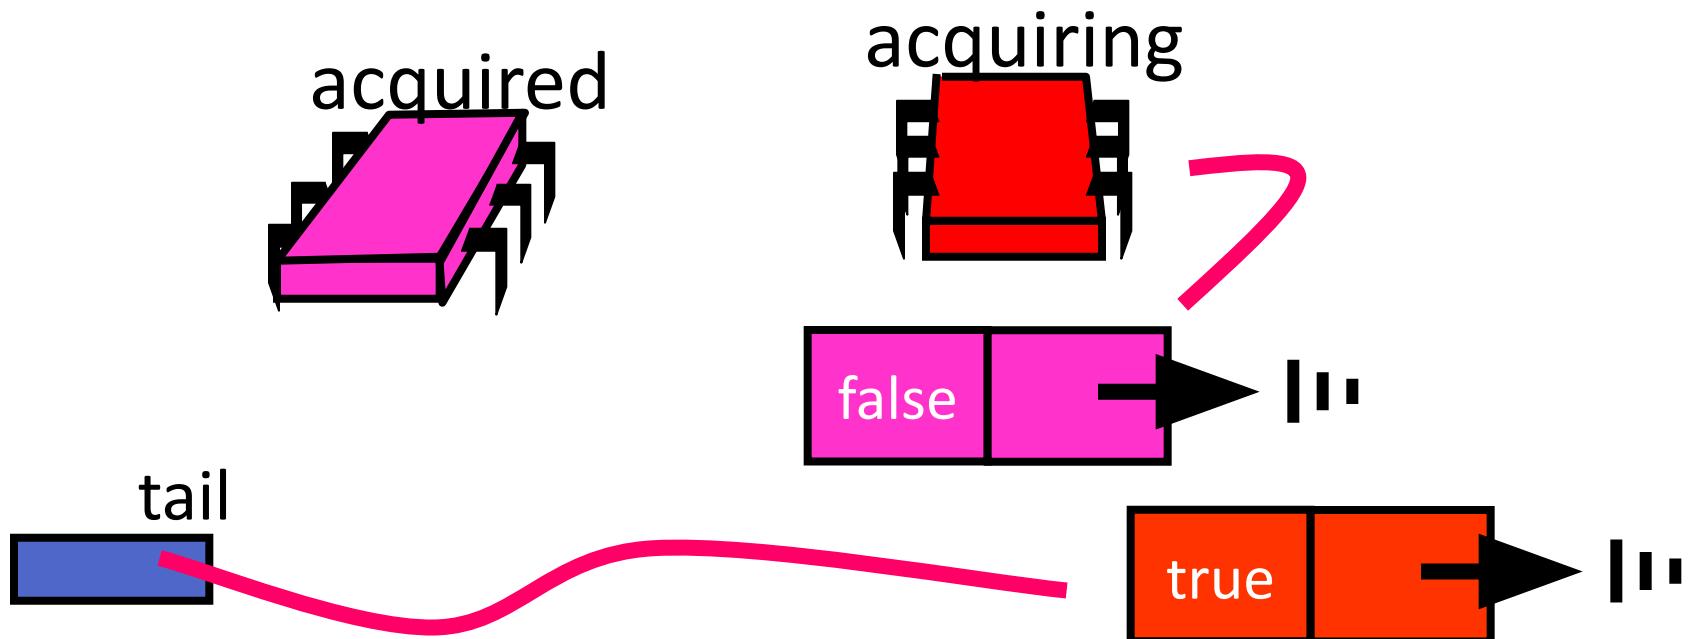
# MCS lock [Mellor-Crummey & Scott '91]



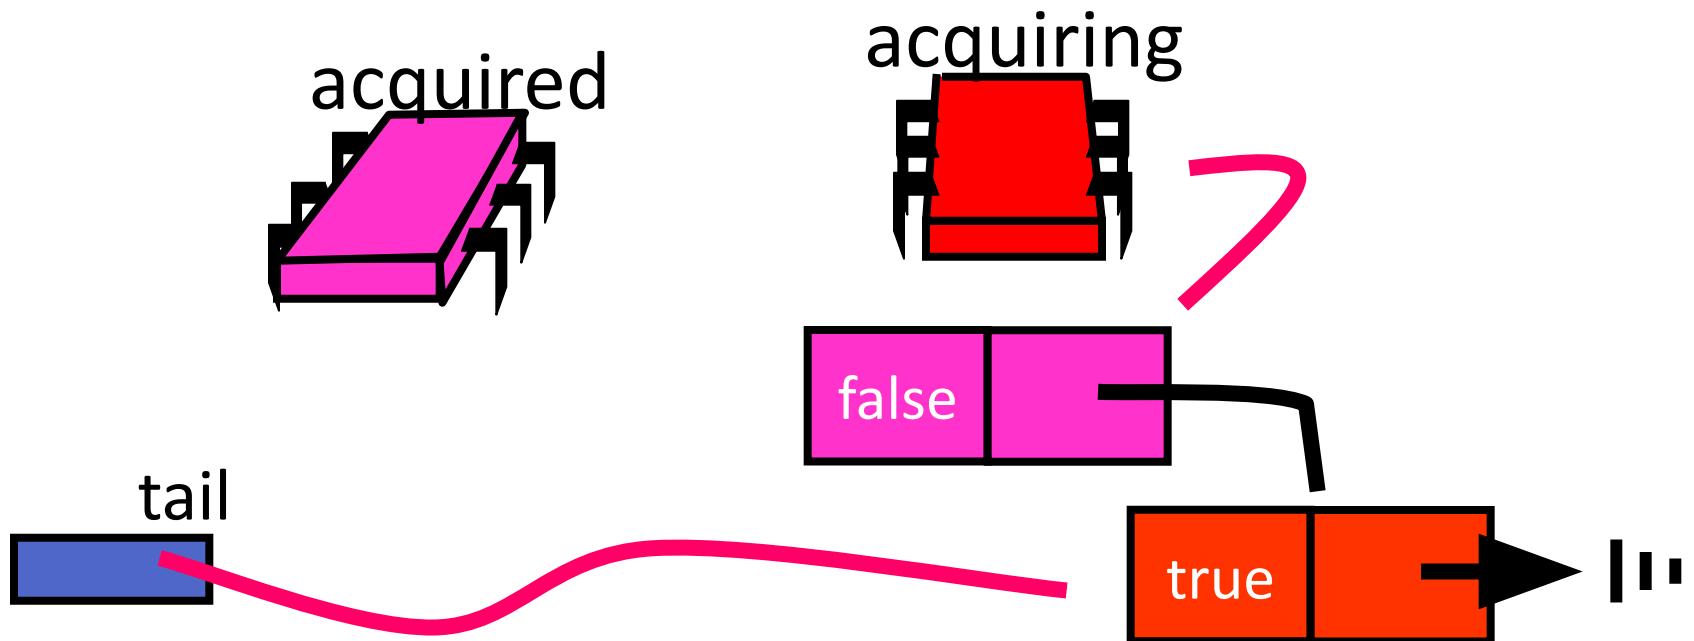
# MCS lock [Mellor-Crummey & Scott '91]



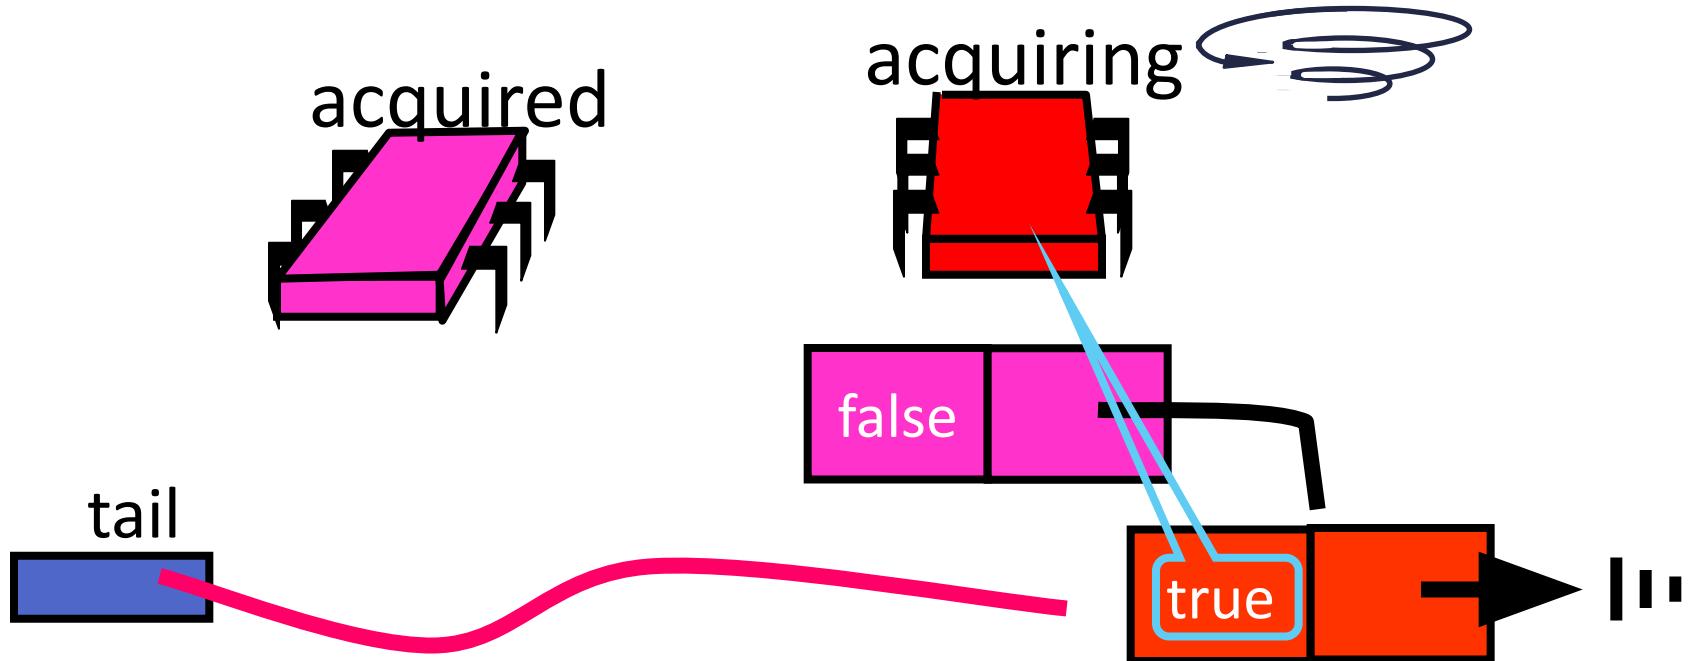
# MCS lock [Mellor-Crummey & Scott '91]



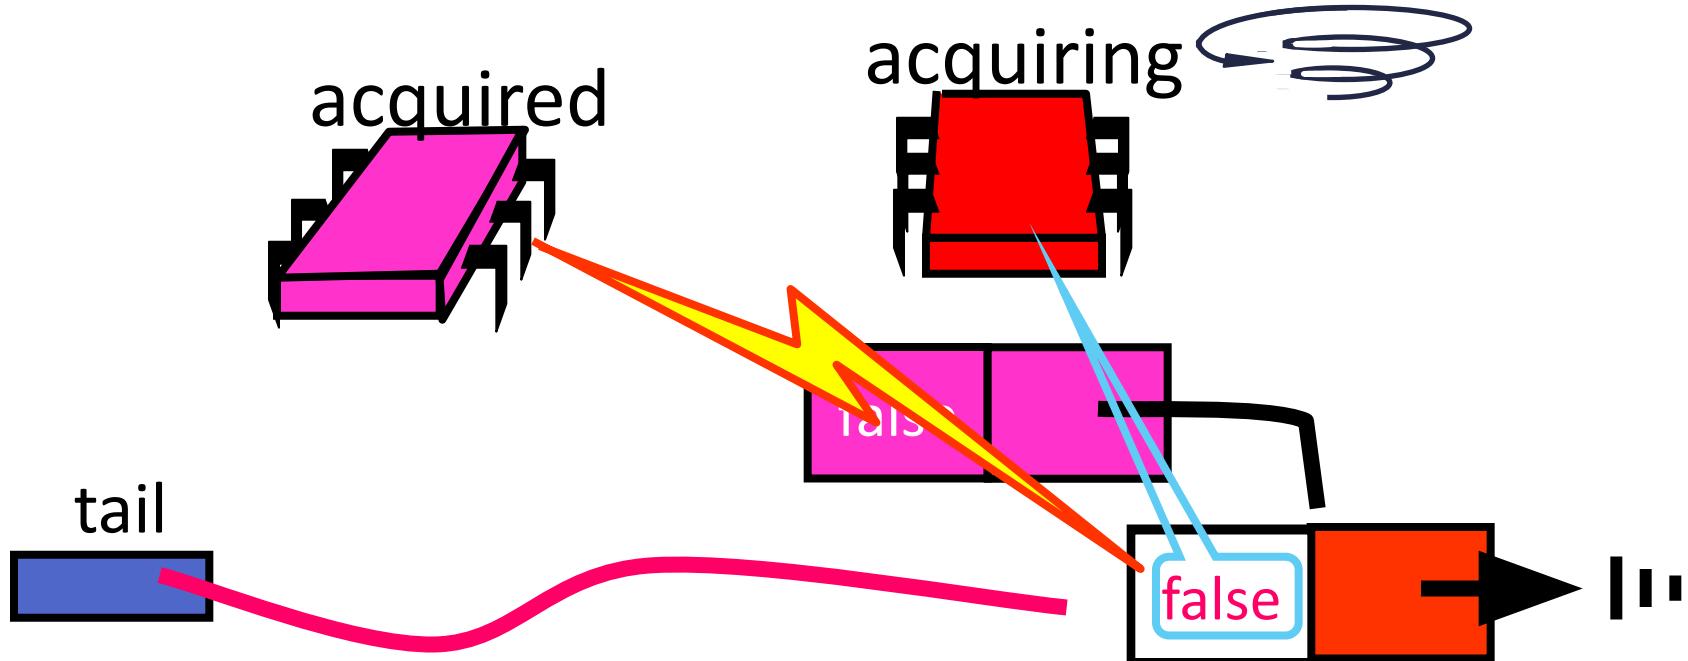
# MCS lock [Mellor-Crummey & Scott '91]



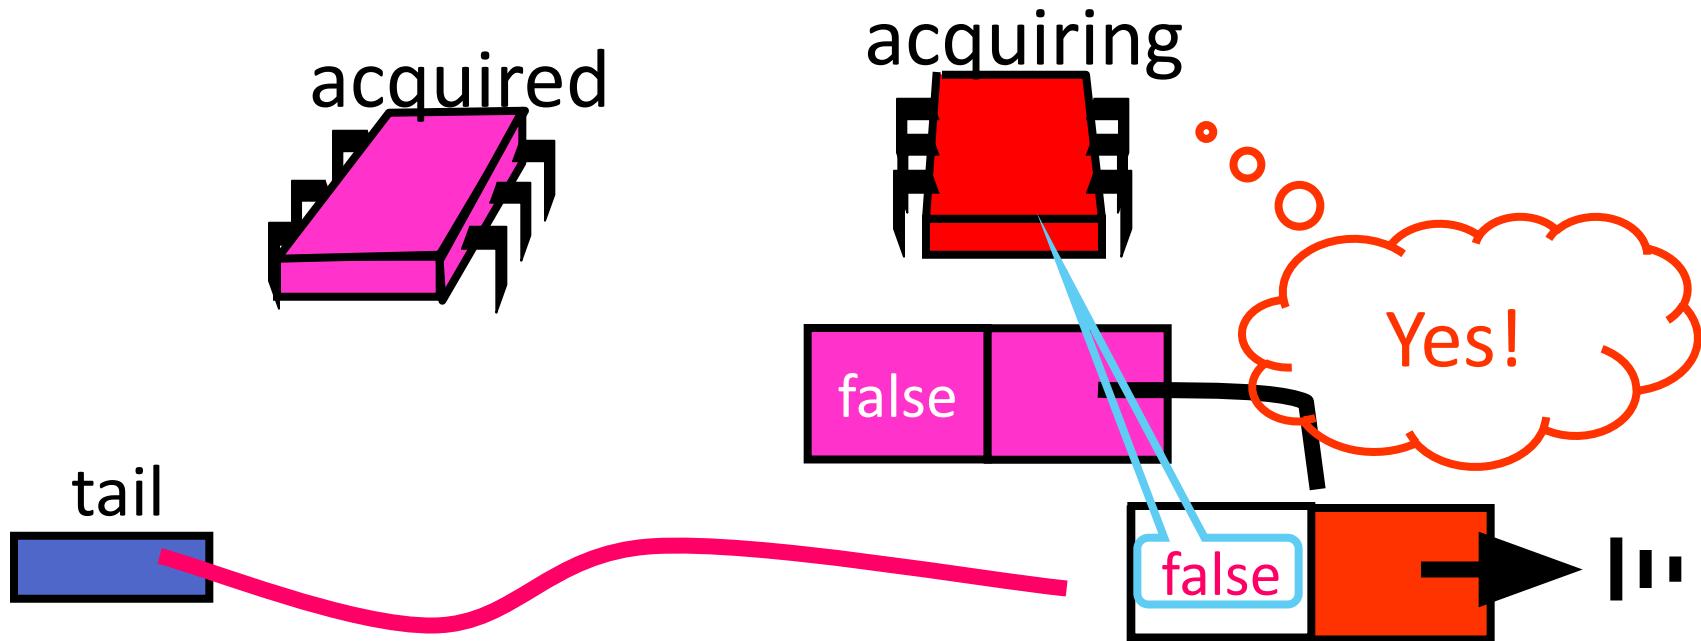
# MCS lock [Mellor-Crummey & Scott '91]



# MCS lock [Mellor-Crummey & Scott '91]



# MCS lock [Mellor-Crummey & Scott '91]

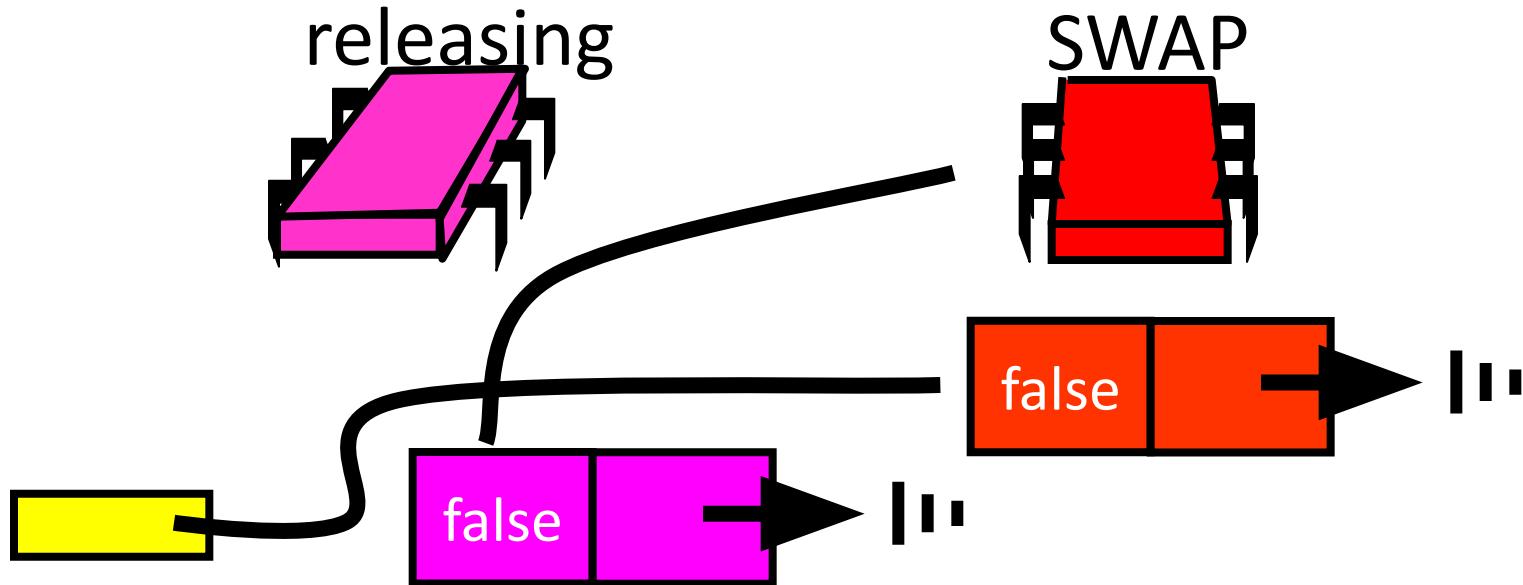


# MCS lock

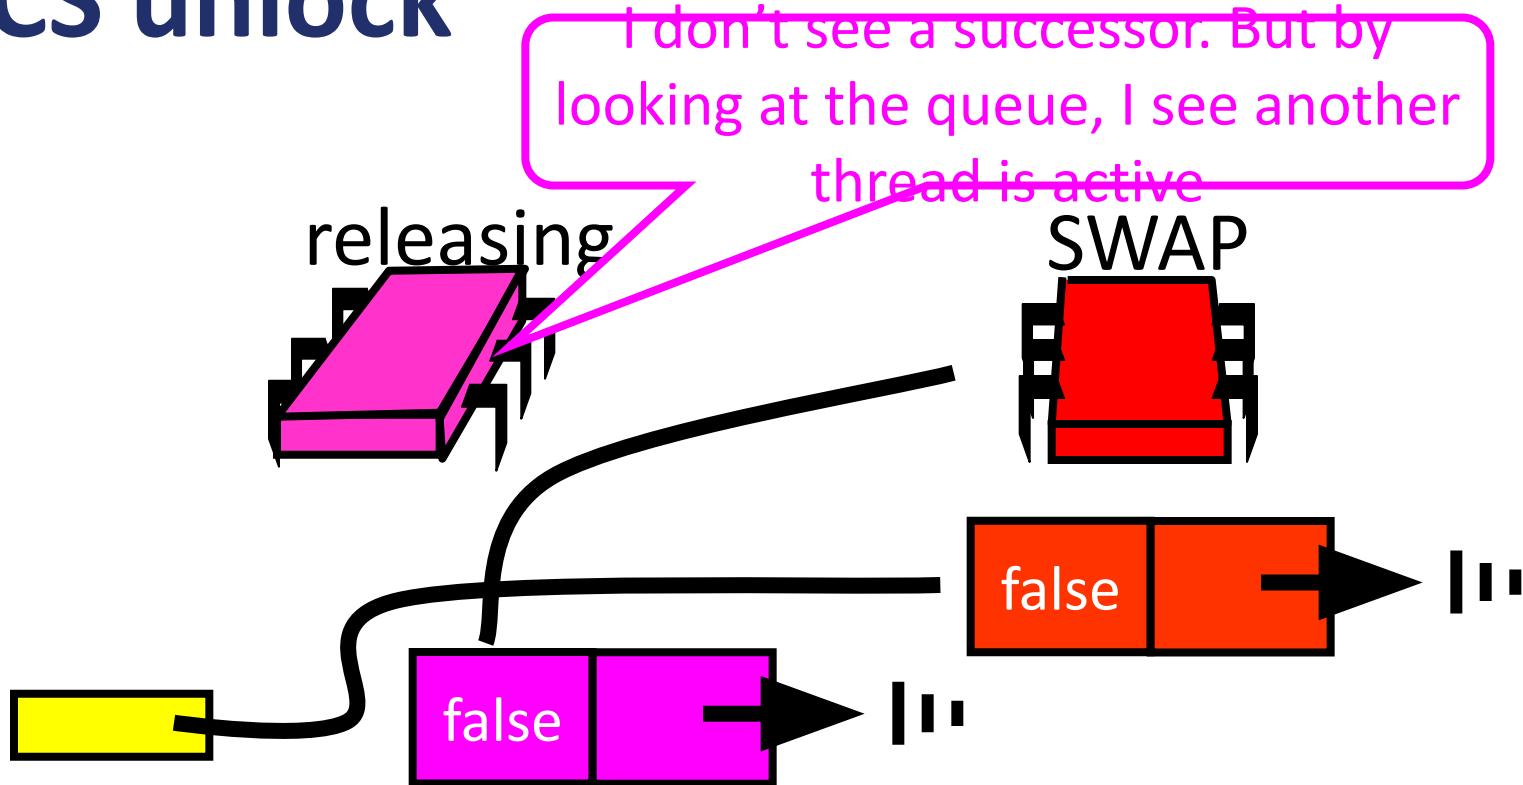
```
qnode : { qnode* next; bool lock; }
```

```
lock(qnode *n) {
    n->next = NULL
    pred = SWAP(&LOCK, n)
    if (pred) {
        n->lock = true
        pred->next = n
        while (n->lock) { }
```

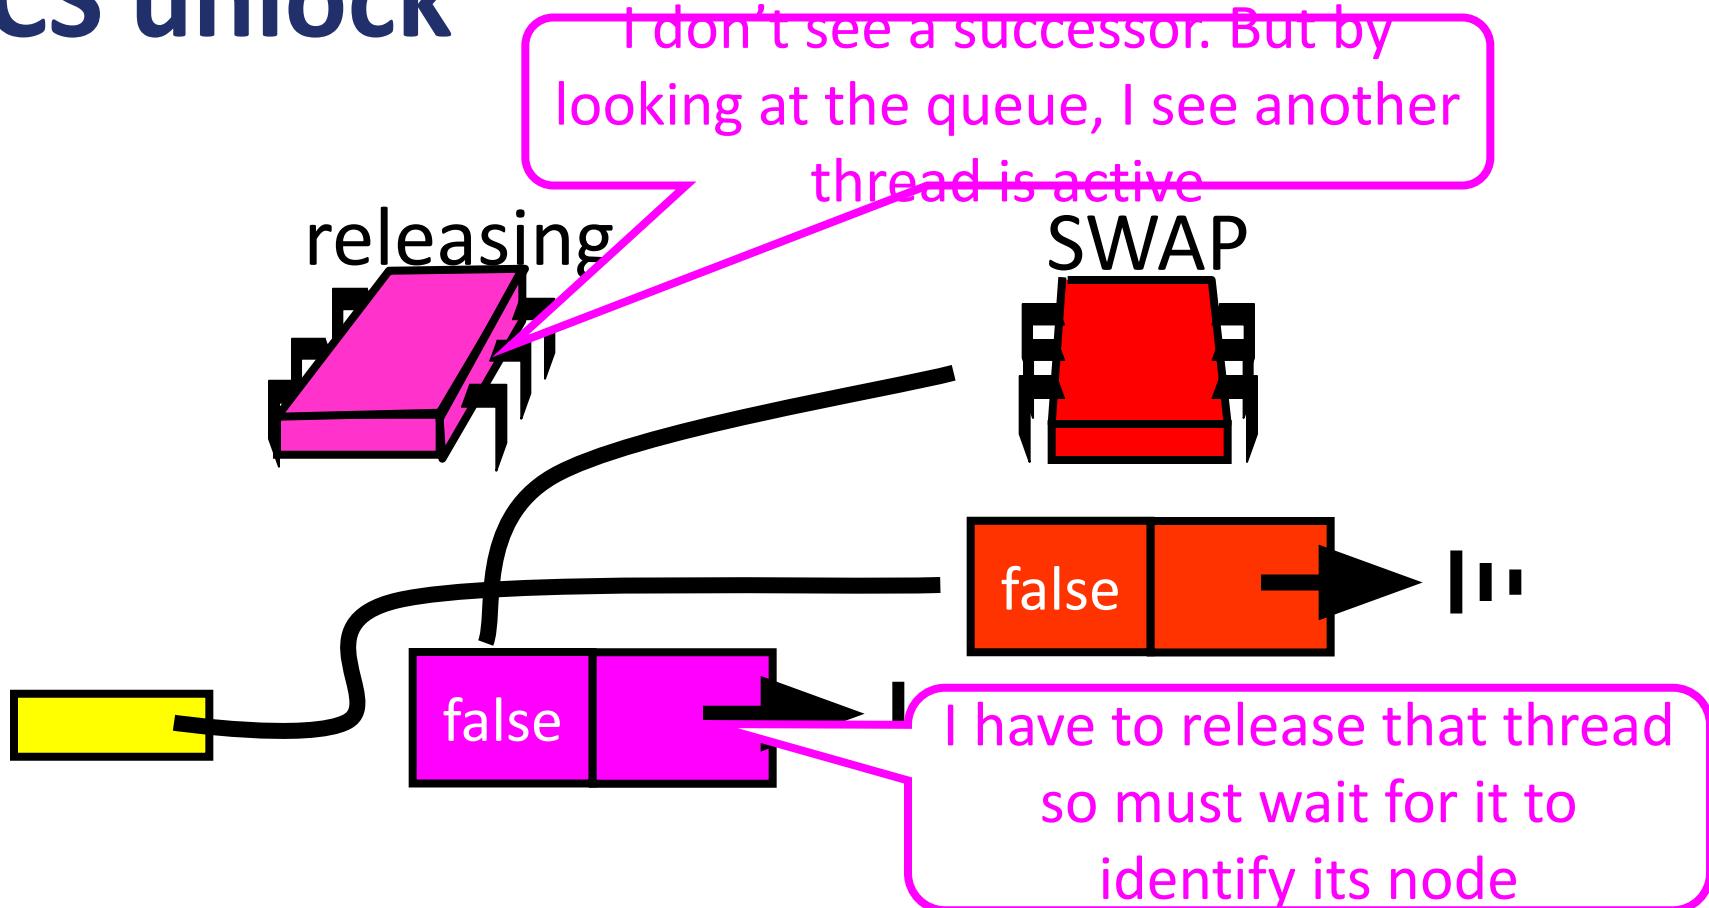
# MCS unlock



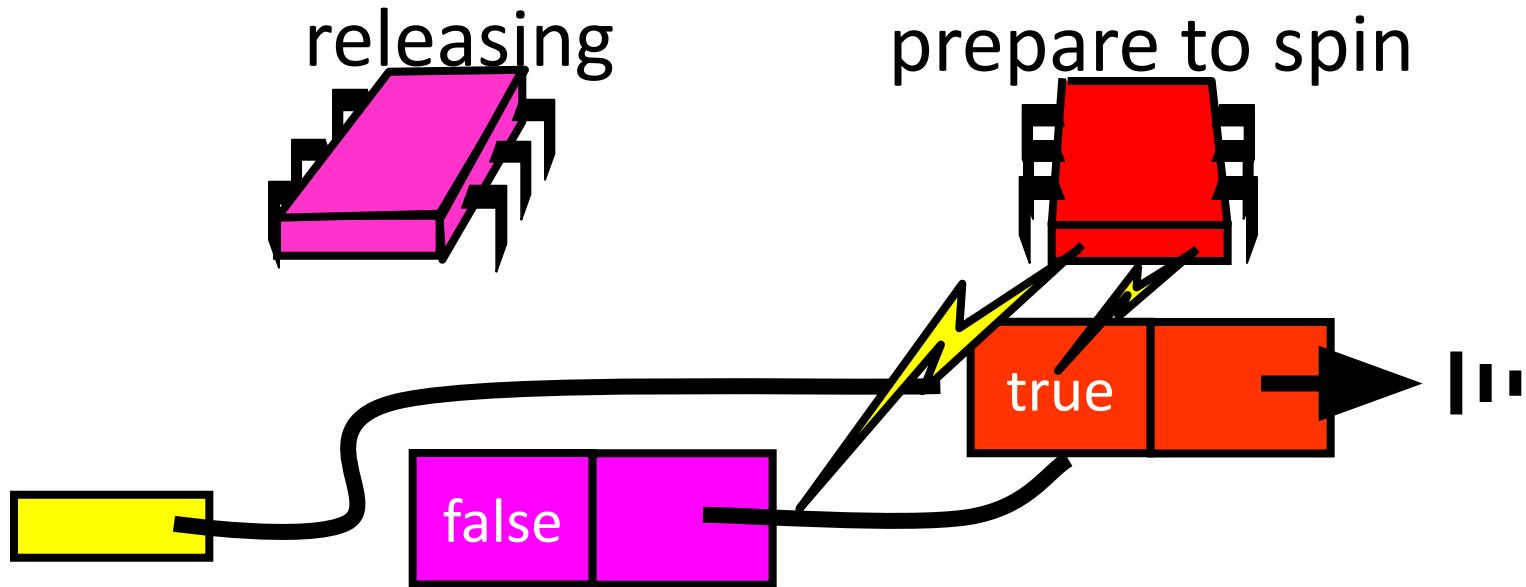
# MCS unlock



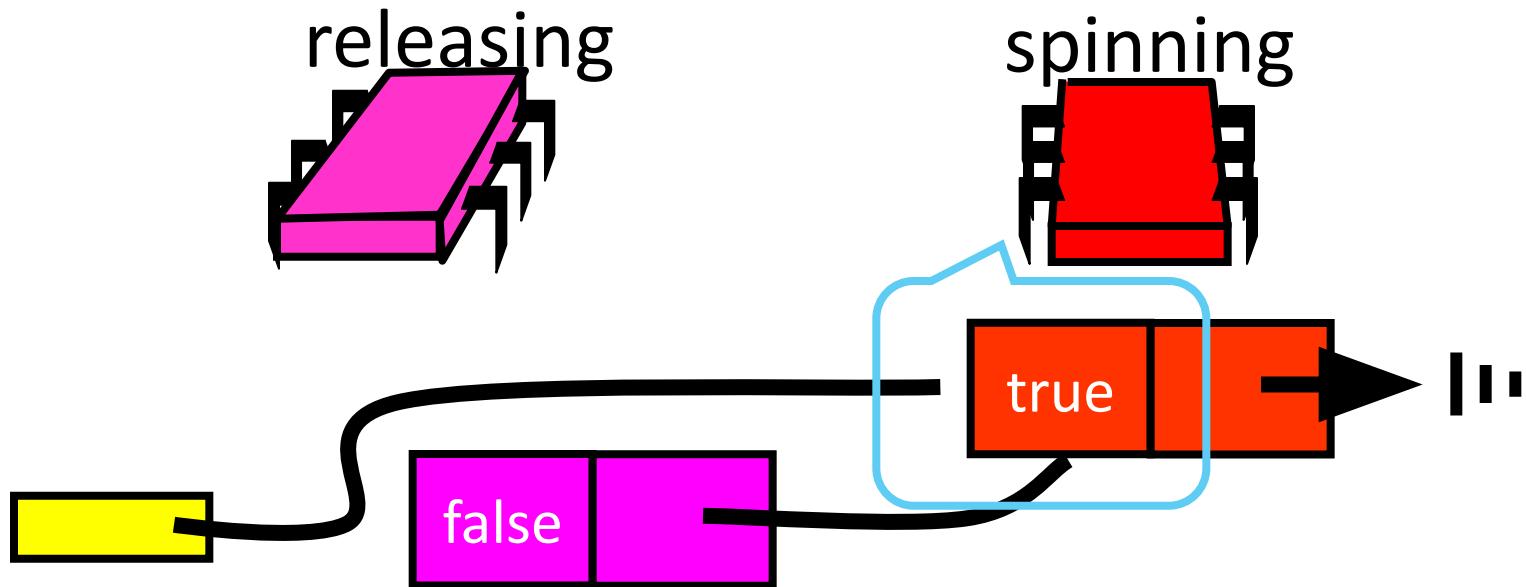
# MCS unlock



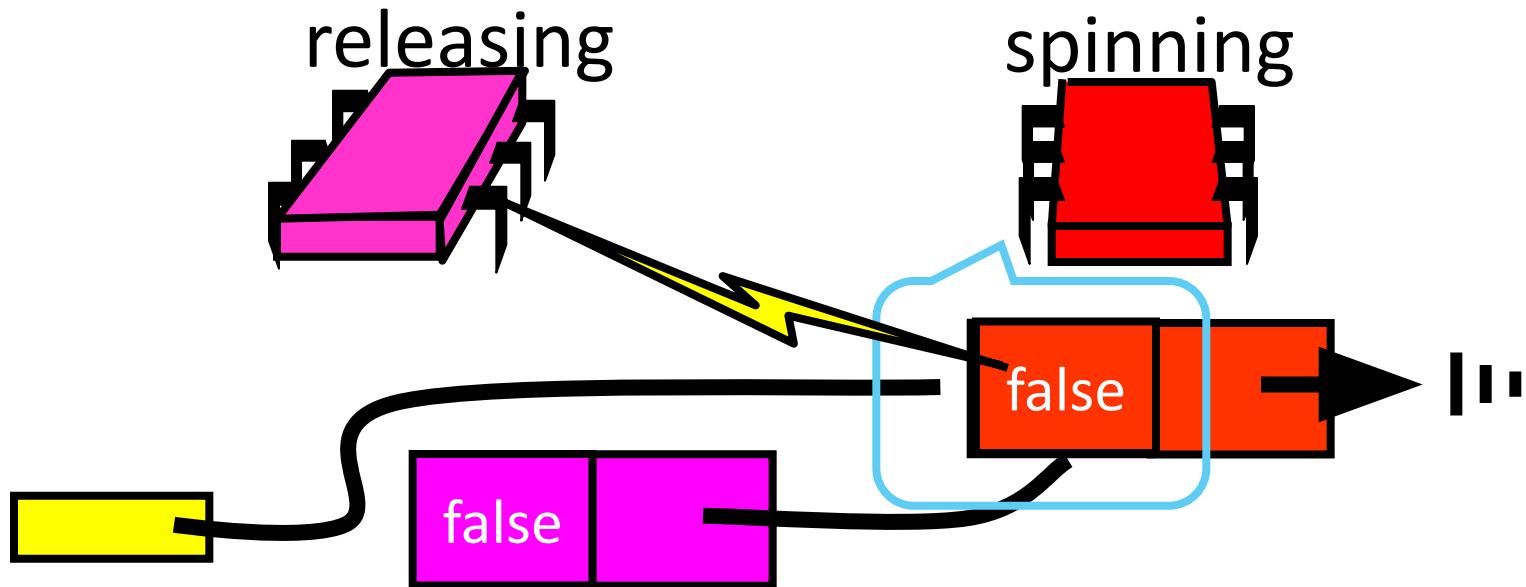
# MCS unlock



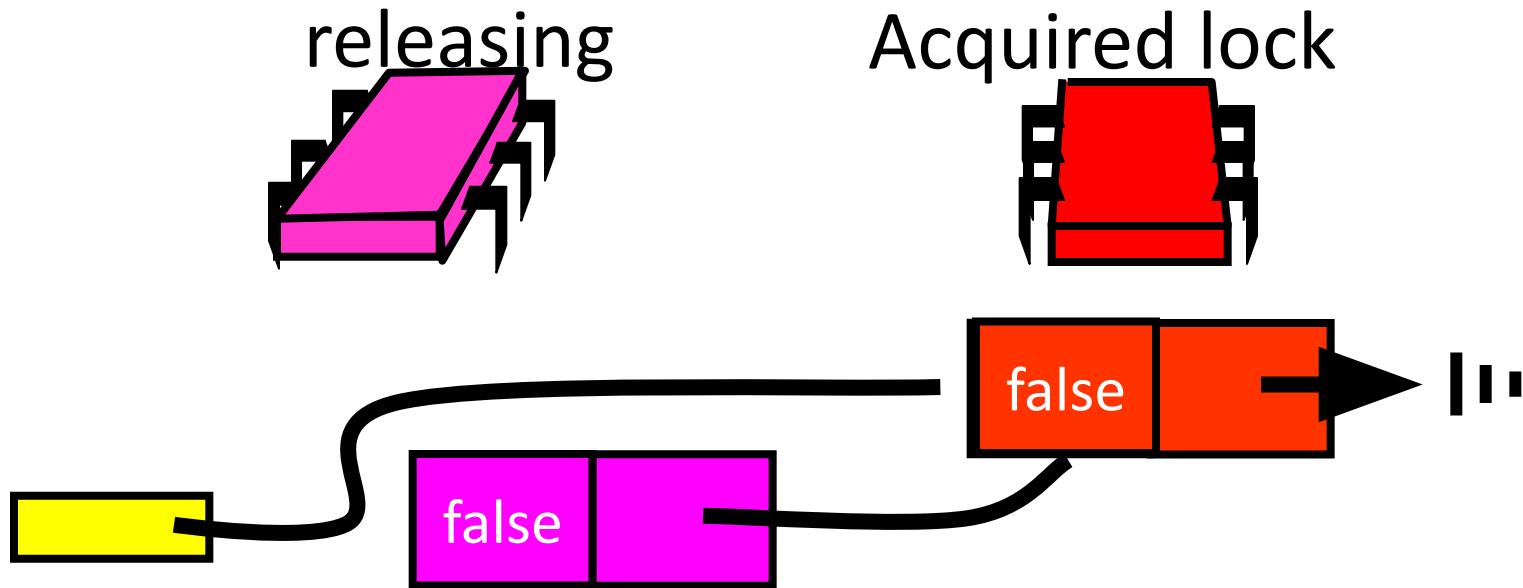
# MCS unlock



# MCS unlock



# MCS unlock



# MCS unlock

```
qnode : { qnode *next; bool lock; }
```

```
unlock(qnode *n) {
    if (n->next == NULL) {
        if (CAS(&LOCK, n, NULL)) return;
        while (n->next == NULL) { }
    }
    n->next->lock = false;
}
```

# MCS lock questions

- Why isn't the contended SWAP in lock() a problem? (Invalidation storm, etc.)
- Exercise: design a lock with queue links in reverse order
  - If B after A in queue, B spins on A's node
  - Problem with this approach: B doesn't spin on its own memory
  - Problem in **DSM** (distributed shared memory) model, where there are no caches and each

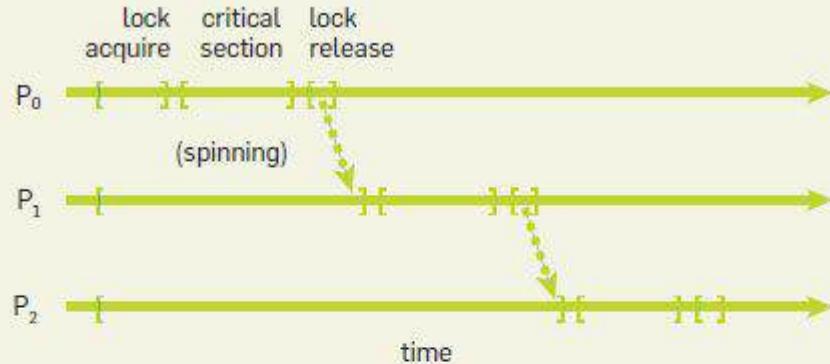
# Delegation

# Delegation

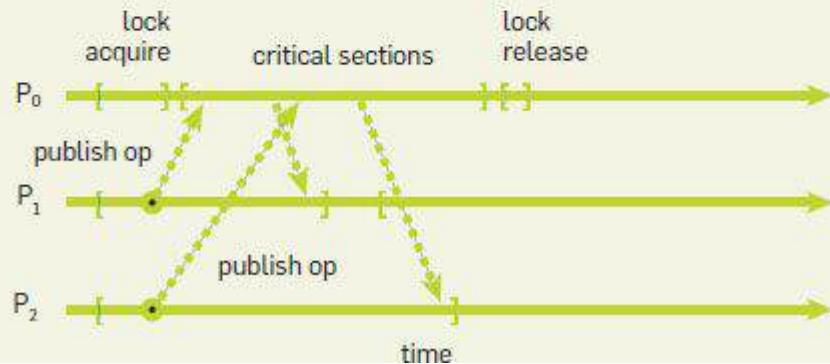
- Idea: Since CS are serialized anyway, just let a single thread run them
  - Thread that gets lock turns into a *server* and executes the operations of the waiting threads on their behalf
- Benefits:
  - Remove many lock acquisition/release from critical path
  - Might speed up CS (data hot in server's cache)
  - Enables *semantic optimizations*

# Delegation intuition

(a) queue-based lock



(b) delegation



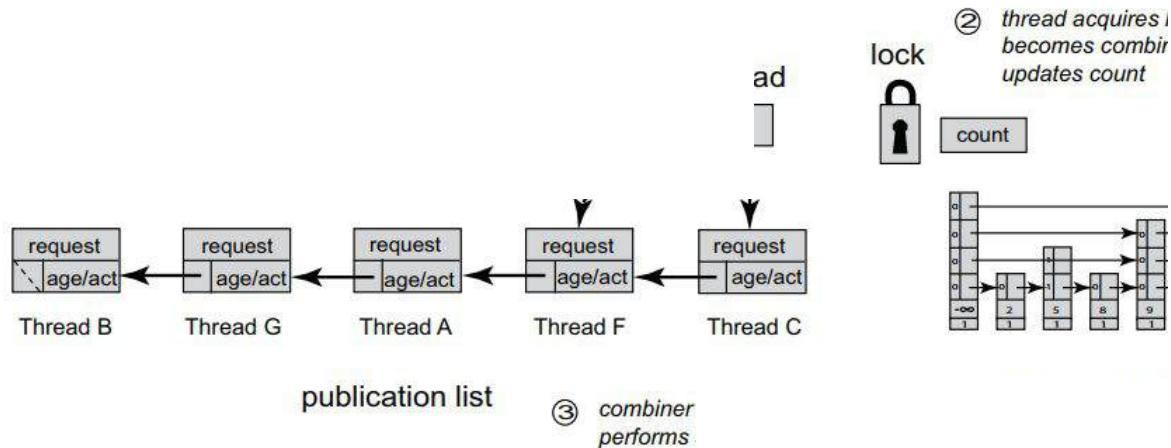
# Flat combining (FC)

[Hendler, Incze, Shavit, Tzafrir '10]

- Maintains *publication list*
  - Nodes contain *request* and *response* fields

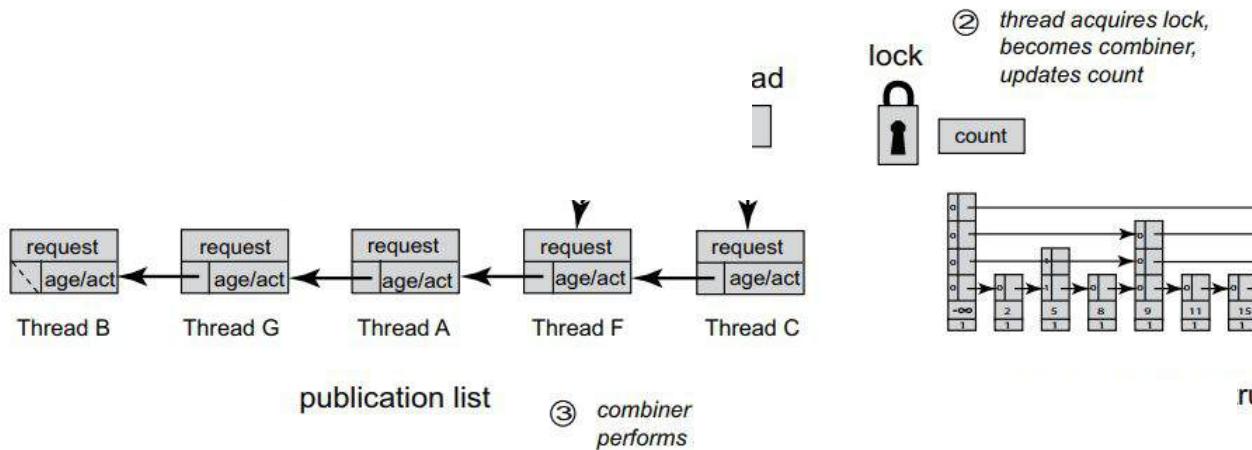
# Flat combining

- Lock:
  - 1) Add node to publication list (Treiber-stack)



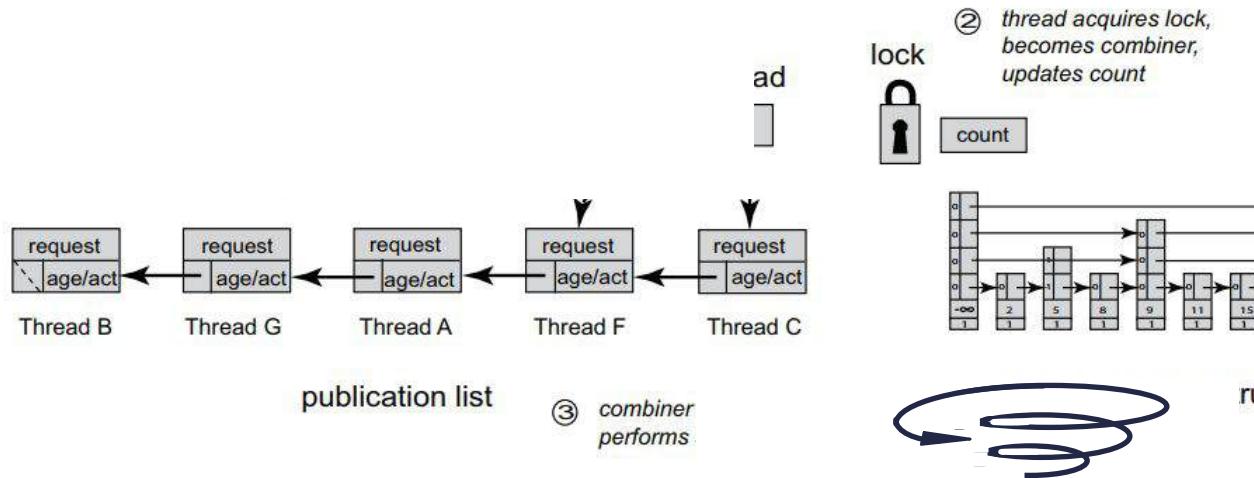
# Flat combining

- Lock:
  - 2) Check if lock taken



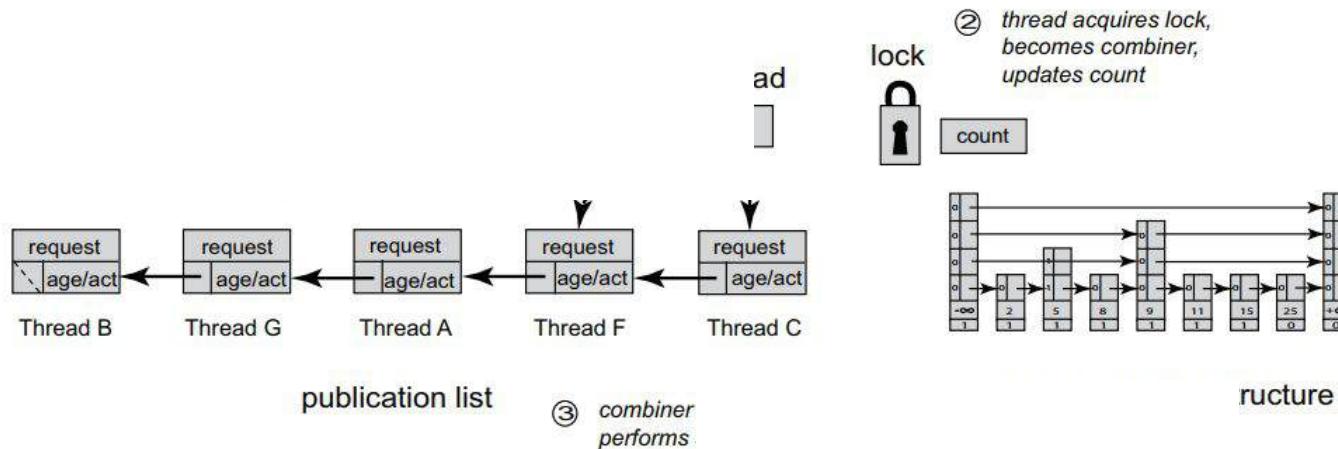
# Flat combining

- Lock:
  - 3) If lock taken, spin waiting for answer



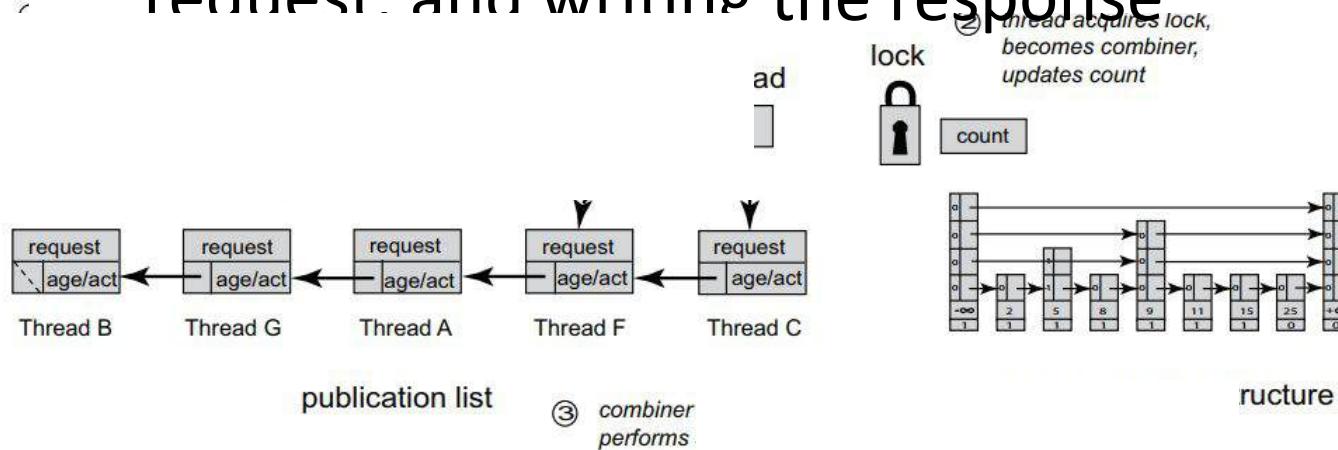
# Flat combining

- Lock:
  - 4) If lock not taken, acquire it



# Flat combining

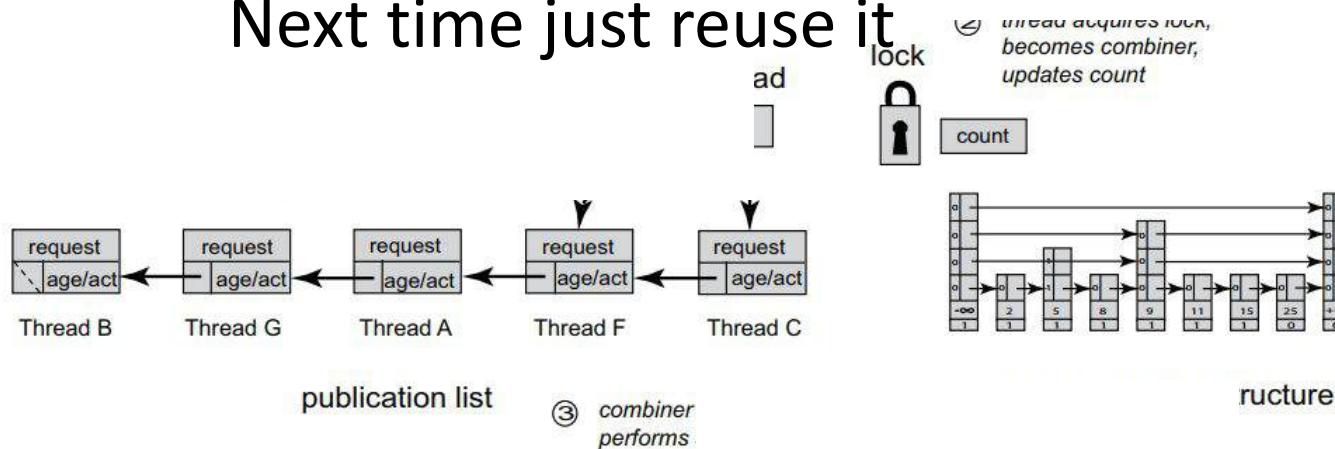
- Lock:
  - 5) Scan publication list, applying each operation request and writing the response



# Flat combining

- CAS contention on publication list head?
  - 6) When done, don't remove your record;

Next time just reuse it



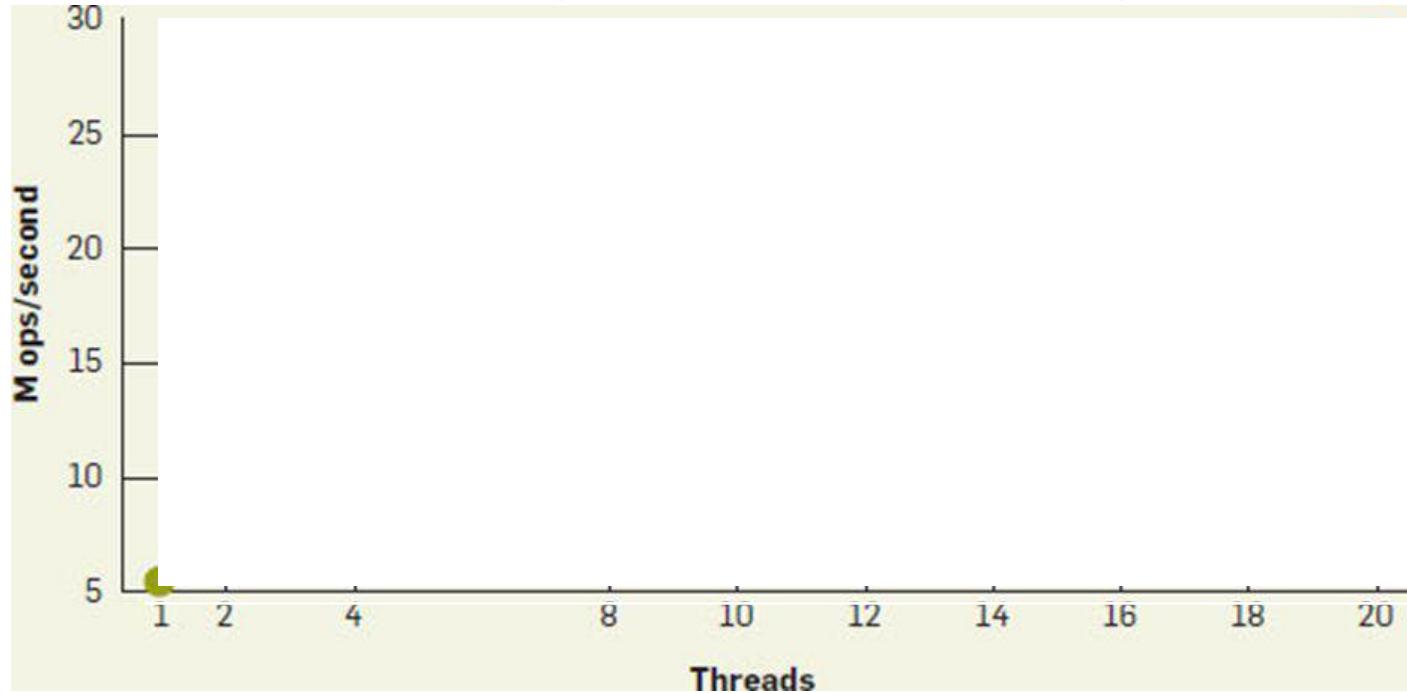
# Flat combining issues

- Starvation
  - Combiner:
    - Stop helping after  $K$  operations
    - Waiting when combiner has left:
      - Check lock every once in a while, if free try to become combiner
  - GC publication list:
    - Combiner does it every once in a while

# Two-lock queue example

- 10-core x 2 hyperthreads CPU
- Enqueue/dequeue benchmark
- Repeatedly
  - Enqueue
  - Do random work
  - Dequeue
  - Do random work
- Measure *throughput* (sum of queue ops per second over all threads)

# Two-lock queue example



# Semantic optimizations

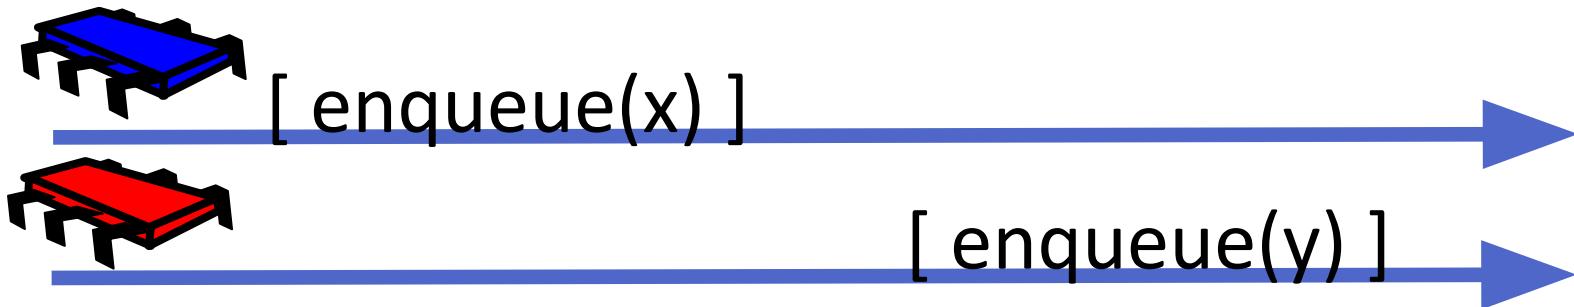
- Combiner's global view allows it to optimize by leveraging data structure semantics
- *Combining*: apply many ops as one
  - `inc()` `inc()` `inc()` -> `add(3)`
- *Elimination*: allow opposing operations to cancel out
  - `add(x)` `remove(x)` => no memory change required
- Deferring

# Deferring

- **Idea:** Operation with void response (like `enqueue()`) can deposit request *and not wait for response*
  - A.k.a. *asynchronous critical section*

# Deferring

- Tricky to do without linearizability violations



- How to maintain order if both are just records in the publication list?

# OpLog [Boyd-Wickizer, Kaashoek, Morris, Zeldovich '14]

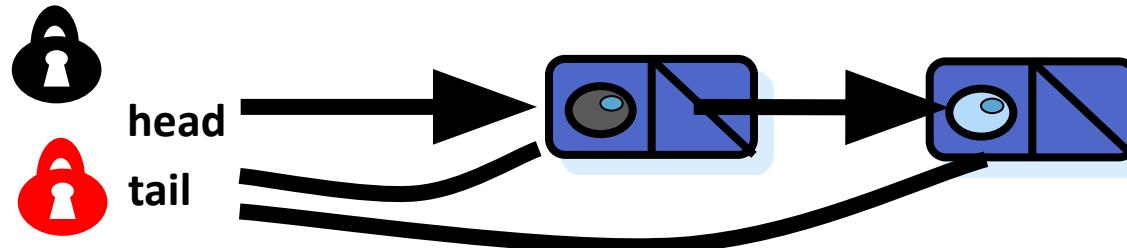
- **Idea:** Use synchronized clock (TSC, timestamp counter) to record *global time* in request, so they can be ordered
- Details in OpLog paper

# Lock-free synchronization

- *Lock-free* algorithm
  - Some operation must eventually complete
- Rules out use of locks
  - If lock holder stops taking steps, nobody can make progress
- Lock-free is often used as proxy for performance
  - Lock holders don't really die, but can be delayed
    - ⇒ If algorithm doesn't have to wait for them, it should perform better

# Lock-freeing the two-lock queue

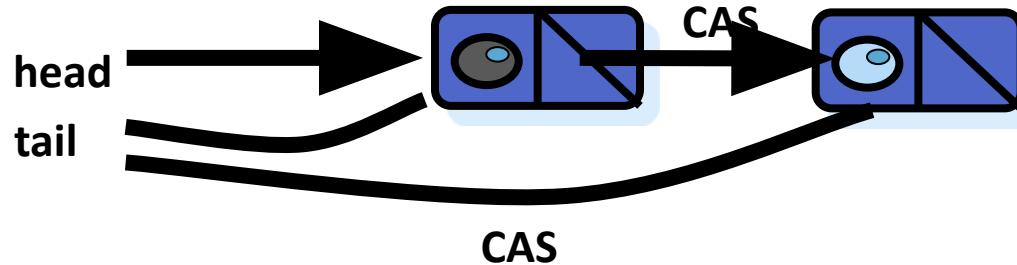
- Enqueue with locks (reminder):



- Locking  
⇒ Atomic node link + tail update

# Lock-freeing the two-lock queue

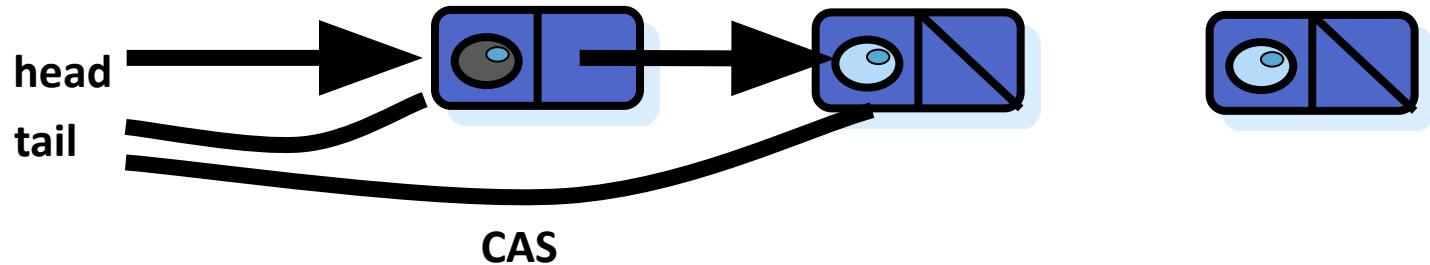
- Lock-free enqueue:



- Logical enqueue (linearization point)
  - Dequeues can now see node
- Physical enqueue

# Lock-freeing the two-lock queue

- Enqueue in inconsistent state:



- Fix it, and then retry

# MS queue [Michael & Scott '96]

```
Node* Head;  
Node* Tail;
```

Initially:

```
Head = Tail = new Node(0);  
Head->next = 0;
```

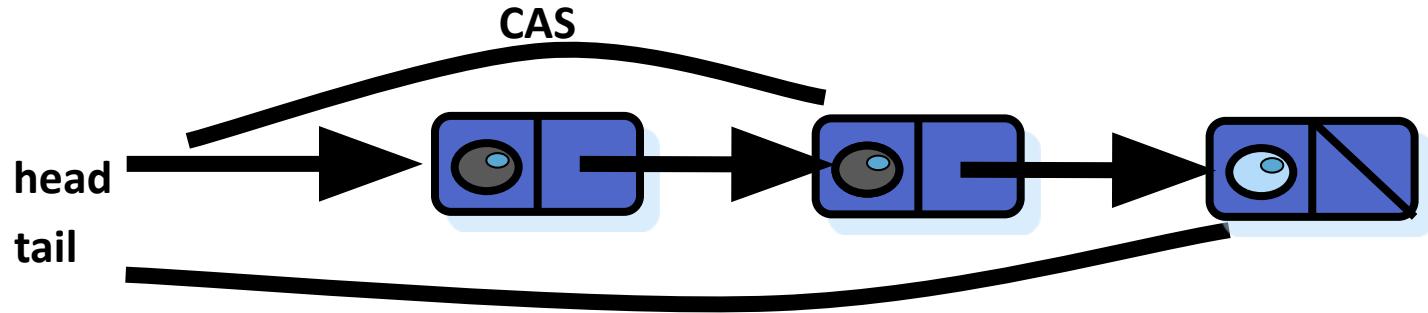
# MS queue [Michael & Scott '96]

```
void enqueue(void *v) {  
    Node *n = new Node(v); // next=NULL  
    while (1) {  
        Node *tail = Tail;  
        Next *next = tail->next;  
  
        if (next == NULL) {  
            if (CAS(&tail->next, 0, n))  
                break;  
        } else {  
            CAS(&Tail, tail, next);  
        }  
    }  
}
```

linearization point

# Lock-freeing the two-lock queue

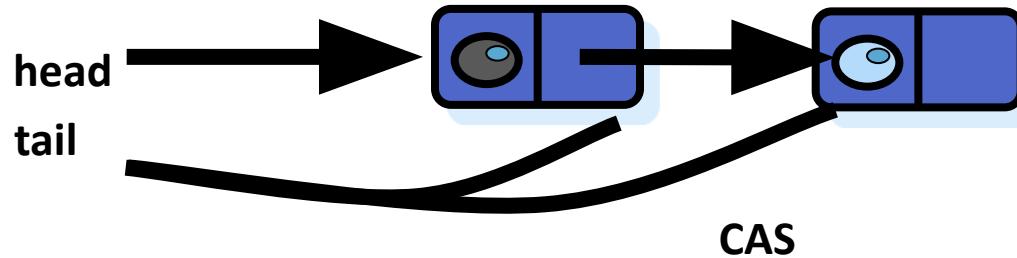
- Dequeue:



- Makes first node the new sentinel
- Old sentinel can be freed

# Lock-freeing the two-lock queue

- Dequeue:



- Also fix inconsistent states

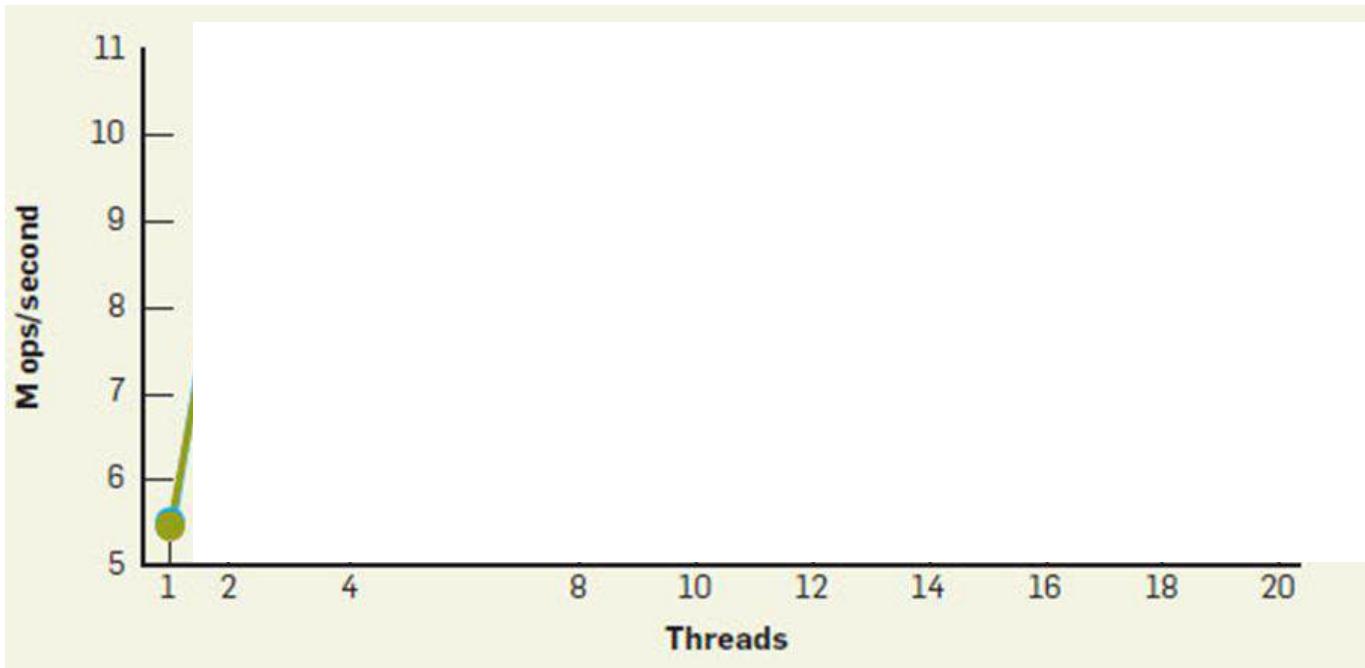
# MS queue [Michael & Scott '96]

```
void* dequeue() {  
    while (1) {  
        Node* head = Head;  
        Node* tail = Tail;  
        Next* next = head->next;  
  
        if (head == tail) {  
            if (next == NULL) return EMPTY;  
            CAS(&Tail, tail, next);  
        } else {  
            void* rv = next->value;  
            if (CAS(&Head, head, next)) return rv;  
        }  
    }  
}
```

Is this necessary?



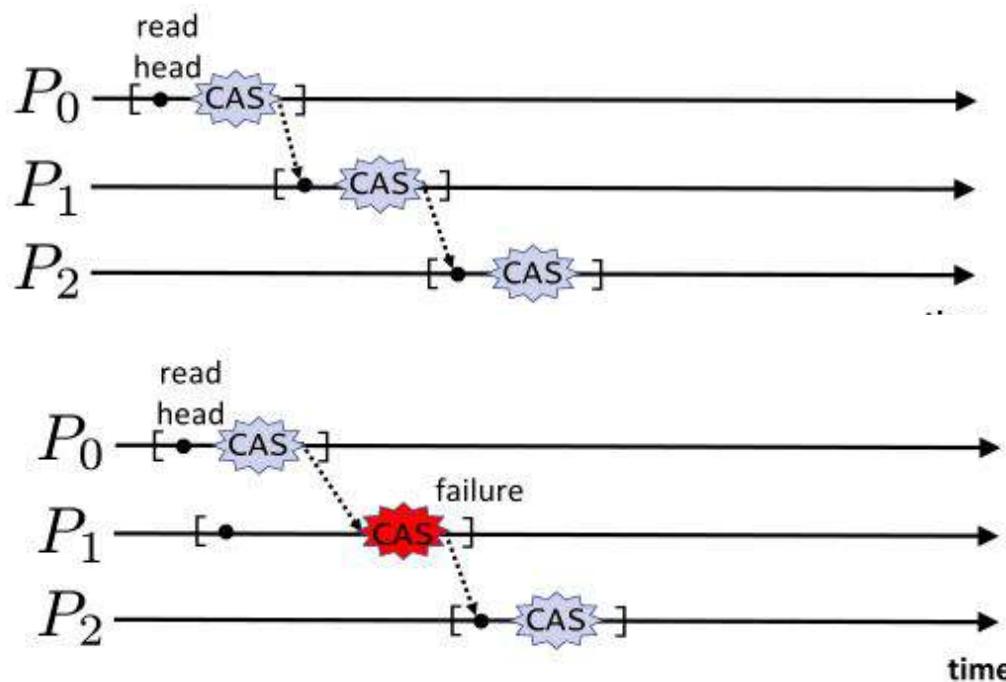
# Lock-freedom to the rescue?



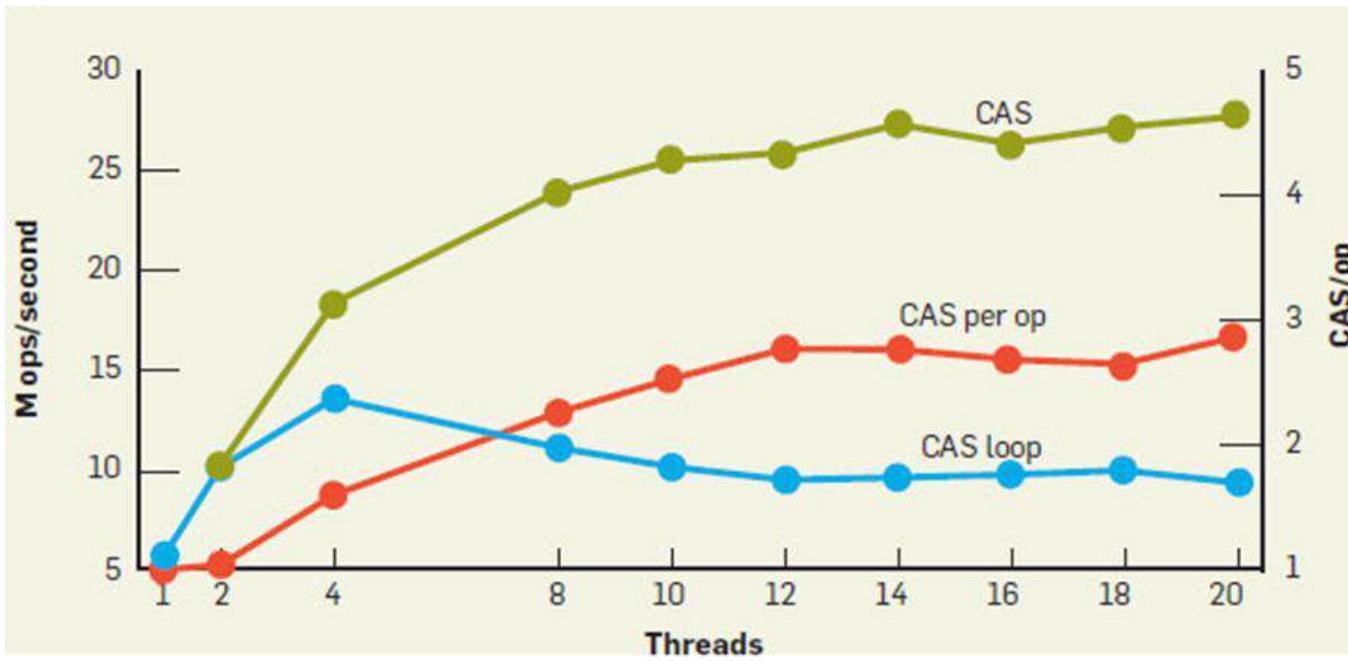
# CAS failures

- MS queue poor performance is due to *CAS failures*, which add *useless work* to the critical path

# CAS failures



# CAS failures



# Avoiding CAS failures

- Idea: Try to replace CAS with an atomic instruction that doesn't fail, so that every atomic op contributes to useful work
  - SWAP, Fetch-And-Add

# FAA queue

```
cell : { void *value; }
```

```
cell Q[] =  $\perp$  .....; uint head = 0, tail = 0;
```

```
enqueue(x) {
```

```
    while (true) {
```

```
        t = F&A(&tail,  $\perp$ )
```

```
        if (CAS(&Q[t],  $\perp$ , x))
```

```
            return
```

```
}
```

```
}
```

```
dequeue(x) {
```

```
    while (true) {
```

```
        h = F&A(&head,  $\perp$ ) $\top$ 
```

```
        if (!CAS(&Q[h],  $\perp$ ,  $\perp$ ))
```

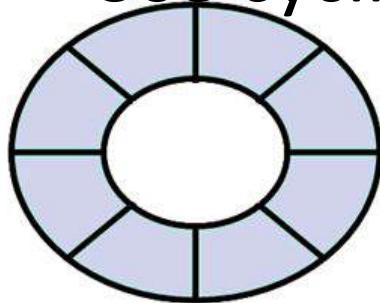
```
            return Q[h]
```

```
        if (tail  $\leq$  h+1) return EMPTY
```

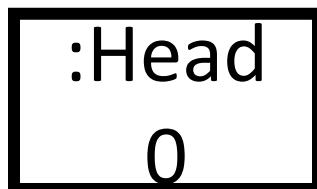
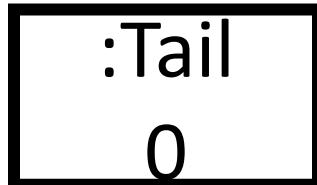
```
} }
```

# Bounding the queue

- Use cyclic array

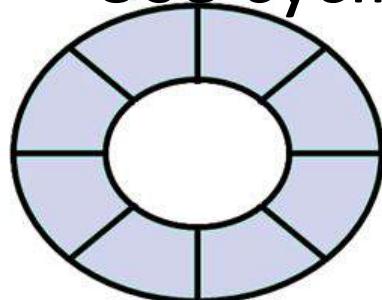


$R$  nodes

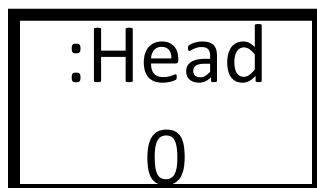
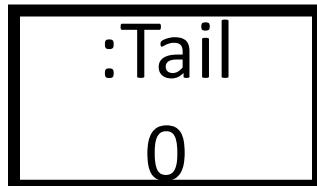
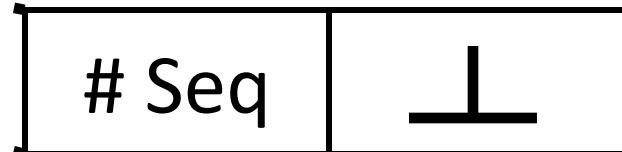


# Bounding the queue

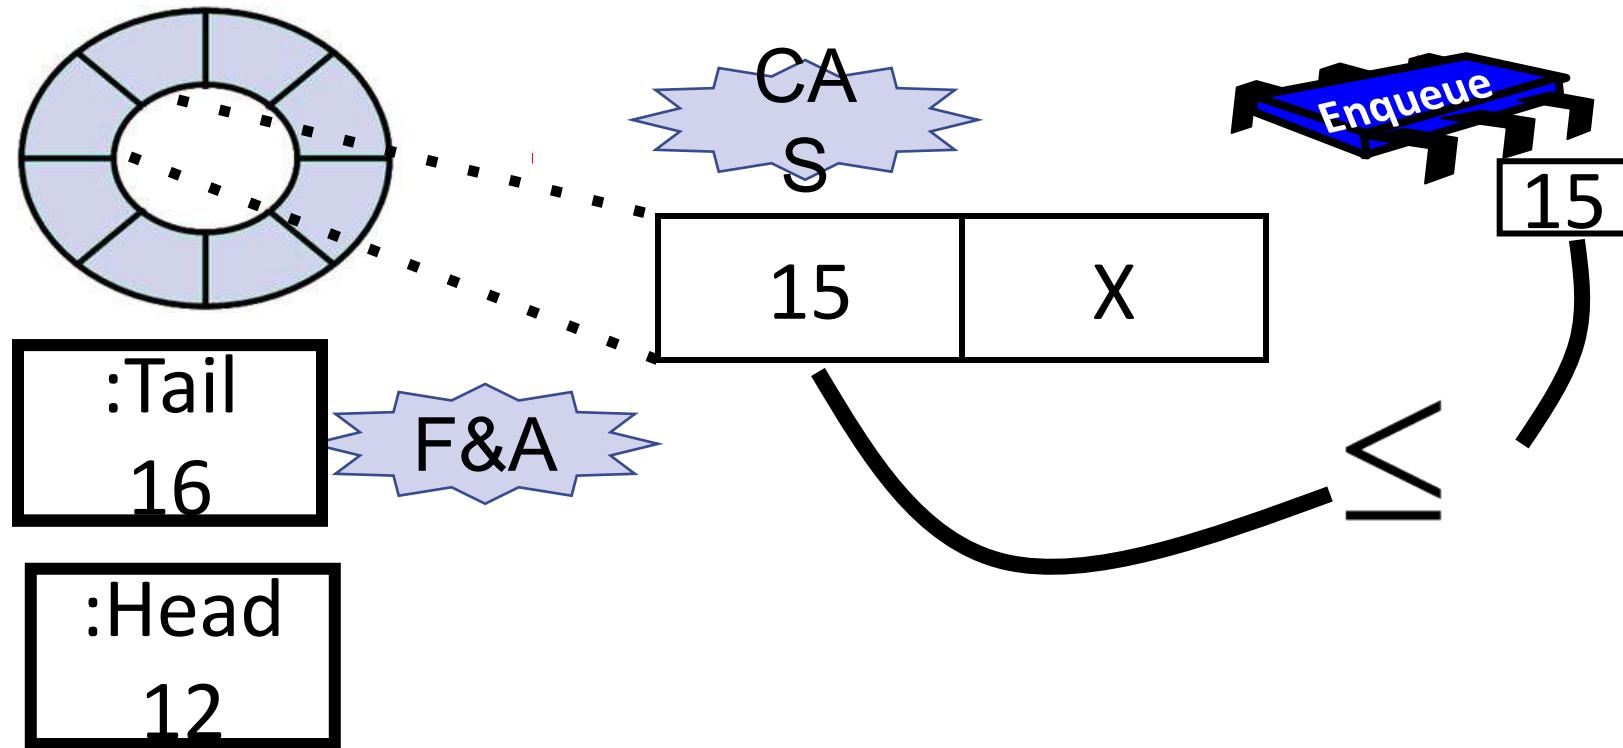
- Use cyclic array



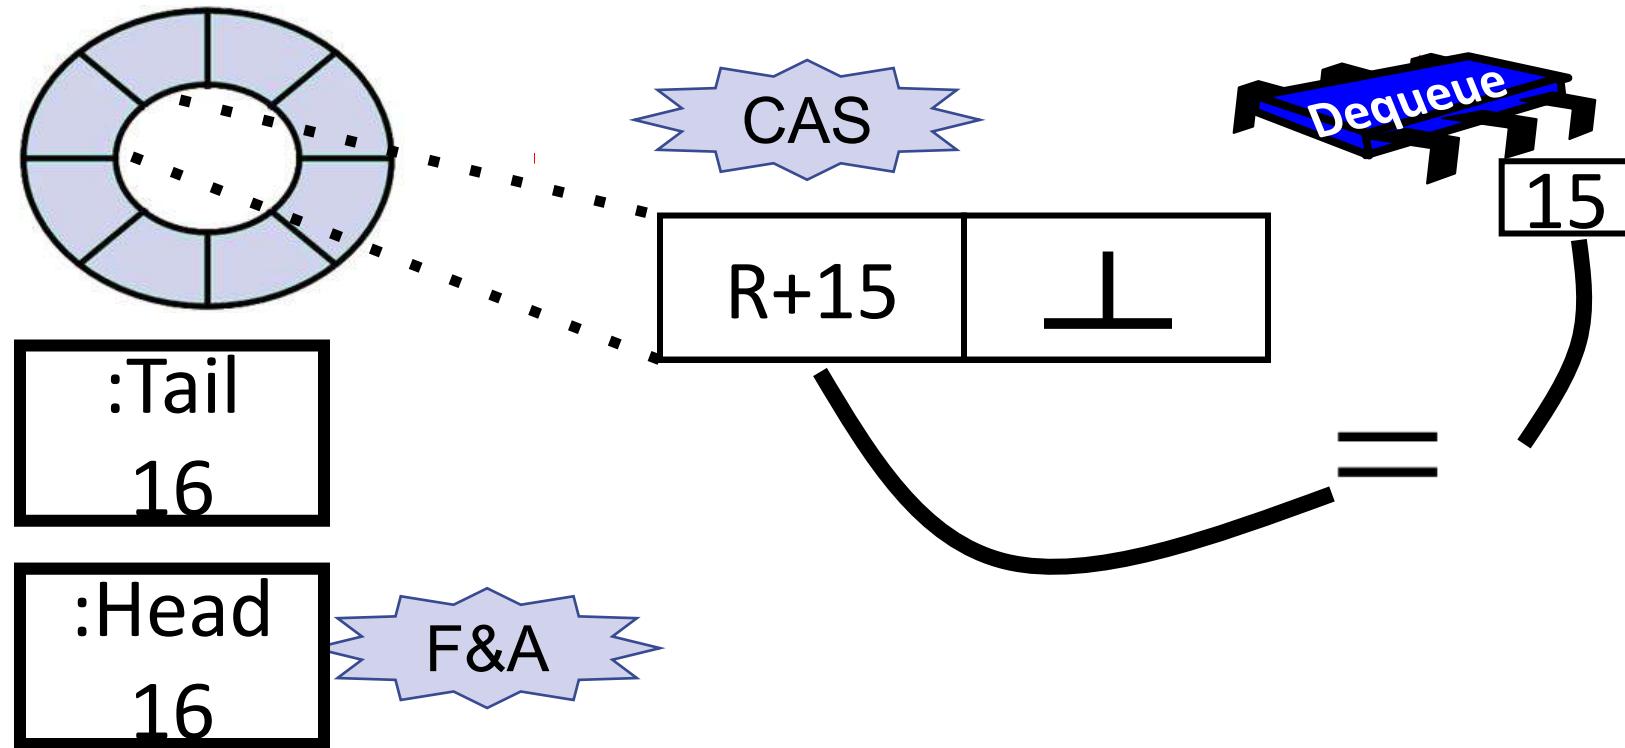
$R$  nodes



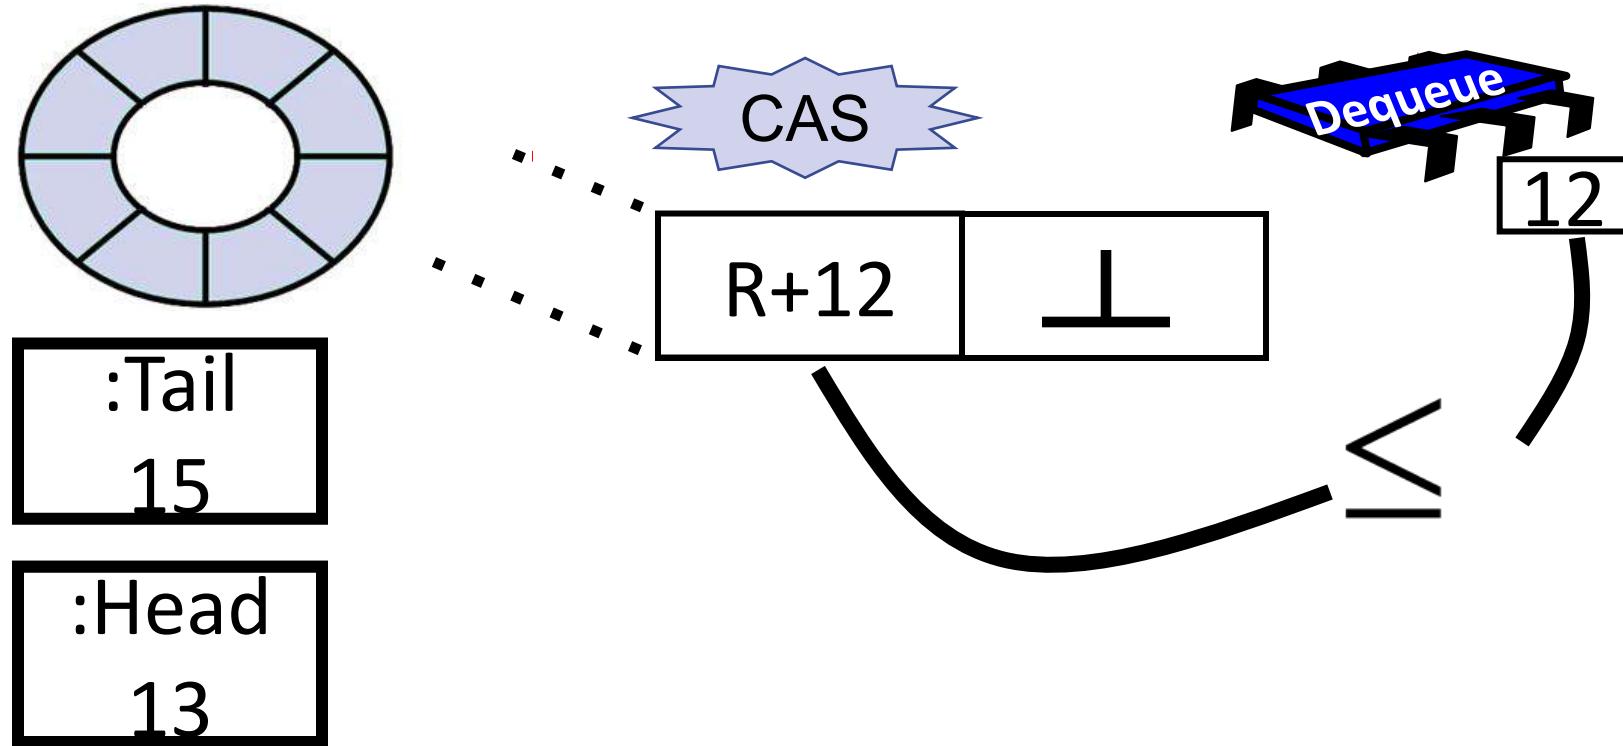
# Enqueue



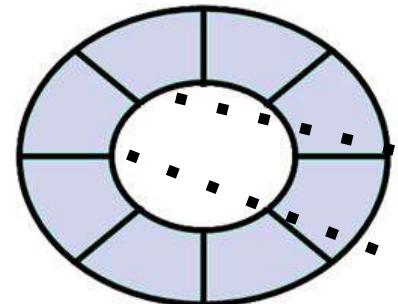
# Dequeue returning a value



# Dequeue before enqueue

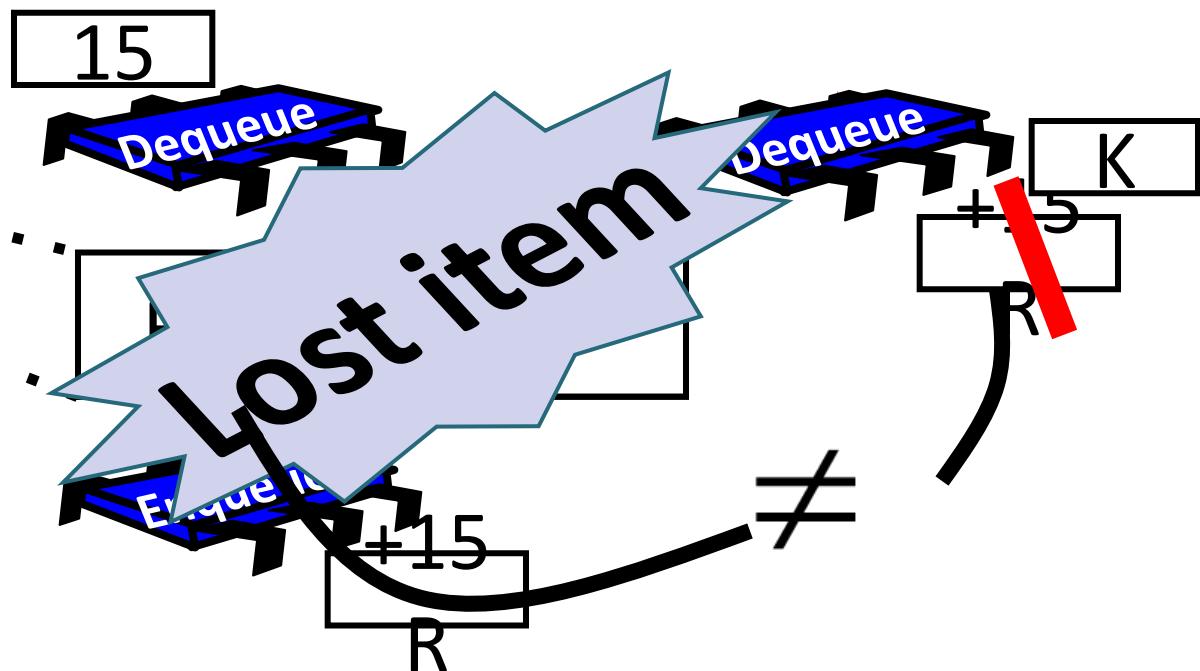


# Deq/Enq mismatch



:Tail  
R+16

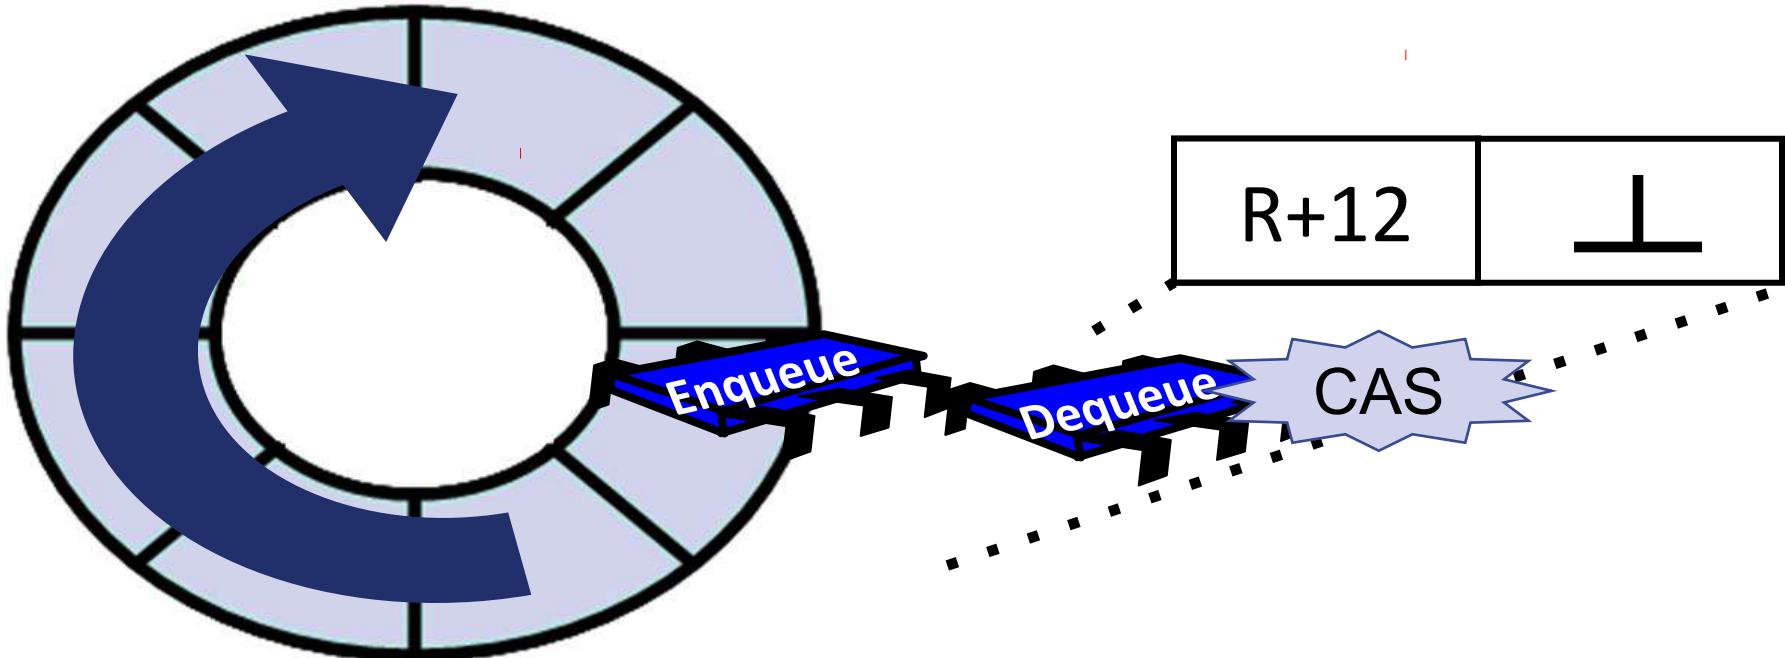
:Head  
K+1



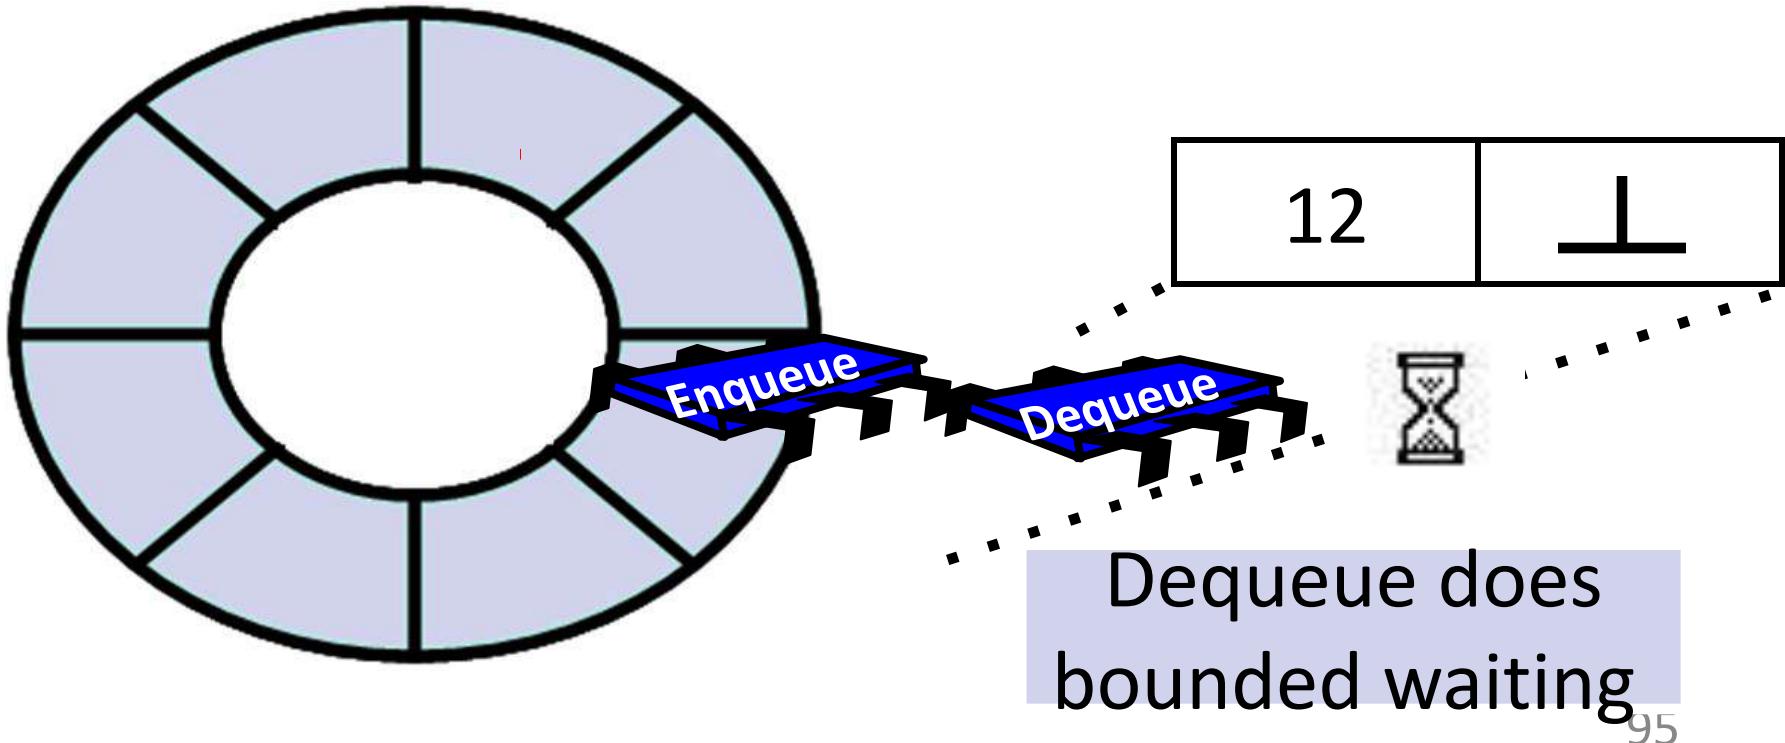
# Deq/Enq mismatch

- Simple solution: if dequeue with index  $y$  sees  $\langle val, x \rangle$  for  $x < y$  in its cell, it waits:
  - Eventually, dequeuer with index  $x$  will remove  $val$
- But this isn't lock-free
- Lock-free solution more complex, see the paper

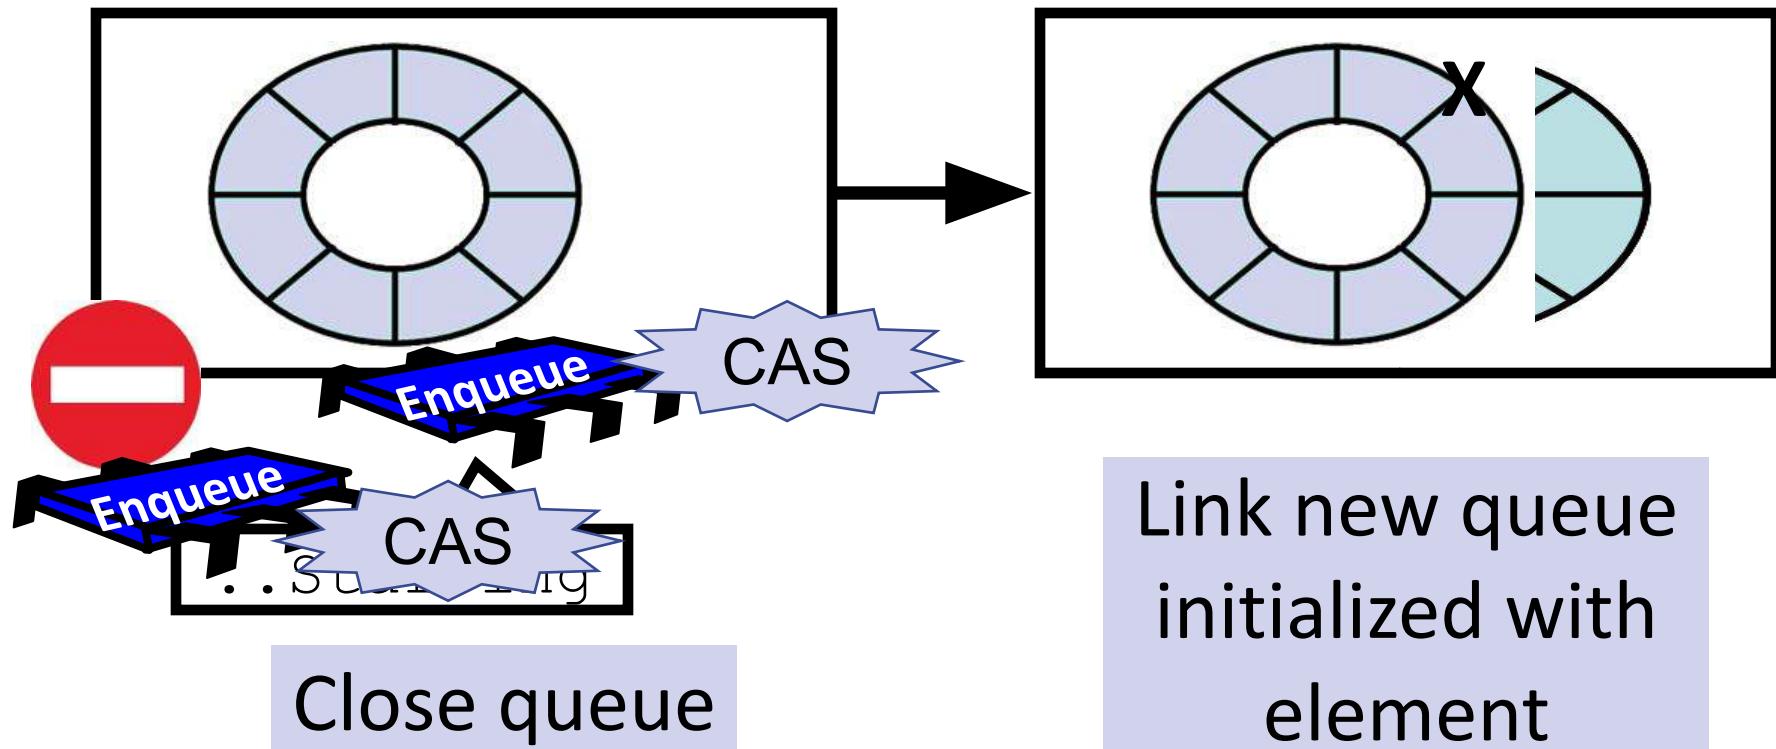
# Algorithm livelocks



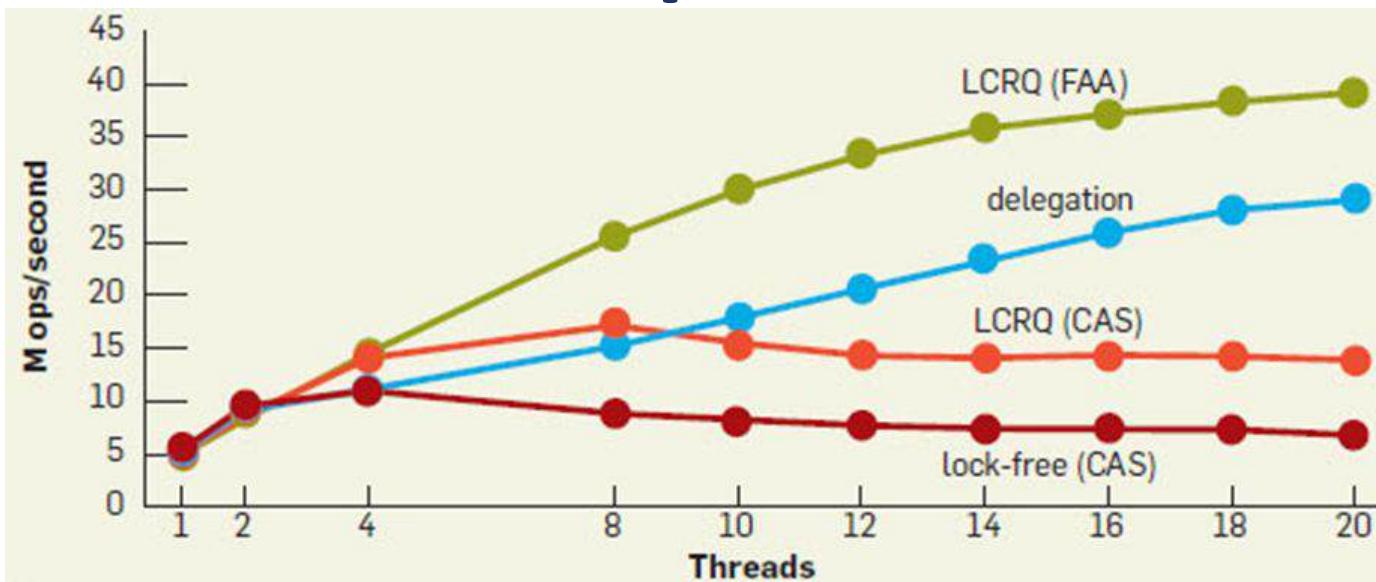
# Practical livelock solution



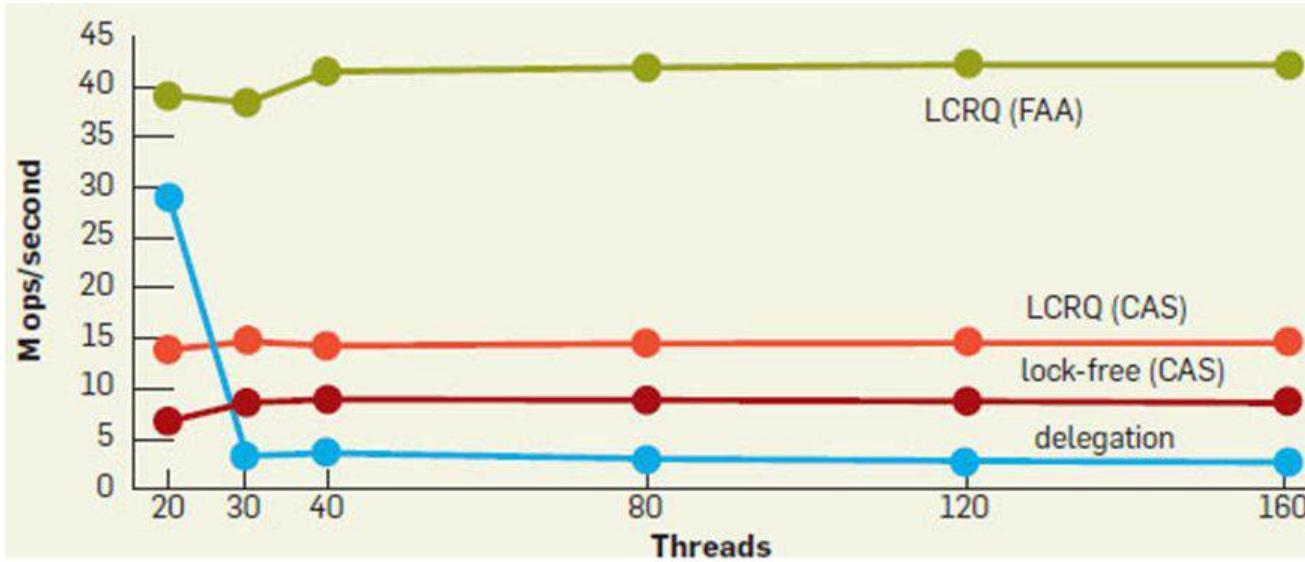
# Solution: list of nodes



# Queue example



# Queue with preemption



# Summary

- Importance of serializing efficiently
- Efficient lock algorithms
- Delegation + optimizations
- Avoiding CAS failures in lock-free algorithms